How Does the Pretraining Distribution Shape In-Context Learning? Task Selection, Generalization, and Robustness

Waïss Azizian*1 and Ali Hasan²

¹Univ. Grenoble Alpes, CNRS, Inria, Grenoble INP, LJK, 38000 Grenoble, France

²Machine Learning Research, Morgan Stanley, New York, USA

waiss.azizian@univ-grenoble-alpes.fr ali.hasan@morganstanley.com

Abstract

The emergence of in-context learning (ICL) in large language models (LLMs) remains poorly understood despite its consistent effectiveness, enabling models to adapt to new tasks from only a handful of examples. To clarify and improve these capabilities, we characterize how the statistical properties of the pretraining distribution (e.g., tail behavior, coverage) shape ICL on numerical tasks. We develop a theoretical framework that unifies task selection and generalization, extending and sharpening earlier results, and show how distributional properties govern sample efficiency, task retrieval, and robustness. To this end, we generalize Bayesian posterior consistency and concentration results to heavy-tailed priors and dependent sequences, better reflecting the structure of LLM pretraining data. We then empirically study how ICL performance varies with the pretraining distribution on challenging tasks such as stochastic differential equations and stochastic processes with memory. Together, these findings suggest that controlling key statistical properties of the pretraining distribution is essential for building ICL-capable and reliable LLMs.

1 Introduction

In-context learning (ICL) is the phenomenon whereby a model generalizes to a new task from a handful of examples provided in the input context without any model weight updates. This emergent behavior has been observed across models in multiple domains, including in language (Brown et al., 2020), vision (Radford et al., 2021), and reinforcement learning (Moeini et al., 2025). ICL is a particularly appealing feature in domains where data for a specific task is scarce such as robotics (Ahn et al., 2023b), healthcare (Singhal et al., 2023), or chemistry (Stokes et al., 2020).

Despite growing interest, the conditions under which ICL emerges are still poorly understood. Several lines of works have emerged to address this question. The algorithmic view focuses on studying which learning algorithms over the context can be implemented by transformer and thereby perform ICL (Garg et al., 2022; Akyürek et al., 2023). Others have suggested modeling ICL as Bayesian inference (Xie et al., 2021; Lin and Lee, 2024; Zhang et al., 2025b; Jeon et al., 2024). Empirical works have sought to design controlled settings in which ICL can be carefully studied, and these works highlight how sensitive to pretraining choices ICL is (Chan et al., 2022; Raventós et al., 2023), indicating that distributional aspects of pretraining play a central role. A crucial line of work also seeks to assess ICL performance on numerical tasks through out-of-distribution robustness of ICL (Wang et al., 2025b; Kwon et al., 2025; Goddard et al., 2025) but its behavior remains poorly understood.

Yet existing modeling frameworks often focus on restricted settings and lack general tools that *links properties* of the pretraining distribution to ICL behavior at test time. Three aspects remain particularly underexplored: (i) heavy-tailed distributions that better reflect real-world pretraining corpora and have been identified as key drivers of ICL (Chan et al., 2022; Singh et al., 2023), (ii) non-i.i.d. and dependent structures (e.g., long-range dependencies in language sequences) that fall outside standard i.i.d. or Markovian ICL modeling (Alabdulmohsin et al., 2024), and (iii) how these distributional properties govern the robustness of ICL under shifts at test time, which is a key feature of ICL (Wang et al., 2025b; Kwon et al., 2025; Goddard et al., 2025).

^{*}Work done during an internship at Morgan Stanley Machine Learning Research.

We thus develop a study of ICL with a focus on the influence of the pretraining distribution. We decompose ICL performance into two components: *task selection* (identifying the right task from the context) and *generalization* (performing well on tasks and sequences unseen during training) and focus on the following questions:

How does the pre-training distribution shape ICL performance on new tasks? How does it affect task selection and generalization errors?

Our contributions are as follows:

- Framework. We develop a general theoretical framework for ICL that focuses on the role of pretraining distributional properties, handling both the task selection error and the ICL generalization error.
- Theory under heavy tails and dependence. We extend Bayesian consistency and concentration guarantees to heavy-tailed priors and dependent sequences, providing conditions that better reflect pretraining data used for LLMs and highlighting the role of these key distributional properties.
- Empirical validation on numerical tasks. We validate the framework on challenging numerical tasks—including stochastic differential equations and processes with memory, assessing ICL via robustness to new tasks and distribution shift, and finding outcomes consistent with our theory.

Together, our results suggest that controlling key statistical properties of the pretraining distribution is essential for building ICL-capable and reliable transformer models.

2 Related Work

A number of works study ICL through varying perspectives and definitions of ICL. We will focus on the perspectives most relevant to what we study.

Conditions for ICL. Other works devoted to studying the conditions under which ICL occurs. From a pre-training perspective, Chan et al. (2022) studied the distribution qualities of a pretraining distribution that leads to ICL while Raventós et al. (2023) studied the influence of regularization and training distribution on linear regression tasks. However, these do not consider a unified theory for predicting how ICL behaves under a particular pre-training distribution and only consider a limited class of experiments. Singh et al. (2023) showed that ICL is transient and conditions must be carefully chosen such that the model performs ICL rather than in-weight learning.

Bayesian Perspectives. From a statistical perspective, a series of questions were raised as to how ICL can be studied through a Bayesian framework where the pre-training distribution acts as a prior. Xie et al. (2021) proposed viewing ICL as Bayesian model averaging. Lin and Lee (2024) studied how ICL involves two modes of operation where one case the model generalizes and the other case the model retrieves similar tasks. Zhang et al. (2025b) considered a theory for a Bayesian perspective of ICL and provided error bounds on the task loss as a function of the number of tasks and the number of points within each task. However, they do not study specific properties of the pre-training distribution that lead to good ICL performance. Jeon et al. (2024) provide an information theoretic perspective on task retrieval for ICL but do not model the distribution of tasks.

Generalization. Several works have studied the generalization properties of ICL. Li et al. (2023) obtain such results by studying the stability of the transformer architecture but they consider the same fixed and finite task distribution during both pre-training and testing. Zhang et al. (2025b); Zekri et al. (2024) both provide generalization bounds for ICL on Markov chains but without modelling the distribution of tasks during pre-training, which is our focus here. Lotfi et al. (2024) provide generalization bounds for transformers on arbitrary sequences but with a restrictive notion of generalization that does not capture the ICL setting.

Numerical Tasks. Related to the experiments we consider is a line of work studies ICL on small transformer models and simple tasks. Zhang et al. (2024); Wu et al. (2024) study ICL on linear regression tasks with a single-linear attention model, characterizing the ICL error of the trained model and the sample complexity of learning ICL. Chan et al. (2025) study a simple model of a Bayesian predictor to understand the different modes of in-weight learning and ICL.

Algorithms and Out of Distribution. Several works focus on the training dynamics of transformers for ICL, as well as how the transformer architecture is expressive enough to implement a wide variety of algorithms for ICL. This is an important and desirable quality since it would allow for generalization across out of distribution tasks. Wang et al. (2025b); Kwon et al. (2025); Goddard et al. (2025) all study this question from different perspectives and ultimately conclude that certain conditions on the pretraining distribution allow for some level of out of distribution performance. We defer a more detailed review of these works to Appendix A.

General Concentration Results. Finally, we briefly review relevant concentration results. The pioneering work of Yu (1994) provides concentration inequalities for dependent processes with a total variation condition, opening up a fruitful line of research, see e.g., Kontorovich and Ramanan (2008); Mohri and Rostamizadeh (2008, 2010); Maurer (2023); Abélès et al. (2025) and, for related coupling techniques, see (Chazottes et al., 2007; Paulin, 2015), as well as references therein. Though these frameworks can handle non-linear functions of dependent sequences, they require boundedness assumptions that are not suitable for our setting. Another line of work has studied so-called functional dependence conditions (Wu, 2005, 2011) and provided concentration inequalities for sums of stationary dependent sequences (Liu et al., 2013). However, our ICL setting requires concentration inequalities for more general function classes and non-stationary sequences, which to the best of our knowledge are not available in the literature. Concerning heavy-tailed concentration bounds, we refer to the recent frameworks of Bakhshizadeh et al. (2023); Li and Liu (2024b); Li et al. (2024); Li and Liu (2024b) which provide concentration inequalities for non-linear functions of independent heavy-tailed random variables and which we extend to the dependent setting.

3 Theoretical framework

3.1 In-context learning setting

In line with existing ICL works, we model the training data as a mixture of tasks, with each task defining its own distribution. Formally, denote by $\Theta \subset \mathbb{R}^d$ the space of tasks θ and by $\pi(\theta)$ the density of the pretraining task distribution. Given a task θ , the data is generated according to a task-specific distribution with density $p(\cdot | \theta)$ The training data is then generated by first sampling a task θ from the task distribution π , and then sampling data points $(x_t)_{t\geq 1}$ according to

$$x_{t+1} \sim p_{t+1}(\cdot | x_{1:t}, \theta)$$
, where $x_{1:t} = (x_1, \dots, x_t)$.

We first present some running examples to illustrate the setting.

Example 3.1 (Classification). Several ICL benchmarks for LLMs such as Bertsch et al. (2025); Zou et al. (2025); Li et al. (2025b) are built on classification tasks. Each task θ represents a small subset of classes from a larger classification problem and the data sequence x_1, \ldots, x_t is a sequence of inputs and labels from these classes. The challenge is therefore to both identify the classes and learn to classify them from the in-context examples.

Example 3.2 (Linear Regression). Introduced by Garg et al. (2022), the regression setting is a popular testbed for ICL. Each task $\theta \in \mathbb{R}^d$ defines a linear model $y = \theta^T q + \epsilon$ where ϵ is some noise. The data sequence x_1, \ldots, x_{2t} is a sequence of input-output pairs $q_1, y_1, \ldots, q_t, y_t$ generated according to the linear model defined by θ .

Example 3.3 (Next-sample prediction for stochastic processes). More generally, we can consider the setting where each task θ defines a stochastic process $x_{t+1} \sim p_{t+1}(\cdot | x_{1:t}, \theta)$. We will consider later the specific case of the Ornstein-Uhlenbeck process: each task $\theta = (\tau, \mu)$ defines a mean-reverting stochastic process with mean μ and reversion speed τ :

$$dX_t = \tau(\mu - X_t)dt + \sigma dW_t, \qquad (1)$$

where W_t is a standard Brownian motion and σ is the volatility parameter. The data sequence x_1, \ldots, x_t is then a discretization of the stochastic process defined by θ . In this setting, the learning objective is to both identify the parameters of the stochastic process and predict the next sample given the previous ones. We will also consider more intricate processes that are not Markovian.

Given a dataset of tasks $\theta_1, \dots, \theta_N$ and associated samples $x_{1:T}^{(1)}, \dots, x_{1:T}^{(N)}$, a model f is trained by minimizing the next-sample prediction loss

$$\widehat{L}(f, (\theta_n, x_{1:T}^n)_{n \le N}) = \frac{1}{NT} \sum_{n=1}^N \sum_{t=1}^T \ell_t(f(x_{1:t-1}^n), x_t^n), \qquad (2)$$

where ℓ_t is a per-sample loss which depend on t to encompass regression and classification tasks. Note that the model is trained to predict the next sample x_t given the previous samples $x_{1:t-1}$, without any explicit supervision on the task θ . This is why ICL is referred to as an emergent ability of large models (Wei et al., 2022).

We consider two kinds of error for ICL: (i) the ability of the model to identify the correct task given some in-context examples, which we refer to as task selection, and (ii) the generalization error of the trained model \hat{f} obtained by minimizing (2) on a training dataset, which we refer to as generalization error. We first study task selection, before turning to the generalization error, which is more involved.

3.2 Task selection

Our first main result concerns the ability of a trained model to perform ICL and in particular to retrieve the correct task given some input sequence. For this, we adopt the Bayesian point of view, similarly to Lin and Lee (2024); Zekri et al. (2024); Jeon et al. (2024); Zhang et al. (2025b); Wang et al. (2025b). Indeed, if f is arbitrarily powerful and trained to optimality, f learns the Bayesian optimal predictor. If we denote the posterior $\hat{p}_t(\theta | x_{1:t-1})$ the posterior distribution over tasks given the input sequence $x_{1:t-1}$, the Bayesian optimal predictor is given by

$$f(x_{1:t-1}) = \arg\min_{\hat{x}_t} \mathbb{E}_{\theta \sim \hat{p}_t(\cdot | x_{1:t-1})} \left[\mathbb{E}_{x_t \sim p_t(\cdot | x_{1:t-1}, \theta)} \left[\ell_t(\hat{x}_t, x_t) \right] \right]. \tag{3}$$

For a model to perform ICL given in-context examples $x_{1:t-1}$ generated from a task θ^* , it is therefore necessary that the posterior $\widehat{p}_t(\theta \mid x_{1:t-1})$ concentrates around the true task θ^* as the number of in-context examples t increases. Our first main result provides a quantitative guarantee of this concentration and highlights the role of the properties of the pretraining distribution π .

For this, we require some mild assumptions on the data generation process only; they do not restrict the prior π . Since our focus is on the influence of the prior π on task identification, in the main text we mainly focus on assumptions and quantities that involve π , and defer the detailed assumptions to Appendix D. We will therefore use the notation poly(x) to denote a quantity that is polynomial in x with coefficients independent of the prior π and the number of samples T.

Assumption 1 (Data generation, informal). Let $\theta^* \in \Theta$ be the true task. We assume:

- (i) **Tail control.** Sequences $x_{1:t}$ generated under the true task θ^* have controlled tails, at most poly(T) on typical tail events and π admits a second moment.
- (ii) Moment bound. For any $T \ge 1$, $\mathbb{E}_{X \sim p_T(\cdot \mid \theta^*)} \left[\log^2 \left(\sup_{\theta \in \Theta} \frac{p_T(x_{1:T} \mid \theta)}{p_T(x_{1:T} \mid \theta^*)} \right) \right]$ is at most poly(T).
- (iii) Local regularity. The prior density π is continuous and, for any R > 0, $t \le T$,

$$\log \frac{\operatorname{p}_t(x_t|x_{1:t-1},\theta)}{\operatorname{p}_t(x_t|x_{1:t-1},\theta')} \leq \operatorname{poly}(R) \|\theta - \theta'\| \quad \text{ for all } x_{1:t},\theta,\theta' \text{ such that } \|x_s\|,\|\theta\|,\|\theta'\| \leq R$$

These assumptions are quite mild and are satisfied by our examples, see Appendix F.2.

As a metric to assess the quality of a given retrieved task θ w.r.t. the true task θ^* , we consider the Rényi divergence (Rényi, 1961) of order $\rho \in (0,1)$ between the distributions $p_T(\cdot | \theta)$ and $p_T(\cdot | \theta^*)$:

$$\mathrm{D}_{\rho}(\theta \parallel \theta^*) = -\frac{1}{T(1-\rho)} \log \mathbb{E}_{X \sim \mathrm{p}_T(\cdot \mid \theta^*)} \left[\prod_{t=1}^T \left(\frac{\mathrm{p}_t(x_t \mid x_{1:t-1}, \theta)}{\mathrm{p}_t(x_t \mid x_{1:t-1}, \theta^*)} \right)^{\rho} \right].$$

We divide by T to obtain a per-sample divergence that does not trivially diverge as T increases.

Our main theorem below shows that, under Assumption 1, the posterior distribution over tasks concentrates around the true task θ^* as the number of in-context examples T increases, at a rate that depends on the properties of the pretraining distribution π .

Theorem 1 (Task selection). Let $\rho \in (0,1)$, under Assumption 1, with $\pi(\theta^*) > 0$ and $x_{1:T} \sim p_T(\cdot \mid \theta^*)$, the posterior distribution over tasks satisfies

$$\mathbb{E}_{x_{1:T}} \left[\mathbb{E}_{\theta \sim \widehat{p}_T(\cdot | x_{1:T})} \left[\mathcal{D}_{\rho}(\theta \| \theta^*) \right] \right] \leq \frac{1+\rho}{(1-\rho)T} \log 1/\pi(\theta^*) + \mathcal{O}\left(\frac{\log T}{T}\right), \tag{4}$$

where the terms in $\mathcal{O}\left(\frac{\log T}{T}\right)$ do not depend on the prior π or are negligible compared to the first term.

To place this result into context, Theorem 1 provides a guarantee on how close the posterior distribution over tasks is to the true task θ^* as the number of in-context examples T increases. The right-hand side (RHS) decays as $\mathcal{O}(1/T)$, which shows that the posterior concentrates around the true task as the number of examples in-context increases. The speed of convergence is governed by the coefficient $\log 1/\pi(\theta^*)$, which quantifies how well the prior π covers the true task θ^* : the smaller $\pi(\theta^*)$, the slower the convergence. Since in ICL we wish to study the capabilities of learning a new task from in-context examples, this result quantifies the speed at which ICL learns this new task θ^* : the further θ^* is from the bulk of the prior π , the slower ICL learns this new task. Thus, when learning with ICL, the ability to learn a new task and its robustness to new tasks therefore crucially depends on the tail of the prior π . This simple statement thus captures a key aspect of ICL that was observed empirically in several works (Chan et al., 2022; Singh et al., 2023).

From a technical viewpoint, Theorem 1 is proven in Appendix D using ideas from Bayesian statistics (Zhang, 2003, 2006) is extremely general, covers discrete and continuous task spaces, and does not require any probabilistic structure on the data sequence $x_{1:t}$ nor specific data distributions. Moreover, unlike most existing results, Theorem 1 provides a guarantee on the posterior distribution given all T in-context examples, and not only on the regret, which bounds the average error of the posterior distributions given $1, \ldots, T$ examples. This better reflects the practical use of ICL, where the user typically only considers the output of the model after all in-context examples have been provided.

Finally, we provide in the appendix, in Appendix D.4 a more refined version of Theorem 1 that involves not just the prior density at the true task $\pi(\theta^*)$ but also the local geometry of the prior π around θ^* , which can provide much sharper bounds in some cases. This refined result also encompasses the case where $\pi(\theta^*) = 0$, in which the ICL error is not vanishing anymore. In this scenario, it shows that ICL can struggle on out-of-distribution tasks, as empirically studied previously (Goddard et al., 2025; Kwon et al., 2025; Yadlowsky et al., 2023).

Takeaway #1: Heavier-tailed priors are beneficial for task identification and its robustness, as they improve the learning speed on new tasks.

We will now examine the generalization error of ICL and see that there is a trade-off.

3.3 Generalization error

The second key statistical question for ICL is its generalization error. For the trained transformer to accurately behave as the Bayesian optimal predictor w.r.t. the prior π , it is necessary that the next-token prediction be minimized on the true data distribution, and not just on the training data.

We therefore study the generalization error of the trained model \hat{f} obtained by minimizing (2) on a training dataset. We consider a dataset consisting of N tasks $\theta_1, \ldots, \theta_N$ sampled independently from the prior π , and for each task θ_n , a sequence of T samples $x_{1:T}^n$ generated according to the task-specific distribution $p_T(\cdot | \theta_n)$: for $n \leq N$, for t < T, $x_{t+1}^{(n)} \sim p_{t+1}(\cdot | x_{1:t}^{(n)}, \theta_n)$.

To the best of our knowledge, existing concentration for dependent sequences do not cover this case. We thus develop our own framework: we encompass non-independent and identically distributed (i.i.d.) and non-Markovian data sequences through a weak dependence assumption in Wasserstein distance, and we handle heavy-tailed task distributions by taking inspiration from the recent framework of Li and Liu (2024a); Li et al.

(2024). The resulting framework is therefore quite general and can be of independent interest beyond ICL, see Appendix E.

Here we again present a simplified version of our assumptions, where we focus on the few key quantities that are relevant in our study: how dependent the data sequence is and how heavy-tailed the prior π is, quantified through the maximal moment of π that exists*. We refer to Appendix E.3 for the complete version of the assumptions. We consider \mathcal{F} a class of models $f: \cup_t (\mathbb{R}^k)^t \to \mathbb{R}^k$ and $\ell_t: \mathbb{R}^k \times \mathbb{R}^k \to \mathbb{R}_+$ a per-sample loss function that can depend on time t.

Assumption 2 (Generalization, informal).

- (i) Moment condition. There is $q \ge 2$ an integer such that $\mathbb{E}_{\theta \sim \pi}[\|\theta\|^q] < \infty$.
- (ii) Influence of the task. There is $A_T > 0$ such that, any $t \le T$, any $\theta, \theta' \in \Theta$,

$$W_1(p_t(dx_t \mid \theta), p_t(dx_t' \mid \theta')) \le A_T \|\theta - \theta'\|. \tag{5}$$

(iii) Weak dependence. There is $B_T > 0$ such that, for any $s < t \le T$, any $\theta \in \Theta$, any $x_{1:s}, x'_{s}$,

$$W_1(p_t(dx_t \mid x_{1:s}, \theta), p_t(dx_t' \mid x_{1:(s-1)}, x_s', \theta)) \le B_T(1 + \|\theta\|).$$
(6)

(iv) Average Lipschitzness. There is an $L_T > 0$ such that, for any $f \in \mathcal{F}$, any $x_{1:T}, x'_t$,

$$\frac{1}{T} \sum_{s=1}^{T} \| f(x_{1:s-1}) - f(x_{1:t-1}, x'_t, x_{t+1:s-1}) \| \le L_T \| x_t - x'_t \|, \tag{7}$$

(v) **Usual conditions.** The losses ℓ_t are 1-Lipschitz; the class of models \mathcal{F} is bounded and uniformly Lipschitz with respect to some metric and x_t conditioned on $x_{1:t-1}$, θ is uniformly sub-Gaussian.

q, A_T , B_T , and L_T are the key quantities that govern the generalization error of ICL. When π has polynomial tails, q quantifies how heavy-tailed the prior π is: the smaller q, the heavier the tail of π . For Student's t-distribution with ν degrees of freedom, $q = \lfloor \nu - 1 \rfloor$. B_T quantifies how dependent the data sequence is while A_T also quantifies how much the task influences the data distribution: in the case of an i.i.d. sequence, both A_T and B_T are bounded w.r.t. T, which might not be the case in general. L_T quantifies how much the model f uses the older examples in context: for transformer with context length at least T, L_T is typically bounded. If, on the contrary, the context length is kept constant and smaller than T, as in Zekri et al. (2024), L_T can decay as 1/T. In particular, Assumption 2 skips the assumptions on the size of the hypothesis class \mathcal{F} since this is not our main focus, and we refer to the appendix for details.

Our main result provides a bound on the generalization error of the trained model \hat{f} :

$$\widehat{\text{gen}} := \mathbb{E}_{\theta \sim \pi} \left[\mathbb{E}_{x_{1:T} \sim p_T(\cdot \mid \theta)} \left[\frac{1}{T} \sum_{t=1}^T \ell_t(\widehat{f}(x_{1:t-1}), x_t) \right] \right] - \widehat{L}(\widehat{f}, (\theta_n, x_{1:T}^n)_{n \leq N}), \tag{8}$$

for \hat{f} being the model obtained using the empirical distribution $(\theta_n, x_{1:T}^n)_{n \leq N}$.

Theorem 2. Under Assumption 2, for any $\delta \in (0, e^{-2})$, with probability at least $1 - \delta$, it holds:

(a) If
$$\delta \ge Ne^{-q}$$
, then
$$\widehat{\text{gen}} \le \mathcal{O}\left(\frac{(\log 1/\delta)^{3/2} L_T \sqrt{T}}{\sqrt{N}} \left(1 + A_T \sqrt{T} + B_T T\right)\right),\tag{9}$$

(b) If
$$\delta < Ne^{-q}$$
, then
$$\widehat{\text{gen}} \le \mathcal{O}\left(\frac{L_T\sqrt{T}}{\delta^{1/q}\sqrt{N}}\left(1 + A_T\sqrt{T} + B_TT\right)\right),\tag{10}$$

where the terms in $\mathcal{O}(\cdot)$ depend polynomially on q, $\log N$, the scale of π and the size of \mathcal{F} .

Like standard concentration inequalities for sums of independent heavy-tailed random variables, Theorem 2 provides two regimes. For small deviations, i.e., δ not arbitrarily small, the generalization error behaves like in a sub-exponential setting. However, for large deviations, i.e., δ very small, the behaviour of the generalization error worsens and depends on the moment q of the prior π .

^{*}We focus here on prior distributions with polynomially decaying tails, such as the Student-t family, since it is the most representative. A similar result could be established for priors with subexponential tails.

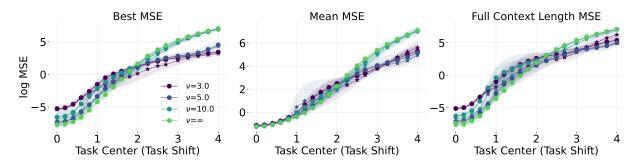


Figure 1: Influence of the degree of freedom parameter of a Student-t pretraining distribution on the ICL error for different task shifts with and without importance weighting. Weighted samples given by $-\star$ marker.

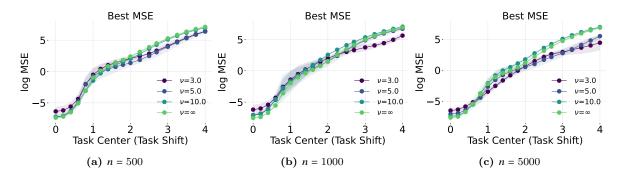


Figure 2: Generalization for linear regression with a Student-t prior of varying ν as a function of n.

The generalization thus depends critically on the moment q of the prior π : the smaller the moment q, the heavier the tail of the prior π and the worse the generalization error. This provides a counterpoint to the task selection result of Theorem 1 that showed that heavier-tailed priors are beneficial for task identification. This highlights a fundamental trade-off in the choice of the pretraining distribution π : heavier-tailed priors are beneficial for task identification, but harm the generalization error.

This bound also highlights how much larger the number of tasks must be compared to the number of incontext examples to ensure good generalization: in general, one needs N to be at least much larger than T to ensure a small generalization error. This is in line with our experiments and previous empirical studies Raventós et al. (2023): to obtain optimal ICL performance with a context length of 16 or 64, one needs thousands of tasks. Moreover, if the data sequence is highly dependent, i.e., A_T and B_T are large, the requirement on the number of tasks N for ICL to generalize well also increases. This will be demonstrated in Section 4.3.

Takeaway #2: Heavier-tailed priors and stronger temporal dependences increase the number of tasks required for reliable ICL generalization.

4 Experiments

We conduct a series of experiments to empirically study the behavior of the pretraining distribution on the performance of ICL[†]. We aim to answer two main questions: do the qualitative characteristics of the proposed bounds hold in practice? and; how do modifications of the pretraining distribution affect performance as the test distribution changes in distance from the pretraining distribution? To do this, we train a transformer under different pretraining distributions to solve different ICL tasks.

ICL evaluation through robustness to distribution shift. The transformer is trained on tasks θ sampled from a pretraining distribution π . To assess the ICL performance, we evaluate the trained model on tasks $\theta' = \theta + \Delta$ where $\theta \sim \pi$ and Δ is a deterministic shift and report the ICL error on these shifted tasks as a function

[†]Additional results and figures are in Appendix B.

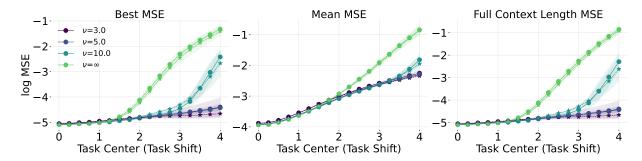


Figure 3: Influence of the degree of freedom parameter of a Student-t pretraining distribution on the ICL error for different task shifts with and without importance weighting for predicting the next step in an OU process with context length of 32. Weighted samples indicated by the $-\star$ marker.

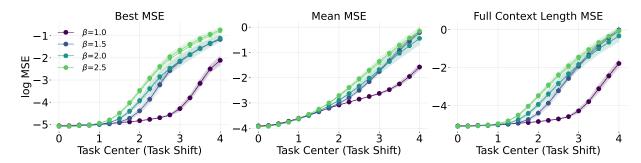


Figure 4: Influence of the shape of a generalized normal pretraining distribution on the ICL error for different task shifts with and without importance weighting for predicting the next step in an OU process.

of the shift magnitude $\|\Delta\|$ Studying this error as a function of the shape of the pretraining distribution allows us to validate the theory in Theorem 1. We also study the performance of ICL as a function of the number of pretraining tasks to test how well the methods generalize, with an emphasis on relating the theory in Theorem 2.

Distributions and Metrics. The pretraining distributions and their parameter values are given in Table 1. The parameters are chosen such that changing them produces a change in the shape of the pretraining distribution. In both cases, lower parameter values indicate heavier tails of the distribution. The scale parameter is chosen such that all pretraining distributions have the same variance. For all experiments, we consider mean squared error (MSE) as the metric we compare. We also consider the best MSE over the context length, which is given by $\min_{t} (\hat{f}(x_t) - x_{t+1})^2$; the mean MSE given by $\frac{1}{T} \sum_{t=1}^{T} (\hat{f}(x_t) - x_{t+1})^2$; and finally the full context length MSE given by $(\hat{f}(x_{T-1}) - x_T)^2$. These allow us to see how the different priors perform while taking into consideration the full context length.

4.1 Linear Regression

We first consider the linear regression setting introduced in Example 3.2 where each $\theta \in \mathbb{R}^d$ defines a linear regression task $y_i = \theta^T q_i + \epsilon_i$ for i = 1, ..., 64 where 64 is the context length. During pretraining, we sample θ according to four different distributions, where the distributions have the same location and scale but different tail decay. We consider Student-t distributions with different shape parameters. In Fig. 1, we see that the performance for small task shifts, the normal distribution prior is the highest performing, but for larger shifts the heavier tailed distributions perform better.

Table 1: Pre-training distribution parameters.

Dist.	Param.
Gen. Normal	$\beta \in \{1, 1.5, 2, 2.5\}$
Student-t	$\nu \in \{3, 5, 10\}$

Reweighting. To further investigate the predictions of Theorem 1, we consider reweighting the pretraining distribution: if we are given samples from a distribution P but know that a pretraining distribution Q exhibits

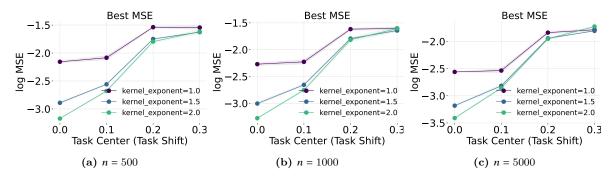


Figure 5: Generalization of a transformer trained to predict the next step of the Volterra as a function of n the number of tasks with context length of 32.

strong performance, can we improve the performance of distribution P by matching Q via importance sampling i.e. $\mathbb{E}_{Q}[\ell(X)] = \mathbb{E}_{P}\left[\ell(Y)\frac{\mathrm{d}Q}{\mathrm{d}P}\right]$? We study this in Fig. 1 where we reweigh samples such that they are approximately uniform over the support of the empirical distribution. The results indicate small improvement in the performance under large shifts using the reweighting as compared to without reweighting.

Generalization. We next consider how the error behaves as the number of pretraining tasks changes. This allows us to see under which regime the tail parameter plays an important role, as Theorem 2 suggests heavier-tailed pretraining distributions require more samples to generalize. We illustrate this in Fig. 2 for $n \in \{500, 1000, 5000\}$. The results validate the theory and indicate that for a smaller number of tasks, the lighter tail prior performs better whereas with more tasks the performance of the heavier tailed prior is improved.

4.2 Linear Stochastic Differential Equations

In the next set of experiments, we follow the setup in Example 3.3 with a stochastic process satisfying (1). For our metric of success, we compare $(\hat{X}_{t+1} - \mathbb{E}[X_{t+1} \mid X_t])^2$ where \hat{X}_{t+1} is conditioned on the context of $X_{s < t}$. We consider θ , μ sampled from different pretraining distributions and again compare the performance of ICL on different test tasks. We study both the Student-t distribution in Fig. 3 and the generalized normal in Fig. 4. In both instances, we see that the heavier tailed pretraining distribution performs better for larger distribution shifts. In the generalized normal case, the effect of reweighting is practically negligible, but in the Student-t case, we see some benefit, particularly in the large shift regime.

4.3 Stochastic Volterra Equations

We finally consider stochastic Volterra equations as a model of nonlinear stochastic processes that have long range dependencies. These processes are, under certain conditions, known to model fractional Brownian motion, which exhibit self-similarity which has been thought to represent the distribution of tokens in LLMs (Alabdulmohsin et al., 2024). Each task θ parametrizes a multi-layer perceptron b_{θ} and induces the process: $X_t = X_0 + \int_0^t (t-s)^{-\alpha} b_{\theta}(X_s) ds + \int_0^t (t-s)^{-\alpha} dW_s$, where W_t is a standard Brownian motion and $\alpha > 0$ controls the temporal dependence of the process: the smaller α is, the more past values influence the current value. The dependency coefficients in Theorem 2 thus depend explicitly on α , they are larger for smaller α , see Appendix F.1. We consider the generalization capabilities as a function of the number of pretraining tasks in Fig. 5 and as a function of α . Theorem 2 predicts that generalization should suffer for smaller α due to the increased dependencies, which is validated in the experiments: the performance gap between the different α is larger for smaller number of tasks.

5 Conclusion

In this work we study ICL through the perspective of task selection and generalization. Our main theoretical contributions describe error bounds of ICL in terms of both task selection and generalization. We show that

a pre-training distribution must be carefully chosen such that the effects of both of these error terms are appropriately balanced. Consequently, the theory allows one to explicitly design a prior distribution based on robustness considerations. We design experiments which consider to what extent ICL can generalize on new tasks that may be out of distribution. The key takeaways are that a heavier tailed prior is appropriate when considering distribution shifts or when many task examples are available. These experiments shed light on how to appropriately pre-train transformers for their use with ICL, with specific emphasis on numerical tasks.

Limitations and Future Directions While our theoretical results are general, the experiments are limited to numerical data: it remains to be seen how this applies to training LLMs when large numbers of documents need to be considered. The reweighting experiments most closely correspond to the possible interventions one may make during pre-training or fine-tuning to improve ICL. A natural follow-up study would consider how to leverage these insights to improve ICL on LLMs with tokens rather than continuous numerical data.

References

- Baptiste Abélès, Eugenio Clerico, and Gergely Neu. Generalization bounds for mixing processes via delayed online-to-pac conversions. In *Algorithmic Learning Theory*, pages 23–40. PMLR, 2025.
- Kwangjun Ahn, Xiang Cheng, Hadi Daneshmand, and Suvrit Sra. Transformers learn to implement preconditioned gradient descent for in-context learning. *Advances in Neural Information Processing Systems*, 36: 45614–45650, 2023a.
- Michael Ahn, Anthony Brohan, Noah Brown, Yevgen Chebotar, Omar Cortes, Byron David, Chelsea Finn, Chuyuan Fu, Keerthana Gopalakrishnan, Karol Hausman, et al. Do as i can, not as i say: Grounding language in robotic affordances. In *Conference on robot learning*, pages 287–318. PMLR, 2023b.
- Ekin Akyürek, Dale Schuurmans, Jacob Andreas, Tengyu Ma, and Denny Zhou. What learning algorithm is incontext learning? investigations with linear models. In *International Conference on Learning Representations* (ICLR), 2023.
- Ibrahim M Alabdulmohsin, Vinh Tran, and Mostafa Dehghani. Fractal patterns may illuminate the success of next-token prediction. Advances in Neural Information Processing Systems, 37:112864–112888, 2024.
- Waïss Azizian, Franck Iutzeler, Jérôme Malick, and Panayotis Mertikopoulos. The global convergence of stochastic gradient descent in non-convex landscapes: Sharp estimates via large deviations. In Forty-second International Conference on Machine Learning, 2025.
- Yu Bai, Fan Chen, Huan Wang, Caiming Xiong, and Song Mei. Transformers as statisticians: Provable incontext learning with in-context algorithm selection. In *Thirty-seventh Conference on Neural Information Processing Systems*, 2023.
- Milad Bakhshizadeh, Arian Maleki, and Victor H De La Pena. Sharp concentration results for heavy-tailed distributions. *Information and Inference: A Journal of the IMA*, 12(3):1655–1685, 2023.
- Raphaël Barboni, Gabriel Peyré, and François-Xavier Vialard. Understanding the training of infinitely deep and wide resnets with conditional optimal transport. Communications on Pure and Applied Mathematics, 2025.
- Andrew Barron, Mark J Schervish, and Larry Wasserman. The consistency of posterior distributions in non-parametric problems. *The Annals of Statistics*, 27(2):536–561, 1999.
- Amanda Bertsch, Maor Ivgi, Emily Xiao, Uri Alon, Jonathan Berant, Matthew R Gormley, and Graham Neubig. In-context learning with long-context models: An in-depth exploration. In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 12119–12149, 2025.
- S Boucheron, O Bousquet, G Lugosi, and P Massart. Moment inequalities for functions of independent random variables. *Annals of Probability*, 33(2), 2005.

- Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. Language models are few-shot learners. Advances in neural information processing systems, 33:1877–1901, 2020.
- Bryan Chan, Xinyi Chen, András György, and Dale Schuurmans. Toward understanding in-context vs. in-weight learning. In *International Conference on Learning Representations (ICLR)*, 2025.
- Stephanie Chan, Adam Santoro, Andrew Lampinen, Jane Wang, Aaditya Singh, Pierre Richemond, James McClelland, and Felix Hill. Data distributional properties drive emergent in-context learning in transformers. Advances in neural information processing systems, 35:18878–18891, 2022.
- J-R Chazottes, Pierre Collet, Christof Külske, and Frank Redig. Concentration inequalities for random fields via coupling. *Probability Theory and Related Fields*, 137(1):201–225, 2007.
- Sjoerd Dirksen. Tail bounds via generic chaining. Electronic Journal of Probability, 20(53):1–29, 2015.
- Takashi Furuya, Maarten V. de Hoop, and Gabriel Peyré. Transformers are universal in-context learners. In *The Thirteenth International Conference on Learning Representations*, 2025.
- Cheng Gao, Yuan Cao, Zihao Li, Yihan He, Mengdi Wang, Han Liu, Jason Klusowski, and Jianqing Fan. Global convergence in training large-scale transformers. *Advances in Neural Information Processing Systems*, 37:29213–29284, 2024.
- Shivam Garg, Dimitris Tsipras, Percy S Liang, and Gregory Valiant. What can transformers learn in-context? a case study of simple function classes. *Advances in Neural Information Processing Systems*, 35:30583–30598, 2022.
- Chase Goddard, Lindsay M. Smith, Vudtiwat Ngampruetikorn, and David J. Schwab. When can in-context learning generalize out of task distribution? In Forty-second International Conference on Machine Learning, 2025.
- Fredrik Hellström, Giuseppe Durisi, Benjamin Guedj, Maxim Raginsky, et al. Generalization bounds: Perspectives from information theory and pac-bayes. Foundations and Trends® in Machine Learning, 18(1):1–223, 2025.
- Hong Jun Jeon, Jason D. Lee, Qi Lei, and Benjamin Van Roy. An information-theoretic analysis of in-context learning. In Forty-first International Conference on Machine Learning, 2024.
- Arych Kontorovich. Concentration in unbounded metric spaces and algorithmic stability. In *International conference on machine learning*, pages 28–36. PMLR, 2014.
- Leonid Aryeh Kontorovich and Kavita Ramanan. Concentration inequalities for dependent random variables via the martingale method. Ann. Probab., 36(1):2126–2158, 2008.
- Anastasis Kratsios and Takashi Furuya. Is in-context universality enough? MLPs are also universal in-context. arXiv preprint arXiv: 2502.03327, 2025.
- Soo Min Kwon, Alec S Xu, Can Yaras, Laura Balzano, and Qing Qu. Out-of-distribution generalization of in-context learning: A low-dimensional subspace perspective. arXiv preprint arXiv:2505.14808, 2025.
- Rafał Latała and Tomasz Tkocz. A note on suprema of canonical processes based on random variables with regular moments. *Electron. J. Probab*, 20(36):1–17, 2015.
- Michel Ledoux and Michel Talagrand. *Probability in Banach Spaces: isoperimetry and processes*. Springer Science & Business Media, 2013.
- Gen Li, Yuchen Jiao, Yu Huang, Yuting Wei, and Yuxin Chen. Transformers meet in-context learning: A universal approximation theory. arXiv preprint arXiv: 2506.05200, 2025a.
- Shaojie Li and Yong Liu. Concentration and moment inequalities for general functions of independent random variables with heavy tails. *Journal of Machine Learning Research*, 25(268):1–33, 2024a.

- Shaojie Li and Yong Liu. Concentration inequalities for general functions of heavy-tailed random variables. In Forty-first International Conference on Machine Learning, 2024b.
- Shaojie Li, Bowei Zhu, and Yong Liu. Algorithmic stability unleashed: generalization bounds with unbounded losses. In Forty-first International Conference on Machine Learning, 2024.
- Tianle Li, Ge Zhang, Quy Duc Do, Xiang Yue, and Wenhu Chen. Long-context llms struggle with long in-context learning. *Transactions on Machine Learning Research*, 2025b.
- Yingcong Li, Muhammed Emrullah Ildiz, Dimitris Papailiopoulos, and Samet Oymak. Transformers as algorithms: Generalization and stability in in-context learning. In Andreas Krause, Emma Brunskill, Kyunghyun Cho, Barbara Engelhardt, Sivan Sabato, and Jonathan Scarlett, editors, *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pages 19565–19594. PMLR, 23–29 Jul 2023.
- Ziqian Lin and Kangwook Lee. Dual operating modes of in-context learning. In Forty-first International Conference on Machine Learning, 2024.
- Weidong Liu, Han Xiao, and Wei Biao Wu. Probability and moment inequalities under dependence. *Statistica Sinica*, pages 1257–1272, 2013.
- Sanae Lotfi, Yilun Kuang, Marc Anton Finzi, Brandon Amos, Micah Goldblum, and Andrew Gordon Wilson. Unlocking tokens as data points for generalization bounds on larger language models. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*, 2024.
- Andreas Maurer. Generalization for slowly mixing processes. arXiv preprint arXiv:2305.00977, 2023.
- Amir Moeini, Jiuqi Wang, Jacob Beck, Ethan Blaser, Shimon Whiteson, Rohan Chandra, and Shangtong Zhang. A survey of in-context reinforcement learning. arXiv preprint arXiv:2502.07978, 2025.
- Mehryar Mohri and Afshin Rostamizadeh. Rademacher complexity bounds for non-iid processes. Advances in neural information processing systems, 21, 2008.
- Mehryar Mohri and Afshin Rostamizadeh. Stability bounds for stationary φ -mixing and β -mixing processes. Journal of Machine Learning Research, 11(2), 2010.
- Daniel Paulin. Concentration inequalities for markov chains by marton couplings and spectral methods. *Electron. J. Probab*, 20(79):1–32, 2015.
- Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, et al. Learning transferable visual models from natural language supervision. In *International conference on machine learning*, pages 8748–8763. PmLR, 2021.
- Allan Raventós, Mansheej Paul, Feng Chen, and Surya Ganguli. Pretraining task diversity and the emergence of non-bayesian in-context learning for regression. *Advances in neural information processing systems*, 36: 14228–14246, 2023.
- Alfréd Rényi. On measures of entropy and information. In *Proceedings of the fourth Berkeley symposium on mathematical statistics and probability, volume 1: contributions to the theory of statistics,* volume 4, pages 547–562. University of California Press, 1961.
- Borja Rodríguez-Gálvez, Ragnar Thobaben, and Mikael Skoglund. An information-theoretic approach to generalization theory. arXiv preprint arXiv:2408.13275, 2024.
- Michaël E Sander, Raja Giryes, Taiji Suzuki, Mathieu Blondel, and Gabriel Peyré. How do transformers perform in-context autoregressive learning? In *Proceedings of the 41st International Conference on Machine Learning*, pages 43235–43254, 2024.
- Michael Eli Sander and Gabriel Peyré. Towards understanding the universality of transformers for next-token prediction. In *The Thirteenth International Conference on Learning Representations*, 2025.

- Aaditya Singh, Stephanie Chan, Ted Moskovitz, Erin Grant, Andrew Saxe, and Felix Hill. The transient nature of emergent in-context learning in transformers. *Advances in Neural Information Processing Systems*, 36: 27801–27819, 2023.
- Karan Singhal, Shekoofeh Azizi, Tao Tu, S Sara Mahdavi, Jason Wei, Hyung Won Chung, Nathan Scales, Ajay Tanwani, Heather Cole-Lewis, Stephen Pfohl, et al. Large language models encode clinical knowledge. *Nature*, 620(7972):172–180, 2023.
- Jonathan M Stokes, Kevin Yang, Kyle Swanson, Wengong Jin, Andres Cubillos-Ruiz, Nina M Donghia, Craig R MacNair, Shawn French, Lindsey A Carfrae, Zohar Bloom-Ackermann, et al. A deep learning approach to antibiotic discovery. *Cell*, 180(4):688–702, 2020.
- Michel Talagrand. Upper and lower bounds for stochastic processes: decomposition theorems, volume 60. Springer Nature, 2022.
- Aditya Varre, Gizem Yüce, and Nicolas Flammarion. Learning in-context n-grams with transformers: Sub-n-grams are near-stationary points. In *International Conference on Machine Learning*, 2025.
- R. Vershynin. High-Dimensional Probability: An Introduction with Applications in Data Science. Cambridge Series in Statistical and Probabilistic Mathematics. Cambridge University Press, 2018. ISBN 9781108244541.
- C. Villani. Optimal Transport: Old and New. Grundlehren der mathematischen Wissenschaften. Springer Berlin Heidelberg, 2008.
- Johannes Von Oswald, Eyvind Niklasson, Ettore Randazzo, Joao Sacramento, Alexander Mordvintsev, Andrey Zhmoginov, and Max Vladymyrov. Transformers learn in-context by gradient descent. In Andreas Krause, Emma Brunskill, Kyunghyun Cho, Barbara Engelhardt, Sivan Sabato, and Jonathan Scarlett, editors, Proceedings of the 40th International Conference on Machine Learning, volume 202 of Proceedings of Machine Learning Research, pages 35151–35174. PMLR, 23–29 Jul 2023.
- Martin J. Wainwright. *High-Dimensional Statistics: A Non-Asymptotic Viewpoint*. Cambridge Series in Statistical and Probabilistic Mathematics. Cambridge University Press, 2019.
- Aaron T Wang, William Convertino, Xiang Cheng, Ricardo Henao, and Lawrence Carin. On understanding attention-based in-context learning for categorical data. In *International Conference on Machine Learning (ICML)*, 2025a.
- Mingze Wang and E. Weinan. Understanding the expressive power and mechanisms of transformer for sequence modeling. *Neural Information Processing Systems*, 2024.
- Qixun Wang, Yifei Wang, Xianghua Ying, and Yisen Wang. Can in-context learning really generalize to out-of-distribution tasks? In *The Thirteenth International Conference on Learning Representations*, 2025b.
- Jason Wei, Yi Tay, Rishi Bommasani, Colin Raffel, Barret Zoph, Sebastian Borgeaud, Dani Yogatama, Maarten Bosma, Denny Zhou, Donald Metzler, Ed H. Chi, Tatsunori Hashimoto, Oriol Vinyals, Percy Liang, Jeff Dean, and William Fedus. Emergent abilities of large language models. *Transactions on Machine Learning Research*, 2022. ISSN 2835-8856. Survey Certification.
- Roderick Wong. Asymptotic approximations of integrals. SIAM, 2001.
- Jingfeng Wu, Difan Zou, Zixiang Chen, Vladimir Braverman, Quanquan Gu, and Peter Bartlett. How many pretraining tasks are needed for in-context learning of linear regression? In *The Twelfth International Conference on Learning Representations*, 2024.
- Wei Biao Wu. Nonlinear system theory: Another look at dependence. *Proceedings of the National Academy of Sciences*, 102(40):14150–14154, 2005.
- Wei Biao Wu. Asymptotic theory for stationary processes. Stat. Interface, 4(2):207–226, 2011.
- Weimin Wu, Maojiang Su, Jerry Yao-Chieh Hu, Zhao Song, and Han Liu. In-context deep learning via transformer models. In *International Conference on Machine Learning (ICML)*, 2025.

- Sang Michael Xie, Aditi Raghunathan, Percy Liang, and Tengyu Ma. An explanation of in-context learning as implicit bayesian inference. *International Conference on Learning Representations*, 2021.
- Steve Yadlowsky, Lyric Doshi, and Nilesh Tripuraneni. Pretraining data mixtures enable narrow model selection capabilities in transformer models. arXiv preprint arXiv: 2311.00871, 2023.
- Bin Yu. Rates of convergence for empirical processes of stationary mixing sequences. *The Annals of Probability*, 22(1):94–116, 1994. ISSN 00911798, 2168894X. URL http://www.jstor.org/stable/2244496.
- Oussama Zekri, Ambroise Odonnat, Abdelhakim Benechehab, Linus Bleistein, Nicolas Boulle, and Ievgen Redko. Large language models as markov chains. arXiv preprint arXiv:2410.02724, 2024.
- Ruiqi Zhang, Spencer Frei, and Peter L Bartlett. Trained transformers learn linear models in-context. *Journal of Machine Learning Research*, 25(49):1–55, 2024.
- Tong Zhang. Learning bounds for a generalized family of bayesian posterior distributions. Advances in Neural Information Processing Systems, 16, 2003.
- Tong Zhang. From ε -entropy to kl-entropy: Analysis of minimum information complexity density estimation. The Annals of Statistics, pages 2180–2210, 2006.
- Yedi Zhang, Aaditya K Singh, Peter E. Latham, and Andrew M Saxe. Training dynamics of in-context learning in linear attention. In Forty-second International Conference on Machine Learning, 2025a.
- Yufeng Zhang, Fengzhuo Zhang, Zhuoran Yang, and Zhaoran Wang. What and how does in-context learning learn? bayesian model averaging, parameterization, and generalization. In The 28th International Conference on Artificial Intelligence and Statistics, 2025b.
- Kaijian Zou, Muhammad Khalifa, and Lu Wang. On many-shot in-context learning for long-context evaluation. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. Association for Computational Linguistics, 2025.

A	Additional Related Work	16
В	Additional Experimental Results B.1 Linear Regression	18
C	Experimental Details C.1 Data Generation	
D	Task Selection D.1 Preliminary Lemmas	25 27 28
E	Generalization bounds E.1 Moment bounds for general functions	39 40 44
F	Additional details on examples F.1 Example: Volterra equation model	

A Additional Related Work

Training dynamics of ICL Varre et al. (2025) shows that *n*-grams are approximate stationary points in the training of two-layers transformers. Zhang et al. (2025a) studies the training dynamics of a one-layer linear transformer with linear attention on linear regression tasks. Sander et al. (2024) characterize the training dynamics of a one-linear layer transformer on auto-regressive tasks, showing how ICL emerges. Ahn et al. (2023a) show that for linear regression problems and a linear transformer, the global minimizer of the training loss corresponds to performing one step of preconditioned gradient descent. In contrast, our approach focuses on the influence of the pre-training distribution on ICL. We therefore assume that the model is sufficiently expressive and trained optimally enough to approximate the Bayes optimal predictor. We refer to recent works on optimization dynamics of transformers Gao et al. (2024); Barboni et al. (2025); Azizian et al. (2025) and on the approximation capabilities of transformers.

Approximation capabilities of transformers The foundational works of Von Oswald et al. (2023); Akyürek et al. (2023) demonstrate that transformers can implement gradient descent. This has led to a fruitful line of work studying the algorithmic capabilities of transformers. Bai et al. (2023) show that transformers can implement a wide variety of statistical methods. Wang et al. (2025a) shows how transformers can implement functional gradient descent on categorical data, generalizing previous works. Wu et al. (2025) shows how attention transformers can implement gradient descent on a ReLU network. Sander and Peyré (2025) explicitly constructs a transformer that implements kernel causal regression. On a more abstract perspective, Furuya et al. (2025); Kratsios and Furuya (2025) show that (causal) transformers can approximate any (causal) map between measures. Wang and Weinan (2024) studies quantitatively the approximation properties of transformers on "sparse memory" target functions. Li et al. (2025a) obtains explicit approximation bounds for numerical ICL tasks.

B Additional Experimental Results

B.1 Linear Regression

We provide comprehensive experimental results for linear regression tasks (detailed in Section 4.1) using Student-t and generalized normal pretraining distributions. This section presents the ICL error as a function of context length (ICL step) for Student-t priors with degrees of freedom $v \in \{3, 5, 10, \infty\}$ and generalized normal priors with shape parameters $\beta \in \{1, 1.5, 2, 2.5\}$, corresponding to the experimental settings in Fig. 1.

The results in Fig. 6 clearly demonstrate the fundamental trade-off in selecting pretraining distributions for ICL: heavy-tailed priors (small ν) achieve superior performance under distribution shift, while light-tailed priors (large ν) excel on in-distribution tasks. In contrast, Fig. 7 shows that varying the shape parameter of generalized normal priors produces more subtle effects on ICL performance in the linear regression setting.

We also notice on Figs. 6 and 7 that longer context lengths are mostly beneficial for in-distribution tasks: as the perturbation magnitude increases, the performance gains from longer contexts diminish. This is in line with Section 3.2: the performance gain per new example is determined by the prior probability of the task, which decreases with larger perturbations.

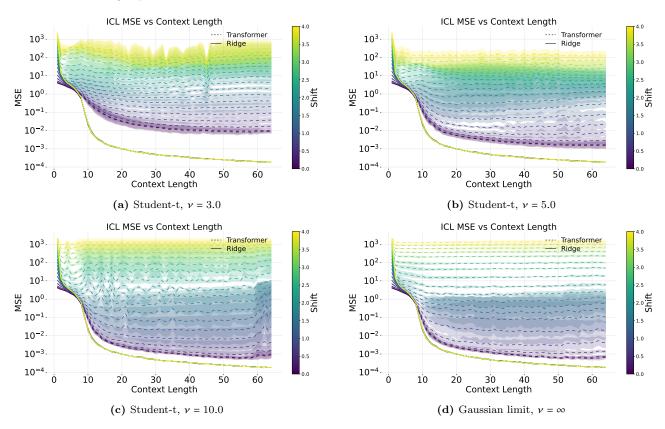


Figure 6: Linear regression with Student-t pretraining distributions: MSE as a function of ICL step for different task shift magnitudes. Heavy-tailed priors ($\nu = 3$) show superior robustness to distribution shift, while light-tailed priors ($\nu = \infty$, Gaussian) perform better on unperturbed tasks. The Ridge regression baseline provides a reference that remains constant across perturbation magnitudes.

We present an extended analysis of the generalization results from Fig. 2 in Fig. 8, examining how the number of pretraining tasks n affects performance across different Student-t tail parameters ν . These results validate Theorem 2, showing that heavy-tailed priors require more training tasks to achieve comparable performance to light-tailed priors.

Finally, we provide an ablation study on the effect of the variance. All other experiments are designed so that the pretraining distribution has unit variance in each dimension. In Fig. 9, we vary the variance of a standard Gaussian pretraining distribution and observe it only changes the ICL performance for in-distribution tasks.

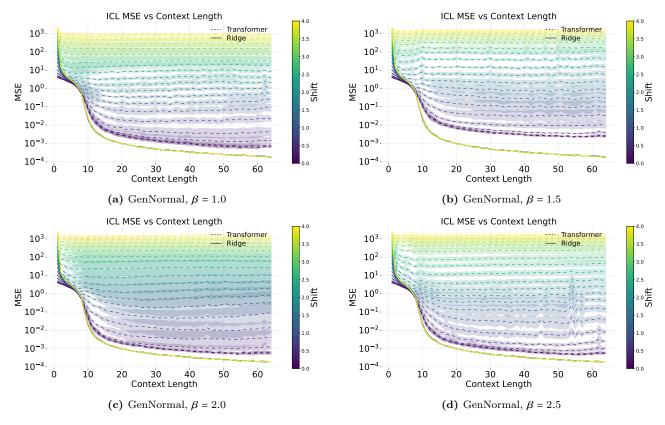


Figure 7: Linear regression with generalized normal pretraining distributions: MSE as a function of ICL step for different task shift magnitudes. The shape parameter β has a more modest impact on performance compared to Student-t distributions, with all variants showing similar convergence patterns across perturbation levels.

B.2 Ornstein-Uhlenbeck Processes

We present detailed experimental results for Ornstein–Uhlenbeck (OU) stochastic processes (described in Section 4.2) using both Student-t and generalized normal pretraining distributions. The figures show ICL error as a function of context length for Student-t priors with degrees of freedom $v \in \{3, 5, 10, \infty\}$ (matching Fig. 3) and generalized normal priors with shape parameters $\beta \in \{1, 1.5, 2, 2.5\}$ (matching Fig. 4) in Figs. 10 and 11, respectively.

Notably, OU processes exhibit different behavior compared to linear regression: the trade-off between indistribution and out-of-distribution performance is less pronounced. As shown in both Figs. 10 and 11, heavytailed priors maintain competitive in-distribution performance while still providing improved robustness to distribution shift.

B.3 Volterra Processes

We present comprehensive results for stochastic Volterra equations (detailed in Section 4.3), which model non-linear processes with long-range dependencies and connections to fractional Brownian motion. Figure 12 shows ICL error as a function of context length for different kernel exponents $\alpha \in \{1, 1.5, 2\}$, where smaller α values correspond to stronger temporal dependencies.

The results confirm our theoretical predictions from Section 3: as the kernel exponent α increases (weaker dependencies), both convergence speed and final performance improve significantly. This validates the dependency structure analysis in Theorem 2.

Figure 13 extends the generalization analysis from Fig. 5, demonstrating how the number of pretraining tasks n interacts with the temporal dependency parameter α . The results show that processes with stronger dependencies (smaller α) require substantially more training data to achieve comparable performance.

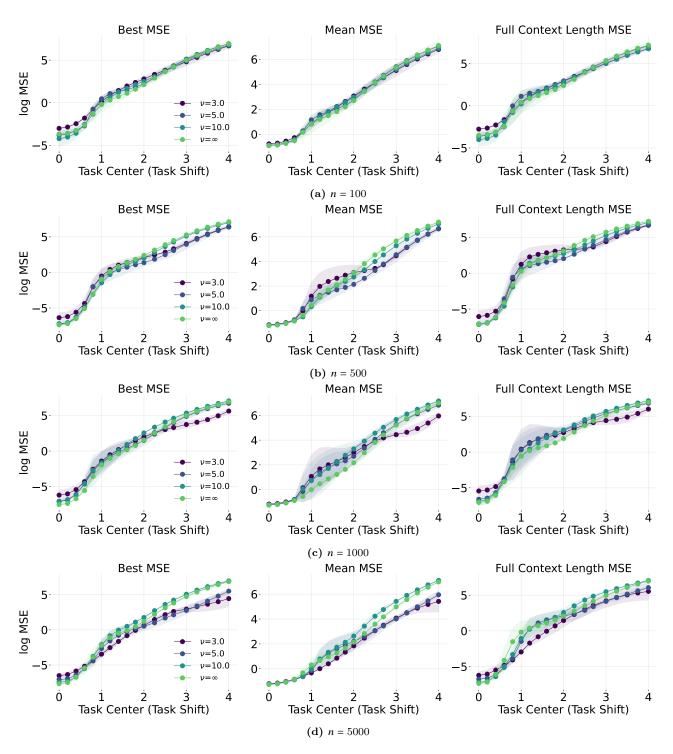


Figure 8: Generalization analysis for linear regression across different numbers of pretraining tasks n for a context length of 64. As predicted by Theorem 2, heavy-tailed priors (small ν) require more tasks to achieve performance comparable to light-tailed priors, but eventually outperform them under distribution shift. The crossover point shifts to larger n for heavier-tailed distributions.

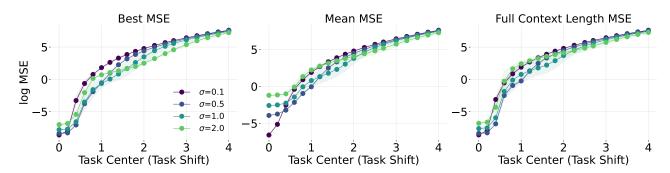


Figure 9: Ablation on the effect of variance for Gaussian pretraining distributions in linear regression. Only indistribution performance is affected by the variance, with larger variances leading to worse performance.

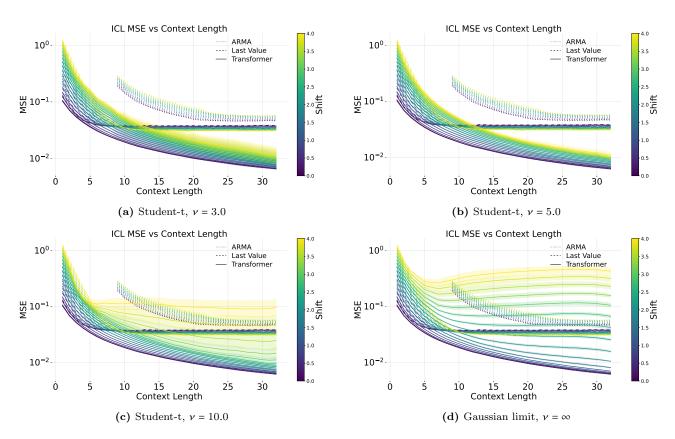


Figure 10: Ornstein—Uhlenbeck processes with Student-*t* pretraining distributions: MSE as a function of ICL step for different task shift magnitudes. Unlike linear regression, heavy-tailed priors maintain strong in-distribution performance while providing superior robustness to perturbations. Baselines include predicting the last observed value and fitting an ARMA(5) model to the context.

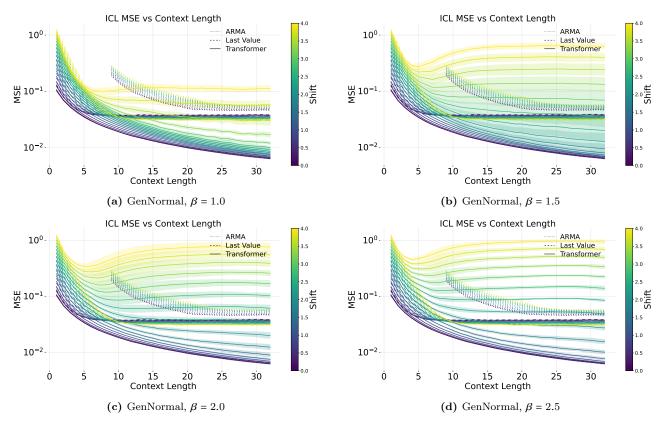


Figure 11: Ornstein–Uhlenbeck processes with generalized normal pretraining distributions (importance weighted): MSE as a function of ICL step for different task shift magnitudes. The shape parameter β shows consistent effects across perturbation levels, with all variants significantly outperforming simple baselines. Importance weighting provides modest improvements in robustness.

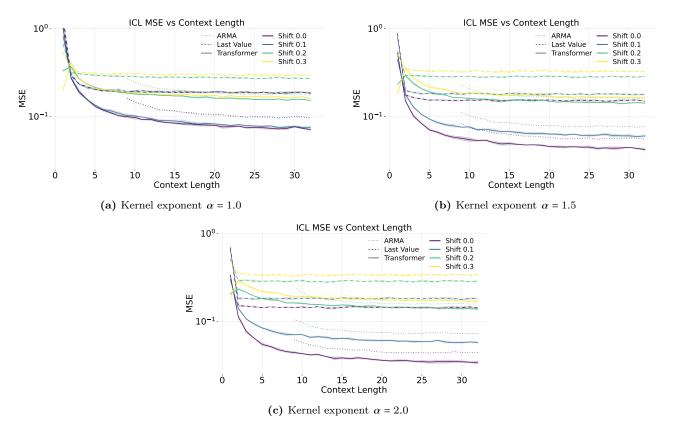


Figure 12: Stochastic Volterra equations: MSE as a function of ICL step across different kernel exponents α . Smaller α values correspond to stronger long-range dependencies, leading to slower convergence and higher final error. The performance gap between different α values demonstrates the impact of temporal dependency structure on ICL learning. Simple baselines provide reference points for comparison.

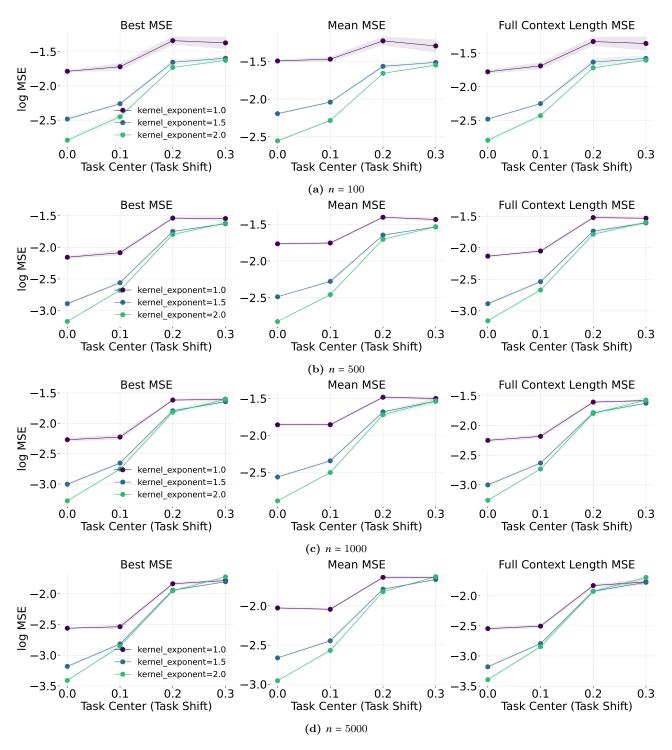


Figure 13: Generalization analysis for Volterra processes across different numbers of pretraining tasks n. Processes with stronger temporal dependencies (smaller α) exhibit larger performance gaps at low n, consistent with Theorem 2. The dependency coefficients in our theory scale with α , explaining why more training tasks are needed to achieve good performance for smaller α values.

C Experimental Details

We roughly follow the experimental setup used by Raventós et al. (2023). Our code is largely based on their implementation given in[‡].

C.1 Data Generation

In all experiments, task parameters $\theta \in \mathbb{R}^d$ are sampled from the distribution mentioned in the main text, data sequences are sampled according to the task. All task distributions during training are zero mean and unit variance in each dimension, except for the Volterra experiments where they are normalized to have standard deviation 0.2. For testing, we sample θ from $\mathcal{N}(\mu \, \mathbb{I}, I)$ where $\mu \in \mathbb{R}$ is the shift value and \mathbb{I} is the all ones vector, and the data is sampled according to this task. Unless otherwise specified, a new set of tasks θ is sampled for each training iteration. Otherwise, when the number of tasks is specified, we sample that many tasks at the start of training and use those same tasks throughout training.

Linear Regression Given a task parameter $\theta \in \mathbb{R}^8$, we sample $x_i \sim \mathcal{N}(0, I_8)$ and $y_i = \langle x_i, \theta \rangle + \epsilon_i$ where $\epsilon_i \sim \mathcal{N}(0, 0.5^2)$. Given a context of $(x_1, y_1), \dots, (x_k, y_k)$, the model is trained to predict y_{k+1} given x_{k+1} with the MSE loss. At evaluation, we evaluate the model output against $x_i^{\mathsf{T}}\theta$. We refer to the linear regression experiments in Raventós et al. (2023) for details.

Ornstein-Uhlenbeck (OU) Process The OU process is given by $dX_t = \tau(\mu - X_t)dt + \sigma dW_t$ and has two parameters: θ and μ . We study a 8-dimensional process where $X_t \in \mathbb{R}^8$ and $\sigma = 0.5I_8$. We consider the initial distribution of $x_0 \sim \mathcal{N}(0, I_8)$. Full paths of X_t are sampled using the Euler-Maruyama method with a step size of $\Delta t = 0.8$. For the sampling of tasks, $\theta \in \mathbb{R}^9$ is sampled from the described distribution, μ is then set to be the first 8 components of θ and τ is set to $0.3 + 0.2 \times \sigma(-0.4\theta_9)$ where σ is the sigmoid function. The model is trained to predict $X_{(k+1)\Delta t}$ given $X_0, X_{\Delta t}, \ldots, X_{k\Delta t}$ with the MSE loss with a maximum context length of 32. For evaluation, we evaluate the model output against $\mathbb{E}[X_{(k+1)\Delta t}|X_0, X_{\Delta t}, \ldots, X_{k\Delta t}]$ which is computable in closed form.

Volterra Process We study a Volterra process in dimension 8 given by

$$X_{t} = X_{0} + \int_{0}^{t} (t - s)^{-\alpha} b_{\theta}(X_{s}) ds + \int_{0}^{t} (t - s)^{-\alpha} \sigma dW_{s},$$
 (C.1)

where the parameter α is chosen according to discrete values in $\{1, 1.5, 2\}$ and $\sigma = 0.6I_8$. X_0 is sampled from $\mathcal{N}(0, I_8)$ again. b_{θ} a clipped two-layer neural network and hidden dimension 16: formally, with $\theta = (W_1, b_1, W_2, b_2)$ then $b_{\theta}(x) = \text{clip}(10(W_2 \tanh(W_1x + b_1) + b_2), -2, 2) - 0.1x$.

We subsample the paths $(X_t)_t$ with step size $\Delta t = 2$ to obtain discrete samples $(X_0, X_{\Delta t}, X_{2\Delta t}, \ldots)$ and each $X_{k\Delta t}$ is computed from past samples using 10 steps of the Euler-Maruyama method with step size $\Delta t/10$. The model is trained to predict $X_{(k+1)\Delta t}$ given $X_0, X_{\Delta t}, \ldots, X_{k\Delta t}$ with the MSE loss with a maximum context length of 32. For evaluation, we evaluate the model output against $\mathbb{E}[X_{(k+1)\Delta t}|X_0, X_{\Delta t}, \ldots, X_{k\Delta t}]$ which is computable in closed form.

C.2 Architecture and Optimization Details

For all experiments, we consider the architecture inspired by GPT-2 as used in Raventós et al. (2023). For linear regression experiments, we use a context length of 64 points, 6 layers, embedding dimension of 32, 8 attention heads and an output dimension of 1. For the other experiments, we use a context length of 32 points, 8 layers, embedding dimension of 128, 2 attention heads and an output dimension of 8.

All models were trained for 5×10^5 iterations. Experiments are run with AdamW optimizer with a weight decay of 0.1 with a cosine learning rate schedule and 50,000 warmup steps. All experiments were run on NVIDIA H100 GPUs. We performed a hyperparameter sweep over learning rate where we considered two learning rates and chose the best model. Experiments are repeated 3 different times with different seeds. LLMs were used to assist in code writing.

[‡]https://github.com/mansheej/icl-task-diversity

D Task Selection

In this section, we study how tasks are selected at test time in ICL. This section is structured as follows. First we consider an abstract setting for Appendices D.1 and D.2 where in Appendix D.1 we state a few preliminary lemmas that will be useful in the analysis, and in Appendix D.2 we prove a template task selection bound under minimal assumptions. Then, in Appendix D.3, we reintroduce the ICL setting along with the detailed assumptions before proving the main task selection bound in Appendix D.4, which is where the main contribution of this section lies.

D.1 Preliminary Lemmas

Definition 1 (Kullback-Leibler divergence). For \mathbb{P} and \mathbb{Q} two probability measures on a measurable space \mathcal{X} , the Kullback-Leibler (KL) divergence from \mathbb{P} to \mathbb{Q} is defined as

$$KL(\mathbb{P} \parallel \mathbb{Q}) = \begin{cases} \int_{\mathcal{X}} \log\left(\frac{d\mathbb{P}}{d\mathbb{Q}}(x)\right) d\mathbb{P}(x) & \text{if } \mathbb{P} \ll \mathbb{Q} \\ +\infty & \text{otherwise.} \end{cases}$$
(D.1)

We now state the Donsker-Varadhan lemma, also known as the Gibbs variational principle.

Lemma D.1 (Donsker-Varadhan lemma, Gibbs variational principle). Consider \mathbb{P} probability measure on a measurable \mathcal{X} and $g \colon \mathcal{X} \to \mathbb{R}$ a measurable function such that $\mathbb{E}_{\mathbb{P}}[\exp(g)] < \infty$. Then, we have

$$\log \mathbb{E}_{\mathbb{P}}[e^{g(x)}] = \sup_{\mathbb{Q}} \{ \mathbb{E}_{\mathbb{Q}}[g(x)] - \mathrm{KL}(\mathbb{Q} \parallel \mathbb{P}) \}, \tag{D.2}$$

with equality attained in particular for $\frac{d\,\mathbb{Q}}{d\,\mathbb{P}}(x)\propto e^{g(x)}.$

See for instance Hellström et al. (2025); Rodríguez-Gálvez et al. (2024) for original references and proofs. Let us state a technical consequence of this lemma that essentially corresponds to Zhang (2003, Lem. 3.1).

Lemma D.2. Consider X a random variable on \mathcal{X} distributed according to \mathbb{P}_X and θ a random variable on Θ with prior distribution $\pi(d\theta)$ and with posterior distribution such that, conditionally on X,

$$\widehat{\mathbb{P}}(d\theta \mid X) = \frac{d \, \mathbb{P}(X \mid \theta)}{d \, \mathbb{P}(X)} \pi(d\theta) \,. \tag{D.3}$$

Consider $L: \mathcal{X} \times \Theta \to \mathbb{R}$ a measurable function. Then,

$$\mathbb{E}_{X,\theta \sim \widehat{\mathbb{P}}(\cdot|X)} [L(X,\theta) - \log \mathbb{E}_X [\exp(L(X,\theta))]] \le \mathbb{E}_X [\mathrm{KL}(\mathbb{P}_{\theta}(\cdot \mid X) \parallel \pi)]. \tag{D.4}$$

Proof. We apply Lemma D.1 with $g(\theta) = L(X, \theta) - \log \mathbb{E}_X[\exp(L(X, \theta))]$ conditionally on X to obtain

$$\mathbb{E}_{\theta \sim \widehat{\mathbb{P}}(\cdot|X)}[L(X,\theta) - \log \mathbb{E}_X[\exp(L(X,\theta))] - \text{KL}(\mathbb{P}_{\theta}(\cdot|X) \| \pi)]$$
(D.5)

$$\leq \log \mathbb{E}_{\theta \sim \pi} \left[\exp(L(X, \theta) - \log \mathbb{E}_X \left[\exp(L(X, \theta)) \right]) \right]. \tag{D.6}$$

We then have

$$\mathbb{E}_{X} \Big[\exp \mathbb{E}_{\theta \sim \widehat{\mathbb{P}}(\cdot|X)} [L(X,\theta) - \log \mathbb{E}_{X} [\exp(L(X,\theta))] - \text{KL}(\mathbb{P}_{\theta}(\cdot|X) \| \pi)] \Big]$$
(D.7)

$$\leq \mathbb{E}_{X \mid \theta \sim \pi} \left[\exp(L(X, \theta) - \log \mathbb{E}_{X} \left[\exp(L(X, \theta)) \right]) \right] = 1, \tag{D.8}$$

and the result follows by Jensen's inequality with the convex function exp.

D.2 Template Task Selection Bound

Let us start with a template task selection bound under minimal assumptions. This proof is adapted from Zhang (2003, Thm. 4.1) to the case of non-i.i.d. data and when the true task is not necessarily in the support of the prior.

Proposition D.1 (Template task selection bound). Consider X a random variable on X distributed according to \mathbb{P}_X and θ a random variable on Θ with prior distribution $\pi(d\theta)$ such that, conditionally on X, θ is distributed according to

$$\widehat{\mathbb{P}}(d\theta \mid X) = \frac{d \, \mathbb{P}(X \mid \theta)}{d \, \mathbb{P}(X)} \pi(d\theta) \,. \tag{D.9}$$

Then, we have, for any $\theta_0 \in \Theta$, for any $\rho \in (0,1)$, $\alpha > 1$,

$$\mathbb{E}_{X,\theta \sim \widehat{\mathbb{P}}(\cdot|X)} \left[-\log \mathbb{E}_X \left[\left(\frac{d \, \mathbb{P}_X(\cdot \mid \theta)}{d \, \mathbb{P}_X(\cdot)} \right)^{\rho} \right] \right] \tag{D.10}$$

$$\leq -\alpha \log \mathbb{E}_{\theta \sim \pi} \left[\exp \left(- \mathbb{E}_X \log \frac{d \, \mathbb{P}_X(\cdot \mid \theta_0)}{d \, \mathbb{P}_X(\cdot \mid \theta)} \right) \right] + \alpha \text{KL}(\mathbb{P}_X(\cdot) \parallel \mathbb{P}_X(\cdot \mid \theta_0)) \tag{D.11}$$

$$+ (\alpha - 1) \mathbb{E}_{X} \left[\log \mathbb{E}_{\theta \sim \pi} \left[\exp \left(-\frac{\alpha - \rho}{\alpha - 1} \log \frac{d \mathbb{P}_{X}(\cdot \mid \theta_{0})}{d \mathbb{P}_{X}(\cdot \mid \theta)} \right) \right] \right]$$
 (D.12)

Proof. To simplify notations in this proof, unless otherwise specified, θ indicates a random variable distributed according to $\widehat{\mathbb{P}}(\cdot \mid X)$. We start from Lemma D.2 with $L(X, \theta) = \rho \log \frac{d \mathbb{P}_X(\cdot \mid \theta)}{d \mathbb{P}_X(\cdot)}$ and rearrange to obtain:

$$\mathbb{E}_{\theta} \left[-\log \mathbb{E}_{X} \left[\left(\frac{d \, \mathbb{P}_{X}(\cdot \mid \theta)}{d \, \mathbb{P}_{X}(\cdot)} \right)^{\rho} \right] \right] \leq \mathbb{E}_{X,\theta} \left[\rho \log \frac{d \, \mathbb{P}_{X}(\cdot)}{d \, \mathbb{P}_{X}(\cdot \mid \theta)} \right] + \mathbb{E}_{X} \left[\mathrm{KL}(\mathbb{P}_{\theta}(\cdot \mid X) \parallel \pi) \right]. \tag{D.13}$$

The left-hand side (LHS) is the quantity we want to bound. We now only need to bound the RHS. Making $\theta_0 \in \Theta$ appear in the bound, we have

$$\mathbb{E}_{X,\theta} \left[\rho \log \frac{d \, \mathbb{P}_X(\cdot)}{d \, \mathbb{P}_X(\cdot \mid \theta)} \right] + \mathbb{E}_X \left[\mathrm{KL}(\mathbb{P}_{\theta}(\cdot \mid X) \parallel \pi) \right] \tag{D.14}$$

$$= \rho \mathbb{E}_{X} \left[\log \frac{d \mathbb{P}_{X}(\cdot)}{d \mathbb{P}_{X}(\cdot \mid \theta_{0})} \right] + \mathbb{E}_{X,\theta} \left[\rho \log \frac{d \mathbb{P}_{X}(\cdot \mid \theta_{0})}{d \mathbb{P}_{X}(\cdot \mid \theta)} \right] + \mathbb{E}_{X} \left[KL(\mathbb{P}_{\theta}(\cdot \mid X) \parallel \pi) \right]$$
(D.15)

$$= \rho \text{KL}(\mathbb{P}_X(\cdot) \parallel \mathbb{P}_X(\cdot \mid \theta_0)) \tag{D.16}$$

$$+ \mathbb{E}_{X,\theta} \left[\rho \log \frac{d \, \mathbb{P}_{X}(\cdot)}{d \, \mathbb{P}_{X}(\cdot \mid \theta)} \right] + \mathbb{E}_{X} [\mathrm{KL}(\mathbb{P}_{\theta}(\cdot \mid X) \parallel \pi)] \,. \tag{D.17}$$

Introducing $\alpha > 1$ and defining $\mu = \frac{\alpha - 1}{\alpha - \rho} < 1$, we now bound the last two terms in (D.17) as follows:

$$\mathbb{E}_{X,\theta} \left[\rho \log \frac{d \, \mathbb{P}_X(\cdot \mid \theta_0)}{d \, \mathbb{P}_X(\cdot \mid \theta)} \right] + \mathbb{E}_X [\mathrm{KL}(\mathbb{P}_{\theta}(\cdot \mid X) \parallel \pi)] \tag{D.18}$$

$$= \alpha \left(\mathbb{E}_{X,\theta} \left[\log \frac{d \, \mathbb{P}_{X}(\cdot \mid \theta_{0})}{d \, \mathbb{P}_{X}(\cdot \mid \theta)} \right] + \mathbb{E}_{X} \left[\mathrm{KL}(\mathbb{P}_{\theta}(\cdot \mid X) \parallel \pi) \right] \right) \tag{D.19}$$

$$-(\alpha - \rho) \left(\mathbb{E}_{X,\theta} \left[\log \frac{d \, \mathbb{P}_X(\cdot \mid \theta_0)}{d \, \mathbb{P}_X(\cdot \mid \theta)} \right] + \mu \, \mathbb{E}_X [\mathrm{KL}(\mathbb{P}_{\theta}(\cdot \mid X) \parallel \pi)] \right). \tag{D.20}$$

Let us first focus on the first term. By the equality case in Lemma D.1 and the definition of $\mathbb{P}(\theta \mid X)$, we have, almost surely,

$$\mathbb{E}_{\theta \sim \mathbb{P}(\cdot|X)} \left[\log \frac{d \, \mathbb{P}(X \mid \theta_0)}{d \, \mathbb{P}(X \mid \theta)} \right] + \text{KL}(\mathbb{P}_{\theta}(\cdot \mid X) \parallel \pi) = \inf_{\mathbb{Q}} \left\{ \mathbb{E}_{\theta \sim \mathbb{Q}} \left[\log \frac{d \, \mathbb{P}(X \mid \theta_0)}{d \, \mathbb{P}(X \mid \theta)} \right] \text{KL}(\mathbb{Q} \parallel \pi) \right\}. \tag{D.21}$$

Passing to the expectation over X we obtain that,

$$\mathbb{E}\left[\log \frac{d \mathbb{P}(X)}{d \mathbb{P}(X \mid \theta)}\right] + \mathbb{E}_{X}[\mathrm{KL}(\mathbb{P}_{\theta}(\cdot \mid X) \parallel \pi)] \tag{D.22}$$

$$= \mathbb{E}_{X} \left[\inf_{\mathbb{Q}} \left\{ \mathbb{E}_{\theta \sim \mathbb{Q}} \left[\log \frac{d \, \mathbb{P}(X \mid \theta_{0})}{d \, \mathbb{P}(X \mid \theta)} \right] + \text{KL}(\mathbb{Q} \parallel \pi) \right\} \right]$$
 (D.23)

$$\leq \inf_{\mathbb{Q}} \left\{ \mathbb{E}_{\theta \sim \mathbb{Q}} \left[\mathbb{E}_{X} \left[\log \frac{d \mathbb{P}(X \mid \theta_{0})}{d \mathbb{P}(X \mid \theta)} \right] \right] + \text{KL}(\mathbb{Q} \parallel \pi) \right\}$$
 (D.24)

$$= -\log \mathbb{E}_{\theta \sim \pi} \left[\exp \left(-\mathbb{E}_X \left[\log \frac{\mathbb{P}_X(\cdot \mid \theta_0)}{\mathbb{P}_X(\cdot \mid \theta)} \right] \right) \right], \tag{D.25}$$

where the last line follows from Lemma D.1 again with $g(\theta) = -\mathbb{E}_X \left[\log \frac{d \mathbb{P}_X(\cdot | \theta_0)}{d \mathbb{P}_X(\cdot | \theta)} \right]$. Let us now bound the second term in (D.20). We have, by Lemma D.1 again,

$$\mathbb{E}_{X,\theta} \left[\log \frac{d \, \mathbb{P}_X(\cdot \mid \theta_0)}{d \, \mathbb{P}_X(\cdot \mid \theta)} \right] + \mu \, \mathbb{E}_X [\text{KL}(\mathbb{P}_{\theta}(\cdot \mid X) \parallel \pi)] \tag{D.26}$$

$$\geq -\mathbb{E}_{X} \left[\log \mathbb{E}_{\theta \sim \pi} \left[\exp \left(-\frac{1}{\mu} \log \frac{d \, \mathbb{P}_{X}(\cdot \mid \theta_{0})}{d \, \mathbb{P}_{X}(\cdot \mid \theta)} \right) \right] \right]. \tag{D.27}$$

Putting together (D.20), (D.25), and (D.27) concludes the proof.

D.3 ICL setting

Let us now re-introduce the ICL setting from Section 3.1 along with the detailed assumptions.

 $\|\cdot\|$ denotes the Euclidean norm on \mathbb{R}^d for any $d \in \mathbb{N}$. Assume that task vectors live in $\Theta \subset \mathbb{R}^d$ the space of tasks θ and by $\pi(\theta)$ the density of the pretraining task distribution. The context sequence is then generated by first sampling a task θ from the task distribution π , and then sampling data points $(x_t)_{t\geq 1}$ according to

$$x_{t+1} \sim p_{t+1}(\cdot | x_{1:t}, \theta)$$
 (D.28)

where $x_{1:t} = (x_1, \dots, x_t)$.

We denote the posterior $\widehat{p}_t(\theta \mid x_{1:t-1})$ the posterior distribution over tasks given the input sequence $x_{1:t-1}$ Assumption 3 combined with Assumption 4 are the detailed version of Assumption 1 from Section 3.1. Recall that we write $\operatorname{poly}(x)$ to denote a quantity that is $\operatorname{polynomial}$ in x with coefficients independent of the prior π and the number of samples T. We also denote by $\overline{\mathbb{B}}(0,R)$ the closed ball of radius R centered at 0 in \mathbb{R}^d for the Euclidean norm $\|\cdot\|$.

Assumption 3 (Data generation). Fix $\theta^* \in \Theta$ the true task and $\theta_0 \in \Theta$ a reference task such that $\pi(\theta_0) > 0$.

• Tail behaviour of $(x_t)_{t\geq 1}$: there is $k\geq 1$ such that for any $T\geq 1,\ R\geq T,$

$$\mathbb{P}_{X \sim p_T(\cdot \mid \theta^*)} \left(\sup_{\theta : \|\theta\| \ge R} p_T(X \mid \theta) \ge p_T(X \mid \theta_0) \right) \le \frac{\text{poly}(T)}{1 + R^{1/k}}$$
(D.29)

$$\mathbb{P}_{X \sim p_T(\cdot | \theta^*)} \Big(\exists t \le T, \|x_t\| \ge R \Big) \le \frac{\text{poly}(T)}{1 + R^{1/k}} + \tag{D.30}$$

• Moment bound on $(x_t)_{t\geq 1}$: for any $T\geq 1$

$$\mathbb{E}_{X \sim p_T(\cdot \mid \theta^*)} \left[\log^2 \left(\sup_{\theta \in \Theta} \frac{p_T(X \mid \theta)}{p_T(X \mid \theta_0)} \right) \right] \le \text{poly}(T). \tag{D.31}$$

• Regularity of the likelihood: for any $t \ge 1$, $\theta, \theta' \in \Theta \cap \overline{\mathbb{B}}(0, R)$,

$$\sup_{x_{1:t} \in \overline{\mathbb{B}}(0,R)^t} \log \frac{p_t(x_t \mid x_{1:t-1}, \theta)}{p_t(x_t \mid x_{1:t-1}, \theta')} \le \text{poly}(R) \|\theta - \theta'\|.$$
(D.32)

For a sequence $(x_t)_{t\geq 1}$, we denote by $x_{a:b}$ the subsequence $(x_a, x_{a+1}, \dots, x_b)$ for $1\leq a\leq b$ with the convention that $x_{a:b}=x_{1:t}$ if a<1.

D.4 Task Selection Bound for ICL

We begin with a discretization argument and first we generalize the bracketing numbers to the non-i.i.d. case. This definition generalizes the bracketing numbers used in Barron et al. (1999); Zhang (2003, 2006) to the non-i.i.d case and the following result generalises the results of Zhang (2006) to the non-i.i.d. case.

Definition 2. Given a sequence of random variables $(x_t)_{t \leq T}$ on a measurable space \mathcal{X} , with parametric densities $p_t(\cdot|\theta)$ parameterized by $\theta \in \Theta$, compact sets $\Theta' \subset \Theta$ and $\mathcal{X}' \subset \mathcal{X}$, the ε -upper bracketing number of Θ' , denoted by $\mathcal{B}(\Theta', \varepsilon, \mathcal{X}', T)$ is the minimum number of sets U_j that cover Θ' such that, for any $t \leq T - 1$, any $x_{1:t+1} \in \mathcal{X}'^{t+1}$, any j,

$$\int_{\mathcal{X}'} \sup_{\theta \in U_i} p_{t+1}(x_{t+1} \mid x_{1:t}, \theta) dx_{t+1} \le 1 + \varepsilon.$$
 (D.33)

Lemma D.3. For $\mu \in (0,1)$, for any $\varepsilon > 0$ and any compact set $\Theta' \subset \Theta$, any set $\mathcal{X}' \subset \mathcal{X}$, it holds

$$\mu \mathbb{E}_{x_{1:T}} \left[\log \mathbb{E}_{\theta \sim \pi} \left[\exp \left(-\frac{1}{\mu} \log \frac{\mathbf{p}_T(x_{1:T} \mid \theta_0)}{\mathbf{p}_T(x_{1:T} \mid \theta)} \right) \right] \right] \tag{D.34}$$

$$\leq 2\log(\mathcal{B}(\Theta', \varepsilon, \mathcal{X}', T)) + 6T\varepsilon + \pi(\theta \notin \Theta')^{\mu} \tag{D.35}$$

$$+ \mathbb{E}_{x_{1:T}} \left[\mathbb{1} \left\{ \sup_{\theta \notin \Theta'} \frac{\mathbf{p}_{T}(x_{1:T} \mid \theta)}{\mathbf{p}_{T}(x_{1:T} \mid \theta_{0})} \ge 1 \right\} \cdot \log \left(1 + \sup_{\theta \notin \Theta'} \frac{\mathbf{p}_{T}(x_{1:T} \mid \theta)}{\mathbf{p}_{T}(x_{1:T} \mid \theta_{0})} \right) \right] \tag{D.36}$$

$$+ \mathbb{E}_{x_{1:T}} \left[\mathbb{1}\left\{ x_{1:T} \notin \mathcal{X}'^T \right\} \cdot \log \left(\sup_{\theta \in \Theta} \frac{\mathbf{p}_T(x_{1:T} \mid \theta)}{\mathbf{p}_T(x_{1:T} \mid \theta_0)} \right) \right]. \tag{D.37}$$

Proof. First, let us consider $\theta \in \Theta'$ and $X = x_{1:T} \in \mathcal{X}'^T$. We have

$$\exp\left(-\frac{1}{\mu}\log\frac{p_T(X\mid\theta_0)}{p_T(X\mid\theta)}\right) = \exp\left(\frac{1}{\mu}\sum_{t=0}^{T-1}\log\frac{p_{t+1}(x_{t+1}\mid x_{1:t},\theta)}{p_{t+1}(x_{t+1}\mid x_{1:t},\theta_0)}\right)$$
(D.38)

Invoking the bracketing definition (Definition 2), we obtain sets U_j , for $j = 1, ..., \mathcal{B}(\Theta', \varepsilon, \mathcal{X}', T)$ such that, for any $t \leq T - 1$, any $x_{1:t+1} \in \mathcal{X}'^{t+1}$, any j, with $g_j(\cdot | \cdot) := \sup_{\theta \in U_j} p_{t+1}(\cdot | \cdot, \theta)$,

$$\int_{\mathcal{X}'} g_j(x_{t+1} \mid x_{1:t}) dx_{t+1} \le 1 + \varepsilon.$$
 (D.39)

Therefore, for any $\theta \in \Theta'$, any $t \ge 1$, any $x_{1:t+1} \in \mathcal{X}'^{t+1}$, there exists $i \in \{1, \dots, \mathcal{B}(\Theta', \varepsilon, \mathcal{X}', T)\}$ such that

$$p_{t+1}(x_{t+1} \mid x_{t-s:t}, \theta) \le g_i(x_{t+1} \mid x_{t-s:t}). \tag{D.40}$$

Hence, we can bound

$$\exp\left(-\frac{1}{\mu}\log\frac{p_{T}(X\mid\theta_{0})}{p_{T}(X\mid\theta)}\right) \leq \exp\left(\frac{1}{\mu}\sum_{t=0}^{T-1}\log\frac{g_{i}(x_{t+1}\mid x_{t-s:t})}{p_{t+1}(x_{t+1}\mid x_{t-s:t},\theta_{0})} + \frac{T}{\mu}\log\frac{1+\varepsilon}{1-\varepsilon}\right). \tag{D.41}$$

We now control the contribution from $\theta \notin \Theta'$ by simply taking the supremum over this set. We have

$$\mathbb{E}_{\theta \sim \pi} \left[\mathbb{1} \{ \theta \notin \Theta' \} \cdot \exp \left(-\frac{1}{\mu} \log \frac{\mathbf{p}_T(X \mid \theta_0)}{\mathbf{p}_T(X \mid \theta)} \right) \right] \tag{D.42}$$

$$= \pi(\theta \notin \Theta') \sup_{\theta \notin \Theta'} \left(\frac{p_T(X \mid \theta)}{p_T(X \mid \theta_0)} \right)^{1/\mu} . \tag{D.43}$$

Combining (D.41) and (D.43), we bound the LHS of the statement as

$$\mu \mathbb{E}_{X} \left[\mathbb{1}\{X \in \mathcal{X}^{\prime T}\} \log \mathbb{E}_{\theta \sim \pi} \left[\exp \left(-\frac{1}{\mu} \log \frac{\mathbf{p}_{T}(X \mid \theta_{0})}{\mathbf{p}_{T}(X \mid \theta)} \right) \right] \right]$$
(D.44)

$$= \mu \mathbb{E}_{X} \left[\mathbb{1}\{X \in \mathcal{X}'^{T}\} \log \mathbb{E}_{\theta \sim \pi} \left[\mathbb{1}\{\theta \in \Theta'\} \exp \left(-\frac{1}{\mu} \log \frac{p_{T}(X \mid \theta_{0})}{p_{T}(X \mid \theta)} \right) \right] + \mathbb{1}\{\theta \notin \Theta'\} \exp \left(-\frac{1}{\mu} \log \frac{p_{T}(X \mid \theta_{0})}{p_{T}(X \mid \theta)} \right) \right]$$
(D.45)

$$\leq \mu \, \mathbb{E}_{X} \left[\, \mathbb{1}\left\{X \in \mathcal{X}^{\prime T}\right\} \log \left(\, \sum_{i=1}^{\mathcal{B}(\Theta^{\prime}, \varepsilon, \mathcal{X}^{\prime}, T)} \exp \left(\frac{1}{\mu} \, \sum_{t=0}^{T-1} \log \frac{g_{i}(x_{t+1} \mid x_{t-s:t})}{p_{t+1}(x_{t+1} \mid x_{t-s:t}, \theta_{0})} + \frac{T}{\mu} \log \frac{1+\varepsilon}{1-\varepsilon} \right) \right]$$

$$(D.46)$$

$$+ \pi(\theta \notin \Theta') \cdot \sup_{\theta \notin \Theta'} \left(\frac{p_T(X \mid \theta)}{p_T(X \mid \theta_0)} \right)^{1/\mu} \right) \right]. \tag{D.47}$$

Since $\mu \in (0,1)$, for any non-negative numbers a_1, \ldots, a_K we have $\left(\sum_{k=1}^K a_k\right)^{\mu} \leq \sum_{k=1}^K a_k^{\mu}$. Using this inequality and that $\log(a+b) \leq \log(1+a) + \log(1+b)$ for $a,b \geq 0$, we obtain

$$\mu \mathbb{E}_{X} \left[\mathbb{1} \{ X \in \mathcal{X}^{\prime T} \} \log \mathbb{E}_{\theta \sim \pi} \left[\exp \left(-\frac{1}{\mu} \log \frac{p_{T}(X \mid \theta_{0})}{p_{T}(X \mid \theta)} \right) \right] \right]$$
(D.48)

$$\leq \mathbb{E}_{X} \left[\mathbb{1}\left\{X \in \mathcal{X}'^{T}\right\} \log \left(\sum_{i=1}^{\mathcal{B}(\Theta', \varepsilon, \mathcal{X}', T)} \exp \left(\sum_{t=0}^{T-1} \log \frac{g_{i}(x_{t+1} \mid x_{t-s:t})}{p_{t+1}(x_{t+1} \mid x_{t-s:t}, \theta_{0})} + T \log \frac{1+\varepsilon}{1-\varepsilon} \right) \right]$$
(D.49)

$$+ \pi(\theta \notin \Theta')^{\mu} \cdot \sup_{\theta \notin \Theta'} \left(\frac{p_T(X \mid \theta)}{p_T(X \mid \theta_0)} \right)$$
 (D.50)

$$\leq \mathbb{E}_{X} \left[\mathbb{1}\left\{X \in \mathcal{X}'^{T}\right\} \log \left(1 + \sum_{i=1}^{\mathcal{B}(\Theta', \varepsilon, \mathcal{X}', T)} \exp \left(\sum_{t=0}^{T-1} \log \frac{g_{i}(x_{t+1} \mid x_{t-s:t})}{\operatorname{p}_{t+1}(x_{t+1} \mid x_{t-s:t}, \theta_{0})} + T \log \frac{1+\varepsilon}{1-\varepsilon} \right) \right) \tag{D.51}$$

$$+\log\left(1+\pi(\theta\notin\Theta')^{\mu}\cdot\sup_{\theta\notin\Theta'}\left(\frac{p_{T}(X\mid\theta)}{p_{T}(X\mid\theta_{0})}\right)\right)\right].\tag{D.52}$$

Using Jensen's inequality on the first term, we have

$$\mu \mathbb{E}_{X} \left[\mathbb{1}\{X \in \mathcal{X}^{\prime T}\} \log \mathbb{E}_{\theta \sim \pi} \left[\exp \left(-\frac{1}{\mu} \log \frac{\mathbf{p}_{T}(X \mid \theta_{0})}{\mathbf{p}_{T}(X \mid \theta)} \right) \right] \right] \tag{D.53}$$

$$\leq \log \left(1 + \mathbb{E}_X \left[\sum_{i=1}^{\mathcal{B}(\Theta', \varepsilon, \mathcal{X}', T)} \exp \left(\sum_{t=0}^{T-1} \log \frac{g_i(x_{t+1} \mid x_{t-s:t})}{p_{t+1}(x_{t+1} \mid x_{t-s:t}, \theta_0)} + T \log \frac{1+\varepsilon}{1-\varepsilon} \right) \right] \right) \tag{D.54}$$

$$+ \mathbb{E}_{X} \left[\log \left(1 + \pi (\theta \notin \Theta')^{\mu} \cdot \sup_{\theta \notin \Theta'} \left(\frac{p_{T}(X \mid \theta)}{p_{T}(X \mid \theta_{0})} \right) \right) \right]$$
 (D.55)

$$\leq \log \left(1 + \mathcal{B}(\Theta', \varepsilon, \mathcal{X}', T) (1 + \varepsilon)^{T} \left(\frac{1 + \varepsilon}{1 - \varepsilon} \right)^{T} \right) + \mathbb{E}_{X} \left[\log \left(1 + \pi (\theta \notin \Theta')^{\mu} \cdot \mathbb{E}_{X} \left[\sup_{\theta \notin \Theta'} \left(\frac{p_{T}(X \mid \theta)}{p_{T}(X \mid \theta_{0})} \right) \right] \right) \right], \tag{D.56}$$

where we used the definition of the bracketing number Definition 2 in the last line. To obtain the final result, we perform additional manipulations on each term. For the first term, we use that $\frac{1}{1-x} \le 1 + 2x$ for $x \in (0, 1/2)$ so that

$$\log\left((1+\varepsilon)^T \left(\frac{1+\varepsilon}{1-\varepsilon}\right)^T\right) \le \log\left((1+2\varepsilon)^{3T}\right) \le 6T\varepsilon\,,\tag{D.57}$$

so that

$$\log\left(1 + \mathcal{B}(\Theta', \varepsilon, \mathcal{X}', T)(1 + \varepsilon)^T \left(\frac{1 + \varepsilon}{1 - \varepsilon}\right)^T\right) \le \log(1 + \mathcal{B}(\Theta', \varepsilon, \mathcal{X}', T)) + 6T\varepsilon \tag{D.58}$$

$$\leq 2\log(\mathcal{B}(\Theta', \varepsilon, \mathcal{X}', T)) + 6T\varepsilon$$
. (D.59)

For the second term, we use that $\log(1+x) \le x$ and distinguish two cases to obtain

$$\mathbb{E}_{X} \left[\log \left(1 + \pi (\theta \notin \Theta')^{\mu} \cdot \mathbb{E}_{X} \left[\sup_{\theta \notin \Theta'} \left(\frac{p_{T}(X \mid \theta)}{p_{T}(X \mid \theta_{0})} \right) \right] \right) \right]$$
 (D.60)

$$\leq \pi(\theta \notin \Theta')^{\mu} + \mathbb{E}_{X} \left[\mathbb{1} \left\{ \sup_{\theta \notin \Theta'} \frac{p_{T}(X \mid \theta)}{p_{T}(X \mid \theta_{0})} \geq 1 \right\} \cdot \log \left(1 + \sup_{\theta \notin \Theta'} \frac{p_{T}(X \mid \theta)}{p_{T}(X \mid \theta_{0})} \right) \right]. \tag{D.61}$$

All that is left to do is to deal with the case $X \notin \mathcal{X}^{\prime T}$. We have, as above,

$$\mu \mathbb{E}_{X} \left[\mathbb{1}\{X \notin \mathcal{X}'^{T}\} \log \mathbb{E}_{\theta \sim \pi} \left[\exp \left(-\frac{1}{\mu} \log \frac{p_{T}(X \mid \theta_{0})}{p_{T}(X \mid \theta)} \right) \right] \right] \leq \mathbb{E}_{X} \left[\mathbb{1}\{X \notin \mathcal{X}'^{T}\} \log \left(\sup_{\theta \in \Theta} \frac{p_{T}(X \mid \theta)}{p_{T}(X \mid \theta_{0})} \right) \right]. \tag{D.62}$$

29

We now leverage Assumption 3 to control the different terms of Lemma D.3.

Lemma D.4. For $\mu \in (0,1)$, under Assumption 3, for any $T \ge 1$, it holds that

$$\mu \mathbb{E}_{x_{1:T}} \left[\log \mathbb{E}_{\theta \sim \pi} \left[\exp \left(-\frac{1}{\mu} \log \frac{p_T(x_{1:T} \mid \theta_0)}{p_T(x_{1:T} \mid \theta)} \right) \right] \right] \le \pi (\theta \notin \Theta')^{\mu} + \mathcal{O}(\log(T)), \tag{D.63}$$

where the $\mathcal{O}(\cdot)$ hides constants that do not depend on π or T.

Proof. Fix R > 0 that will be chosen later and take $\mathcal{X}' = \overline{\mathbb{B}}(0, R)$ and $\Theta' = \overline{\mathbb{B}}(0, R)$. Let us consider a δ -cover of Θ' with $\delta > 0$ that will be chosen later: there are K sets U_j , $j = 1, \ldots, K$ that cover Θ' such that for any $\theta, \theta' \in U_j$, we have $\|\theta - \theta'\| \le \delta$. By e.g., Wainwright (2019, Ex. 5.2), we can take K such that $\log K \le d \log(1 + 2R/\delta)$.

Assumption 3 ensures that the sets U_j satisfy the bracketing condition of Definition 2 with $\varepsilon = \exp(\text{poly}(R)\delta)$ –1. Therefore, we have, with this choice of ε ,

$$\log \mathcal{B}(\Theta', \varepsilon, \mathcal{X}', T) \le d \log(1 + 2R/\delta). \tag{D.64}$$

Using Cauchy-Schwarz inequality and Assumption 3, we have that, both

$$\mathbb{E}_{x_{1:T}} \left[\mathbb{1} \left\{ \sup_{\theta \notin \Theta'} \frac{\mathbf{p}_T(x_{1:T} \mid \theta)}{\mathbf{p}_T(x_{1:T} \mid \theta_0)} \ge 1 \right\} \cdot \log \left(1 + \sup_{\theta \notin \Theta'} \frac{\mathbf{p}_T(x_{1:T} \mid \theta)}{\mathbf{p}_T(x_{1:T} \mid \theta_0)} \right) \right] \le \frac{\mathrm{poly}(T)}{1 + R^{1/k}}$$
(D.65)

$$\mathbb{E}_{x_{1:T}} \left[\mathbb{1}\left\{ x_{1:T} \notin \mathcal{X}'^T \right\} \cdot \log \left(\sup_{\theta \in \Theta} \frac{\mathbf{p}_T(x_{1:T} \mid \theta)}{\mathbf{p}_T(x_{1:T} \mid \theta_0)} \right) \right] \le \frac{\mathbf{poly}(T)}{1 + R^{1/k}} . \tag{D.66}$$

Choose R = poly(T) so that both (D.65) and (D.66) are $\mathcal{O}(1)$. Finally, we choose $\delta = (\text{poly}(T))^{-1}$ so that $\varepsilon = \exp(\text{poly}(R)\delta) - 1 = \mathcal{O}(1/T)$. Combining this (D.64)–(D.66) with Lemma D.3 concludes the proof.

We can now state our main result for ICL. As a metric to asses the quality of a given retrieved task θ w.r.t. the true task θ^* , we consider the Rényi divergence (Rényi, 1961) of order $\rho \in (0,1)$ between the distributions $p_T(\cdot | \theta)$ and $p_T(\cdot | \theta^*)$:

$$D_{\rho}(\theta \parallel \theta^*) = -\frac{1}{T(1-\rho)} \log \mathbb{E}_{X \sim p_T(\cdot \mid \theta^*)} \left[\prod_{t=1}^T \left(\frac{p_t(x_t \mid x_{1:t-1}, \theta)}{p_t(x_t \mid x_{1:t-1}, \theta^*)} \right)^{\rho} \right].$$
 (D.67)

Theorem D.1. Under Assumption 3, for any $\rho \in (0,1)$, $T \ge 1$, it holds that, for $x_{1:T} \sim p_T(\cdot \mid \theta^*)$,

$$\mathbb{E}_{x_{1:T}} \left[\mathbb{E}_{\theta \sim \widehat{p}_{T}(\cdot \mid x_{1:T})} \left[\mathcal{D}_{\rho}(\theta \parallel \theta^{*}) \right] \right]$$
 (D.68)

$$\leq -\frac{1+\rho}{(1-\rho)T}\log\left(\mathbb{E}_{\theta\sim\pi}\left[\exp\left(-\mathbb{E}_{x_{1:T}}\left[\log\frac{p_T(x_{1:T}\mid\theta_0)}{p_T(x_{1:T}\mid\theta)}\right]\right)\right]\right) \tag{D.69}$$

$$+\frac{1+\rho}{1-\rho}\frac{\mathrm{KL}(\mathrm{p}_{T}(\cdot\,|\,\theta^{*})\,\|\,\mathrm{p}_{T}(\cdot\,|\,\theta_{0}))}{T}\tag{D.70}$$

$$+\mathcal{O}\left(\frac{\log(T)}{T}\right),$$
 (D.71)

where the $\mathcal{O}(\cdot)$ hides constants that do not depend on π or T.

Proof. This is a direct consequence of Proposition D.1 combined with Lemma D.4 with $\lambda = 1 + \rho$ and bounding $\pi(\theta \notin \Theta')^{\mu} \leq 1$.

A few comments are in order. The first term of (D.69) captures how much the prior π covers the reference task θ_0 . When $\theta_0 = \theta^*$, this term thus quantifies how well the prior covers the true task θ^* . When θ_0 is inside the support of π , this term is vanishing as T grows large, see the next results below.

The second term of (D.70) captures how well the reference task θ_0 approximates the true task θ^* . When $\theta_0 = \theta^*$, the term of (D.70) is 0. Otherwise, consider the case the KL will typically be of order T so that this term is $\mathcal{O}(1)$: it represents the best ICL error one can hope for when the true task θ^* is not in the support of the prior π .

D.5 Laplace Approximation

We will make use of the following version of the Laplace approximation, see Wong (2001, Chap. 9, Thm. 3) for a proof.

Lemma D.5 (Laplace approximation). Let μ be a probability measure on \mathbb{R}^d with density $g : \mathbb{R}^d \to [0, \infty)$. Fix $x^* \in \mathbb{R}^d$ such that g is continuous at x^* and $g(x^*) > 0$. Then, as $\varepsilon \to 0$,

$$\int_{\mathbb{R}^d} \exp\left(-\frac{1}{2\varepsilon} \|x - x^*\|\right) g(x) \, dx, = g(x^*) \, C \, \varepsilon^d + o(\varepsilon^d).$$

where $C := \int_{\mathbb{R}^d} \exp\left(-\frac{1}{2} \|y\|\right) dy \in (0, \infty).$

Assumption 4. Consider the following additional assumptions to Assumption 3:

• Tail behaviour: for any $T \ge 1$, R > 0,

$$\mathbb{P}_{X \sim p_T(\cdot \mid \theta^*)} \left(\sup_{\theta : \|\theta\| \ge R} p_T(X \mid \theta) \ge p_T(X \mid \theta_0) \right) \le \operatorname{poly}(T) e^{-R}$$
(D.72)

$$\mathbb{P}_{X \sim p_T(\cdot \mid \theta^*)} \Big(\exists t \le T, \|x_t\| \ge R \Big) \le \text{poly}(T) e^{-R} . \tag{D.73}$$

- Regularity of π : π is continuous and positive at θ_0 .
- Second moment of π :

$$\mathbb{E}_{\theta \sim \pi} \left[\|\theta\|^2 \right] < \infty. \tag{D.74}$$

Proposition D.2. Under Assumptions 3 and 4, then, for T large enough,

$$-\log\left(\mathbb{E}_{\theta \sim \pi}\left[\exp\left(-\mathbb{E}_{x_{1:T}}\left[\log\frac{p_T(x_{1:T}\mid\theta_0)}{p_T(x_{1:T}\mid\theta)}\right]\right)\right]\right) \leq \log 1/\pi(\theta_0) + \mathcal{O}(\operatorname{poly}(\log T)). \tag{D.75}$$

Proof. For some $R_T \ge r_T > 0$, we split the term as

$$-\log\left(\mathbb{E}_{\theta \sim \pi}\left[\exp\left(-\mathbb{E}_{x_{1:T}}\left[\log\frac{\mathbf{p}_{T}(x_{1:T}\mid\theta_{0})}{\mathbf{p}_{T}(x_{1:T}\mid\theta)}\right]\right)\right]\right) \tag{D.76}$$

$$=-\log\left(\mathbb{E}_{\theta\sim\pi}\left[\mathbb{1}\{\|\theta\|\leq R_T\}\exp\left(-\mathbb{E}_{x_1:T}\left[\log\frac{\mathbf{p}_T(x_1:T\mid\theta_0)}{\mathbf{p}_T(x_1:T\mid\theta)}\right]\right)+\mathbb{1}\{\|\theta\|>R_T\}\exp\left(-\mathbb{E}_{x_1:T}\left[\log\frac{\mathbf{p}_T(x_1:T\mid\theta_0)}{\mathbf{p}_T(x_1:T\mid\theta)}\right]\right)\right]\right)$$
(D.77)

$$\leq -\log \left(\mathbb{E}_{\theta \sim \pi} \left[\mathbb{1}\{\|\theta\| \leq r_T\} \exp \left(-\mathbb{E}_{x_{1:T}} \left[\log \frac{\mathbf{p}_T(x_{1:T} \mid \theta_0)}{\mathbf{p}_T(x_{1:T} \mid \theta)}\right]\right) + \mathbb{1}\{\|\theta\| > R_T\} \exp \left(-\mathbb{E}_{x_{1:T}} \left[\log \frac{\mathbf{p}_T(x_{1:T} \mid \theta_0)}{\mathbf{p}_T(x_{1:T} \mid \theta)}\right]\right)\right) \right) \tag{D.78}$$

Using Cauchy-Schwarz inequality and Assumption 3 and its refinement in the statement, we bound the second term as, for θ such that $\|\theta\| > R_T$, so that

$$\left| \mathbb{E}_{x_{1:T}} \left[\log \frac{\mathbf{p}_T(x_{1:T} \mid \theta_0)}{\mathbf{p}_T(x_{1:T} \mid \theta)} \right] \right| \le e^{-R_T/2} \operatorname{poly}(T) . \tag{D.79}$$

so that

$$\mathbb{E}_{\theta \sim \pi} \left[\mathbb{1}\{\|\theta\| > R_T\} \exp\left(-\mathbb{E}_{x_{1:T}} \left[\log \frac{\mathbf{p}_T(x_{1:T} \mid \theta_0)}{\mathbf{p}_T(x_{1:T} \mid \theta)} \right] \right) \right] \tag{D.80}$$

$$\leq \exp\left(e^{-R_T/2}\operatorname{poly}(T)\right)\pi(\|\theta\| > R_T) \tag{D.81}$$

$$\leq \exp\left(e^{-R_T/2}\operatorname{poly}(T)\right) \frac{\mathbb{E}_{\theta \sim \pi}\left[\|\theta\|^2\right]}{R_T^2},\tag{D.82}$$

where we used Markov's inequality in the last line. Take $R_T = T^{(d+1)}/2$ so that (D.82) is $\mathcal{O}(1/T^{d+1})$. We now focus on the first term of (D.78) and bound it as:

$$\mathbb{E}_{x_{1:T}}\left[\log\frac{\mathbf{p}_{T}(x_{1:T}\mid\theta_{0})}{\mathbf{p}_{T}(x_{1:T}\mid\theta)}\right] = \mathbb{E}_{x_{1:T}}\left[\mathbb{1}\left\{\max_{t}||x_{t}|| \leq r_{T}\right\}\log\frac{\mathbf{p}_{T}(x_{1:T}\mid\theta_{0})}{\mathbf{p}_{T}(x_{1:T}\mid\theta)}\right] + \mathbb{E}_{x_{1:T}}\left[\mathbb{1}\left\{\max_{t}||x_{t}|| > r_{T}\right\}\log\frac{\mathbf{p}_{T}(x_{1:T}\mid\theta_{0})}{\mathbf{p}_{T}(x_{1:T}\mid\theta)}\right]$$
(D.83)

$$\leq \operatorname{poly}(r_T)T\|\theta - \theta_0\| + \operatorname{poly}(T)e^{-r_T/2} \tag{D.84}$$

where we used the regularity assumption of Assumption 3 for the first term and Cauchy-Schwarz inequality combined with Assumption 4 for the second term.

Take $r_T = \text{poly}(\log T)$ so that $\text{poly}(T)e^{-r_T/2} = \mathcal{O}(1)$ and assume that T is large enough so that $r_T \ge \|\theta_0\| + 1$. Putting everything together, we have

$$-\log\left(\mathbb{E}_{\theta \sim \pi}\left[\exp\left(-\mathbb{E}_{x_{1:T}}\left[\log\frac{\mathbf{p}_{T}(x_{1:T}\mid\theta_{0})}{\mathbf{p}_{T}(x_{1:T}\mid\theta)}\right]\right)\right]\right) \tag{D.85}$$

$$\leq -\log \left(\mathbb{E}_{\theta \sim \pi} \left[\mathbb{1}\{\|\theta\| \leq r_T\} \exp(-\operatorname{poly}(r_T)T\|\theta - \theta_0\| + \mathcal{O}(1)) + \mathcal{O}\left(\frac{1}{T^{d+1}}\right) \right] \right) \tag{D.86}$$

$$\leq -\log\left(\mathbb{E}_{\theta \sim \pi}\left[\mathbb{1}\{\|\theta\| \leq \|\theta_0\| + 1\}\exp(-\operatorname{poly}(\log T)T\|\theta - \theta_0\| + \mathcal{O}(1)) + \mathcal{O}\left(\frac{1}{T^{d+1}}\right)\right]\right),\tag{D.87}$$

where we used that we assumed that $r_T = \text{poly}(\log T) \ge ||\theta_0|| + 1$.

Applying Lemma D.5 with $\varepsilon = 1/(\text{poly}(\log T)T)$ yields:

$$\mathbb{E}_{\theta \sim \pi} [\mathbb{1}\{\|\theta\| \le \|\theta_0\| + 1\} \exp(-\operatorname{poly}(\log T)T\|\theta - \theta_0\|)] = \operatorname{poly}(\log T)T^{-d}(\pi(\theta_0)C + o(1)), \tag{D.88}$$

where C is the constant of Lemma D.5 and this concludes the proof.

We can now combine Theorem D.1 and Proposition D.2 to obtain the final result in the main text.

Theorem D.2. Under Assumptions 3 and 4, for any $\rho \in (0,1)$, $T \ge 1$, it holds that, for $x_{1:T} \sim p_T(\cdot \mid \theta^*)$,

$$\mathbb{E}_{x_{1:T}} \left[\mathbb{E}_{\theta \sim \widehat{p}_{T}(\cdot \mid x_{1:T})} \left[\mathcal{D}_{\rho}(\theta \parallel \theta^{*}) \right] \right]$$
 (D.89)

$$\leq -\frac{1+\rho}{(1-\rho)T}\log 1/\pi(\theta_0) \tag{D.90}$$

$$+\frac{1+\rho}{1-\rho}\frac{\mathrm{KL}(\mathrm{p}_{T}(\cdot\,|\,\theta^{*})\,\|\,\mathrm{p}_{T}(\cdot\,|\,\theta_{0}))}{T}\tag{D.91}$$

$$+\mathcal{O}\left(\frac{\log(T)}{T}\right),$$
 (D.92)

where the $\mathcal{O}(\cdot)$ hides constants that do not depend on π or T.

Proof. This is a direct consequence of Theorem D.1 and Proposition D.2.

E Generalization bounds

E.1 Moment bounds for general functions

In this subsection, we generalize the heavy-tail concentration results of Li and Liu (2024a) to allow for non-i.i.d. data. This section can also be seen as extending concentration results for dependent sequences to the case where the function of interest does not necessarily admit bounded differences but only bounded moments. In particular, Lemma E.1 extends the coupling argument of Chazottes et al. (2007) to our setting, in particular not requiring bounded differences but only bounded moments. Indeed, for this, we replace the total variation distance by the Wasserstein-1 distance. It can also be seen as an extension of the bounded differences result of Kontorovich and Ramanan (2008) to our setting (see Mohri and Rostamizadeh (2010) for a presentation of the results of Kontorovich and Ramanan (2008) in a setting closer to ours). Moreover, note that even the handling of the subGaussian increments is much more trickier than in Kontorovich (2014), since we have to carefully apply a convex domination argument to handle the conditional dependence. The main result of this section is Theorem E.1, which is of independent interest.

As in the previous section, $\|\cdot\|$ denotes the Euclidean norm on \mathbb{R}^d for any $d \in \mathbb{N}$.

At multiple places, we will use the Wasserstein-1 distance[§] with respect to a cost function $\rho: \mathbb{Z} \times \mathbb{Z} \to [0, \infty)$, defined as

$$W_{\rho}(\mu,\nu) := \inf_{\pi \in \Pi(\mu,\nu)} \int \rho(z,z') d\pi(z,z'), \tag{E.1}$$

where $\Pi(\mu, \nu)$ is the set of couplings of μ and ν . We refer to the textbook Villani (2008) for more details.

Lemma E.1. Consider \mathcal{Z} measurable space. Let Z_1, \ldots, Z_m be \mathcal{Z} -valued random variables with natural filtration $\mathcal{F}_i := \sigma(Z_1, \ldots, Z_i)$. For each i, assume there is Z_i' such that

$$Z_i' \sim Law(Z_i \mid \mathcal{F}_{i-1}), \quad Z_i' \perp \!\!\! \perp Z_i \mid \mathcal{F}_{i-1}.$$
 (E.2)

Let $g: \mathbb{Z}^m \to \mathbb{R}$ be measurable and coordinate-wise Lipschitz with respect to cost functions $\rho_i: \mathbb{Z} \times \mathbb{Z} \to [0, \infty)$ such that $\rho_i(z_i, z_i) = 0$, with constants $L_i \geq 0$: for any $z, z' \in \mathbb{Z}^m$ differing only in the i-th coordinate,

$$|g(z) - g(z')| \le L_i \rho_i(z_i, z_i'). \tag{E.3}$$

With $W_{\rho_i}(\cdot, \cdot)$ the Wasserstein-1 distance with respect to ρ_j , define, for i < j,

$$\delta_{i,j}(z_{1:i}, z_i') = W_{\rho_i}(Law(Z_i \mid Z_{1:i} = z_{1:i}), Law(Z_i \mid Z_{1:i-1} = z_{1:i-1}, Z_i = z_i')). \tag{E.4}$$

for $i \in \{1, \ldots, m\}$,

$$|\mathbb{E}[g(Z_{1:m}) | \mathcal{F}_i] - \mathbb{E}[g(Z_{1:i-1}, Z_i', Z_{i+1:m}) | \mathcal{F}_{i-1}, Z_i']| \le L_i \rho_i(Z_i, Z_i') + \sum_{i=i+1}^m L_j \delta_{i,j}(Z_{1:i}, Z_i')$$
(E.5)

Proof. Fix $i \in \{1, ..., m\}$. We condition on \mathcal{F}_{i-1} . Let $u := Z_i$ and $u' := Z'_i$. Not to overburden notations, all expectations and probabilities in the following are conditional on $\mathcal{F}_{i-1}, Z_i = u, Z'_i = u'$. Define the tail functions

$$\psi(z_{i+1:m}) := g(Z_{1:(i-1)}, u, z_{i+1:m}), \tag{E.6}$$

$$\psi'(z_{i+1:m}) := q(Z_{1:(i-1)}, u', z_{i+1:m}). \tag{E.7}$$

 $\text{Denote } Z_{(i+1):m} \sim \operatorname{Law}(Z_{(i+1):m} \mid \mathcal{F}_{i-1}, Z_i = u) \text{ and } Z'_{(i+1):m} \sim \operatorname{Law}(Z_{(i+1):m} \mid \mathcal{F}_{i-1}, Z_i = u'). \text{ We decompose } Z_{(i+1):m} \sim \operatorname{Law}(Z_{(i+1):m} \mid \mathcal{F}_{i-1}, Z_i = u').$

$$|\mathbb{E}[g(Z_{1:m})] - \mathbb{E}[g(Z_{1:(i-1)}, Z'_{i:m})]|$$
 (E.8)

$$= \left| \mathbb{E}[\psi(Z_{(i+1):m})] - \mathbb{E}[\psi'(Z'_{(i+1):m})] \right|$$
 (E.9)

$$\leq \mathbb{E}\left[\left|\psi(Z_{(i+1):m}) - \psi'(Z_{(i+1):m})\right|\right] + \left|\mathbb{E}\left[\psi'(Z_{(i+1):m})\right] - \mathbb{E}\left[\psi'(Z_{(i+1):m}')\right]\right|. \tag{E.10}$$

We bound the two terms separately.

[§] This is a slight abuse of terminology, since the Wasserstein-1 distance is usually defined for metric spaces, while we only assume ρ to be a cost function. However, this slight abuse of terminology will not cause any confusion in the following.

By the coordinate-wise Lipschitz condition at i,

$$\mathbb{E}_{P}[|\psi(Z_{(i+1):m}) - \psi'(Z_{(i+1):m})|] \le L_{i}\rho_{i}(u, u') = L_{i}\rho_{i}(Z_{i}, Z'_{i}). \tag{E.11}$$

We write the following telescoping decomposition:

$$\left| \mathbb{E} \left[\psi'(Z_{(i+1):m}) \right] - \mathbb{E} \left[\psi'(Z'_{(i+1):m}) \right] \right| \leq \sum_{j=i}^{m-1} \left| \mathbb{E} \left[\psi'(Z'_{(i+1):j}, Z_{(j+1):m}) \right] - \mathbb{E} \left[\psi'(Z'_{(i+1):(j+1)}, Z_{(j+1):m}) \right] \right|. \tag{E.12}$$

By the definition of the Wasserstein-1 distance, there exists a coupling of (Z_{j+1}, Z'_{j+1}) such that

$$\mathbb{E}\left[\rho_{j+1}(Z_{j+1}, Z'_{j+1}) \middle| \mathcal{F}_{i}, Z'_{i}\right] = W_{\rho_{j+1}}(\text{Law}(Z_{j+1} | \mathcal{F}_{i}), \text{Law}(Z_{j+1} | \mathcal{F}_{i-1}, Z'_{i})) \le \delta_{i,j+1}(Z_{1:i-1}, Z'_{i}). \tag{E.13}$$

We obtain a bound on the increment at coordinate j by combining the coupling with the coordinate-wise Lipschitz condition at j:

$$\left| \mathbb{E} \left[\psi'(Z'_{(i+1):j}, Z_{(j+1):m}) \right] - \mathbb{E} \left[\psi'(Z'_{(i+1):(j+1)}, Z_{(j+1):m}) \right] \right|$$
 (E.14)

$$\leq \mathbb{E}\left[\left|\psi'(Z'_{(i+1):j}, Z_{(j+1):m}) - \psi'(Z'_{(i+1):(j+1)}, Z_{(j+1):m})\right|\right] \tag{E.15}$$

$$\leq L_{j+1} \mathbb{E} \Big[\rho_{j+1}(Z_{j+1}, Z'_{j+1}) \Big]$$
 (E.16)

$$= L_{i+1} W_{\rho_{i+1}} \left(\text{Law}(Z_{i+1} \mid \mathcal{F}_i), \text{Law}(Z_{i+1} \mid \mathcal{F}_{i-1}, Z_i') \right) = L_{i+1} \delta_{i, j+1} (Z_{1:i}, Z_i'). \tag{E.17}$$

Combining the above estimates gives

$$\left| \mathbb{E} \left[\psi'(Z_{(i+1):m}) \right] - \mathbb{E} \left[\psi'(Z'_{(i+1):m}) \right] \right| \le \sum_{j=i}^{m-1} L_{j+1} \delta_{i,j+1}(Z_{1:i}, Z'_{i}). \tag{E.18}$$

which yields the desired result.

We now state a classic convex domination lemma which is a slight variant of Ledoux and Talagrand (2013, Lem. 4.6).

Lemma E.2 (Convex domination). Consider X, Z a zero-mean symmetric random variables such that

$$\mathbb{P}(|X| > t) \le C \, \mathbb{P}(|Z| > t) \,, \tag{E.19}$$

for some C > 0 and all t > 0.

Then, for any convex function $h: \mathbb{R} \to \mathbb{R}$,

$$\mathbb{E}[h(X)] \le \mathbb{E}[h(CZ)]. \tag{E.20}$$

Proof. Let $\delta \sim \text{Bernoulli}(1/C)$ be independent of (X, Z). Then, for all t > 0, $\mathbb{P}(|Z| > t) \geq \frac{1}{C} \mathbb{P}(|X| > t) = \mathbb{P}(|\delta X| > t)$. Hence $|\delta X|$ is stochastically dominated by |Z| and we may construct a coupling such that

$$|\delta X| \le |Z|$$
 a.s. (E.21)

Since X is symmetric, we may write in distribution $X \stackrel{d}{=} \varepsilon |X|$ where ε is a Rademacher variable independent of |X|. Likewise, $Z \stackrel{d}{=} \varepsilon' |Z|$ with an independent Rademacher ε' .

Condition on (δ, X, Z) and define

$$\Phi(a) := \mathbb{E}\left[h(a \,\varepsilon \,|Z|) \mid \delta, X, Z\right], \qquad a \in [-1, 1]. \tag{E.22}$$

The map $a \mapsto \Phi(a)$ is convex (as an average of convex functions). By convexity, its maximum on [-1,1] is attained at an extreme point $\{-1,1\}$. On the coupling where (E.21) holds, define

$$a := \begin{cases} \frac{\delta |X|}{|Z|}, & \text{if } Z \neq 0, \\ 0, & \text{if } Z = 0, \end{cases}$$
 (E.23)

so that $a \in [-1, 1]$ almost surely thanks to $|X| \leq |\delta Z|$. Therefore,

$$\mathbb{E}\left[h(\varepsilon|X|\delta) \mid \delta, X, Z\right] = \Phi(a) \le \max\{\Phi(-1), \Phi(1)\} = \mathbb{E}\left[h(\varepsilon|Z|) \mid \delta, |X|, Z\right]. \tag{E.24}$$

Taking expectations and using $X \stackrel{d}{=} \varepsilon |X|$ and $Z \stackrel{d}{=} \varepsilon |Z|$,

$$\mathbb{E}[h(\delta X)] \le \mathbb{E}[h(Z)]. \tag{E.25}$$

Since h is convex and $\mathbb{E}[\delta \mid X, Z] = 1/C$, we have, by Jensen's inequality,

$$\mathbb{E}[h(X/C)] = \mathbb{E}[h(\mathbb{E}[\delta X \mid X, Z])] \le \mathbb{E}[\mathbb{E}[h(\delta X) \mid X, Z]] = \mathbb{E}[h(\delta X)] \le \mathbb{E}[h(Z)], \tag{E.26}$$

Finally, apply the previous inequality with the convex function $u \mapsto h(Cu)$ to obtain

$$\mathbb{E}[h(X)] = \mathbb{E}[h(C \cdot (X/C))] \le \mathbb{E}[h(CZ)].$$

This is exactly the desired bound.

We now state a fact of subGaussian random variables, which can be found in Wainwright (2019, Thm. 2.6) for instance.

Lemma E.3 (Convex domination). Consider X a zero-mean real-valued σ^2 -sub-Gaussian random variable, which is, in addition, symmetric, i.e., $X \stackrel{d}{=} -X$. Then, for $Z \sim \mathcal{N}(0, \sigma^2)$,

$$\mathbb{P}(|X| > t) \le 8\,\mathbb{P}(|Z| > t). \tag{E.27}$$

Lemma E.4 (Causal symmetrization). Let $m \in \mathbb{N}$ and $(\mathcal{Z}, \mathcal{A})$ be a standard Borel measurable space. Let Z_1, \ldots, Z_m be \mathcal{Z} -valued random with natural filtration $(\mathcal{F}_i)_{i=0,\ldots,m}$ Let $h \colon \mathbb{R} \to \mathbb{R}$ be convex.

Consider $g: \mathbb{Z}^m \to \mathbb{R}$ be measurable. Set $S := g(Z_1, \ldots, Z_m)$. For each $i \in \{1, \ldots, m\}$, assume there exists a conditionally independent resample

$$Z_i' \sim \text{Law}(Z_i \mid \mathcal{F}_{i-1}), \quad Z_i' \perp \!\!\! \perp Z_i \mid \mathcal{F}_{i-1}.$$
 (E.28)

Let $\varepsilon_{1:m}$, $\varepsilon'_{1:m}$ be independent Rademacher variables, independent of all Z, Z' and \mathcal{F}_m .

Assume there exist measurable functions $c_i \colon \mathcal{Z} \times \mathcal{Z} \to [0, \infty)$, $d_i \colon \mathcal{Z} \to [0, \infty)$ and $J \subset \{1, \dots, m\}$ such that, the following conditions hold:

(i) For any i, there exists $j(i) \in J$, such that, for any $z_{1:i-1} \in \mathbb{Z}^{i-1}$ and $z_i, z_i' \in \mathbb{Z}$,

$$\left| \mathbb{E}[S \mid Z_{1:i} = z_{1:i}] - \mathbb{E}[S \mid Z_{1:i-1} = z_{1:i-1}, Z_i = z_i'] \right| \le c_i(z_i, z_i') + d_i(z_{j(i)}) \, \mathbb{1}\{i \notin J\} \,. \tag{E.29}$$

- (ii) For any $i \notin J$, $\varepsilon_i c_i(Z_i, Z_i')$ is σ_i^2 -sub-Gaussian conditionally on \mathcal{F}_{i-1} .
- (iii) For any $j \in J$, Z_j is independent of \mathcal{F}_{j-1} .

Then, there are Gaussian random variables $G_j, G'_j \sim \mathcal{N}(0, 8\sigma_j^2)$ independent and independent of all $Z, Z', \varepsilon, \mathcal{F}_m$ such that

$$\mathbb{E}[h(S - \mathbb{E}[S])] \le \mathbb{E}\left[h\left(\sum_{i \notin J} \operatorname{Sym}_{j(i)}\left(\varepsilon_i(|G_i| + d_i(Z_{j(i)}))\right) + \sum_{i \in J} \varepsilon_i c_j(Z_j, Z_j')\right)\right],\tag{E.30}$$

where we use the notation:

$$\operatorname{Sym}_{j(i)}\left(\varepsilon_{i}\left(|G_{i}|+d_{i}(Z_{j(i)})\right)\right) \coloneqq \varepsilon_{j(i)}\left(\varepsilon_{i}\left(|G_{i}|+d_{i}(Z_{j(i)})\right)-\varepsilon'_{i}(|G'_{i}|+d_{i}(Z'_{j(i)}))\right). \tag{E.31}$$

Proof. Define $\mathcal{G} = \sigma(\varepsilon_{1:m}, G_{1:m})$.

We show the result by induction on k: our goal is to show that, for any $k \in \{0, ..., m\}$,

$$\mathbb{E}[h(S - \mathbb{E}[S])] \leq \mathbb{E}\left[h\left(\sum_{\substack{i \notin J \\ i \geq k+1}} \left(\mathbb{1}\{j(i) \leq k\}\varepsilon_i\left(|G_i| + d_i(Z_{j(i)})\right) + \mathbb{1}\{j(i) \geq k+1\}\operatorname{Sym}_{j(i)}\left(\varepsilon_i\left(|G_i| + d_i(Z_{j(i)})\right)\right)\right)\right]\right]$$

(E.32)

$$+ \sum_{\substack{i \in J \\ i \ge k+1}} \varepsilon_i c_i(Z_i, Z_i') + \mathbb{E}[S \mid Z_{1:k}] - \mathbb{E}[S] \bigg) \bigg], \tag{E.33}$$

where $G_i, G'_i \sim \mathcal{N}(0, 8\sigma_i^2)$ are independent and independent of all $Z, Z', \varepsilon, \varepsilon', \mathcal{F}_m$. (E.33) holds trivially for k=m. We now show that if it holds for some $k\in\{1,\ldots,m\}$, then it also holds for k-1.

Note that we can rewrite

$$\sum_{\substack{i \notin J \\ i \geq k+1}} \left(\mathbb{1}\{j(i) \leq k\} \varepsilon_i \left(|G_i| + d_i(Z_{j(i)}) \right) + \mathbb{1}\{j(i) \geq k+1\} \operatorname{Sym}_{j(i)} \left(\varepsilon_i \left(|G_i| + d_i(Z_{j(i)}) \right) \right) \right) \tag{E.34}$$

$$+\sum_{\substack{i\in J\\i\geq k+1}}\varepsilon_ic_i(Z_i,Z_i')\tag{E.35}$$

$$= \underbrace{\sum_{\substack{i \notin J \\ i \ge k+1}}} \mathbb{1}\{j(i) \ge k+1\} \operatorname{Sym}_{j(i)} \left(\varepsilon_{i} \left(|G_{i}| + d_{i}(Z_{j(i)})\right)\right) + \sum_{\substack{i \in J \\ i \ge k+1}}} \varepsilon_{i} c_{i}(Z_{i}, Z_{i}')$$

$$=: Y_{\perp}$$

$$+ \sum_{\substack{i \in J \\ i \ge k+1}} \mathbb{1}\{j(i) \le k\} \varepsilon_{i} \left(|G_{i}| + d_{i}(Z_{j(i)})\right)$$
(E.36)

$$+ \underbrace{\sum_{\substack{i \notin J \\ i \ge k+1}} \mathbb{1}\{j(i) \le k\} \varepsilon_i (|G_i| + d_i(Z_{j(i)}))}_{=:Y_k} \tag{E.37}$$

$$=Y_{\perp}+Y_{k}, \qquad (E.38)$$

where Y_{\perp} is independent of \mathcal{F}_k and Y_k is \mathcal{F}_k -measurable. More precisely, we show that

$$\mathbb{E}[h(Y_{\perp} + Y_k + \mathbb{E}[S \mid Z_{1:k}] - \mathbb{E}[S]) \mid Y_{\perp}]$$
(E.39)

$$\leq \mathbb{E}\left[h\left(Y_{\perp} + Y_{k-1} + \mathbb{1}\{k \notin J\}\varepsilon_{k}(|G_{k}| + d_{k}(Z_{j(k)}))\right)\right] \tag{E.40}$$

$$+ \mathbb{1}\{k \in J\}(\varepsilon_k c_k(Z_k, Z_k')) \tag{E.41}$$

$$+ \sum_{\substack{i \notin J \\ i \geq k+1 \\ j(i) = k}} \operatorname{Sym}_{k}(\varepsilon_{i}(|G_{i}| + d_{i}(Z_{k})))) \mathbb{E}[S \mid Z_{1:k-1}] - \mathbb{E}[S])|Y_{\perp}|, \qquad (E.42)$$

with $Y_{k-1} := \sum_{i \notin J, i \geq k+1} \varepsilon_i \mathbb{1}\{j(i) \leq k-1\}(|G_i| + d_i(Z_{j(i)}))$, which will imply the induction step (E.33) with $k \leftarrow k-1$ by taking expectations over Y_{\perp} . Since Y_{\perp} is considered constant in (E.42), we may assume without loss of generality that $Y_{\perp} = 0$, at the potential cost of replacing h by $h(\cdot + Y_{\perp})$, which is still convex. Therefore, it suffices to show

$$\mathbb{E}[h(Y_k + \mathbb{E}[S \mid Z_{1:k}] - \mathbb{E}[S]) \mid Y_{\parallel}] \tag{E.43}$$

$$\leq \mathbb{E}\left[h\left(Y_{k-1} + \mathbb{1}\left\{k \notin J\right\}\varepsilon_k(|G_k| + d_k(Z_{j(k)}))\right] \tag{E.44}$$

$$+ 1\{k \in J\}(\varepsilon_k c_k(Z_k, Z_k')$$
(E.45)

$$+ \sum_{\substack{i \notin J \\ i \ge k+1 \\ j(i)=k}} \operatorname{Sym}_{k}(\varepsilon_{i}(|G_{i}| + d_{i}(Z_{k})))) \mathbb{E}[S | Z_{1:k-1}] - \mathbb{E}[S])|Y_{\perp}], \qquad (E.46)$$

We first consider the case of $k \notin J$. Define $\Phi(z_{1:k}) := \mathbb{E}[S \mid Z_{1:k} = z_{1:k}]$. We rewrite the RHS of (E.46) as

$$\mathbb{E}[h(Y_k + \mathbb{E}[S \mid Z_{1:k}] - \mathbb{E}[S]) \mid Y_{\perp}] \tag{E.47}$$

$$= \mathbb{E} \left[h(Y_k + \Phi(Z_{1:k}) - \mathbb{E} \left[\Phi(Z_{1:k-1}, Z_k') \mid Z_{1:k-1} \right] + \mathbb{E} [S \mid Z_{1:k-1}] - \mathbb{E} [S]) | Y_{\perp} \right]$$
(E.48)

$$= \mathbb{E} \left[h(Y_k + \mathbb{E} \left[\Phi(Z_{1:k}) - \Phi(Z_{1:k-1}, Z_k') \, \middle| \, Z_{1:k} \right] + \mathbb{E} \left[S \, \middle| \, Z_{1:k-1} \right] - \mathbb{E} \left[S \right] \right) | Y_{\perp \perp} \right]$$
(E.49)

$$= \mathbb{E} \left[h(Y_k + \mathbb{E} \left[\Phi(Z_{1:k}) - \Phi(Z_{1:k-1}, Z_k') \mid Z_{1:k}, \mathcal{G} \right] + \mathbb{E} [S \mid Z_{1:k-1}] - \mathbb{E} [S] \right) | Y_{\perp \perp} \right]$$
(E.50)

(E.51)

where we used the fact that $\mathbb{E}[S \mid Z_{1:k-1}] = \mathbb{E}[\Phi(Z_{1:k-1}, Z_k') \mid Z_{1:k-1}] = \mathbb{E}[\Phi(Z_{1:k_1}, Z_k') \mid Z_{1:k}] = \mathbb{E}[\Phi(Z_{1:k-1}, Z_k') \mid Z_{1:k}] = \mathbb{E}[\Phi(Z_{1:k-1}, Z_k') \mid Z_{1:k}, \mathcal{G}],$ since $Z_k' \sim \text{Law}(Z_k \mid Z_{1:k-1})$ and $Z_k' \perp Z_k \mid Z_{1:k-1}$ and \mathcal{G} is independent of all Z, Z'. Since both Y_k and $\mathbb{E}[S \mid Z_{1:k-1}] - \mathbb{E}[S]$ are $\sigma(\mathcal{F}_k, \mathcal{G})$ -measurable, by Jensen's inequality (convexity of h) applied to the conditional expectation w.r.t. $Z_{1:k}, \mathcal{G}$, we have

$$\mathbb{E}[h(Y_k + \mathbb{E}[S \mid Z_{1:k}] - \mathbb{E}[S]) \mid Y_{\perp}] \tag{E.52}$$

$$\leq \mathbb{E} \left[h \left(Y_k + \Phi(Z_{1:k}) - \Phi(Z_{1:k-1}, Z_k') + \mathbb{E}[S \mid Z_{1:k-1}] - \mathbb{E}[S] \right) \mid Y_{\perp} \right]. \tag{E.53}$$

Since $k \notin J$, then Y_k is $\sigma(\mathcal{F}_{k-1}, \mathcal{G})$ -measurable. The following argument will now be made conditionally on $\mathcal{F}_{k-1}, \mathcal{G}, Y_{\perp \! \! \perp}$.

We have that $\Phi(Z_{1:k}) - \Phi(Z_{1:k-1}, Z'_k)$ is symmetric. Moreover, since $|\Phi(Z_{1:k}) - \Phi(Z_{1:k-1}, Z'_k)| \le c_k(Z_k, Z'_k) + c_k(Z_k, Z'_k)$ $d_k(Z_{i(k)})$ by assumption (i), we have that, for any t > 0,

$$\mathbb{P}(|\Phi(Z_{1:k}) - \Phi(Z_{1:k-1}, Z_k')| > t \mid \mathcal{F}_{k-1}, \mathcal{G}, Y_{\perp})$$
(E.54)

$$\leq \mathbb{P}\left(c_k(Z_k, Z_k') + d_k(Z_{j(k)}) > t \mid \mathcal{F}_{k-1}, \mathcal{G}, Y_{\perp}\right) \tag{E.55}$$

$$\leq \mathbb{P}\left(c_k(Z_k, Z_k') > t - d_k(Z_{j(k)}) \mid \mathcal{F}_{k-1}, \mathcal{G}, Y_{\perp}\right) \tag{E.56}$$

$$\leq 8 \mathbb{P}(|G_k| > t - d_k(Z_{i(k)}) \mid \mathcal{F}_{k-1}, \mathcal{G}, Y_{\perp}), \tag{E.57}$$

where we used that $\varepsilon_k c_k(Z_k, Z'_k)$ is σ_k^2 -sub-Gaussian conditionally on \mathcal{F}_{k-1} by assumption (ii) and Lemma E.3. Therefore, we can apply Lemma E.2 with $X \leftarrow \Phi(Z_{1:k}) - \Phi(Z_{1:k-1}, Z'_k)$ and $Z \leftarrow \varepsilon_k(|G_k| + d_k(Z_{j(k)}))$ with C = 8conditionally on $\mathcal{F}_{k-1}, Y_{\perp}$ to obtain

$$\mathbb{E}[h(Y_k + \mathbb{E}[S \mid Z_{1:k}] - \mathbb{E}[S]) \mid Y_{\perp}]$$
(E.58)

$$\leq \mathbb{E}\left[h(Y_k + \varepsilon_k(|G_k| + d_k(Z_{j(k)})) + \mathbb{E}[S \mid Z_{1:k-1}] - \mathbb{E}[S]) \mid Y_{\perp}\right],\tag{E.59}$$

which is (E.46) in the case $k \notin J$.

For the case $k \in J$, we use a similar argument. We now have, as before,

$$\mathbb{E}[S \mid Z_{1:k-1}] = \mathbb{E}[\Phi(Z_{1:k-1}, Z_k') \mid Z_{1:k-1}]$$
(E.60)

$$= \mathbb{E}[\Phi(Z_{1:k-1}, Z_k') + \sum_{\substack{i \notin J \\ i \ge k+1 \\ i(i) = k}} \varepsilon_i'(|G_i'| + d_i(Z_k)) | Z_{1:k-1}]$$
(E.61)

$$= \mathbb{E}[\Phi(Z_{1:k-1}, Z'_k) + \sum_{\substack{i \notin J \\ i \ge k+1 \\ j(i) = k}} \varepsilon'_i(|G'_i| + d_i(Z_k)) | Z_{1:k-1}]$$

$$= \mathbb{E}[\Phi(Z_{1:k-1}, Z'_k) + \sum_{\substack{i \notin J \\ i \ge k+1 \\ i(i) = k}} \varepsilon'_i(|G'_i| + d_i(Z_k)) | Z_{1:k}, \mathcal{G}],$$
(E.62)

by construction.

Since both Y_k and $\mathbb{E}[S \mid Z_{1:k-1}] - \mathbb{E}[S]$ are $\sigma(\mathcal{F}_k, \mathcal{G})$ -measurable, by Jensen's inequality (convexity of h) applied to the conditional expectation w.r.t. $Z_{1:k}$, \mathcal{G} , we have

$$\mathbb{E}[h(Y_k + \mathbb{E}[S \mid Z_{1:k}] - \mathbb{E}[S]) \mid Y_{\perp \perp}] \tag{E.63}$$

$$\leq \mathbb{E} \left[h \left(Y_{k} + \Phi(Z_{1:k}) - \Phi(Z_{1:k-1}, Z'_{k}) - \sum_{\substack{i \notin J \\ i \geq k+1 \\ i(i) = k}} \varepsilon'_{i} (|G'_{i}| + d_{i}(Z_{k})) + \mathbb{E}[S \mid Z_{1:k-1}] - \mathbb{E}[S] \right] \right| Y_{\perp} \right]. \tag{E.64}$$

We write Y_k as

$$Y_{k} = Y_{k-1} + \sum_{\substack{i \notin J \\ i \ge k+1 \\ i(i) = k}} \varepsilon_{i}(|G_{i}| + d_{i}(Z_{k})), \qquad (E.65)$$

where Y_{k-1} is $\sigma(\mathcal{F}_{k-1}, \mathcal{G})$ -measurable and obtain,

$$\mathbb{E}[h(Y_{k-1} + \mathbb{E}[S \mid Z_{1:k}] - \mathbb{E}[S]) \mid Y_{\perp}]$$

$$\leq \mathbb{E}\left[h\left(Y_{k-1} + \Phi(Z_{1:k}) - \Phi(Z_{1:k-1}, Z'_{k}) + \sum_{\substack{i \notin J \\ i \geq k+1 \\ j(i) = k}} \varepsilon_{i}(\mid G_{i}\mid + d_{i}(Z_{k})) - \varepsilon'_{i}(\mid G'_{i}\mid + d_{i}(Z_{k})) + \mathbb{E}[S \mid Z_{1:k-1}] - \mathbb{E}[S]\right) \middle| Y_{\perp}\right].$$
(E.66)

We now make the following domination argument conditionally on $\mathcal{F}_{k-1}, Y_{k-1}, Y_{\perp}$. The random variable

$$\Phi(Z_{1:k}) - \Phi(Z_{1:k-1}, Z'_k) + \sum_{\substack{i \notin J \\ i \ge k+1 \\ j(i) = k}} \varepsilon_i(|G_i| + d_i(Z_k)) - \varepsilon'_i(|G'_i| + d_i(Z_k))$$
(E.68)

is symmetric and, by assumption (i) and the triangle inequality, bounded in absolute value by

$$\left| \varepsilon_k c_k(Z_k, Z_k') + \sum_{\substack{i \notin J \\ i \ge k+1 \\ j(i) = k}} \operatorname{Sym}_k(\varepsilon_i(|G_i| + d_i(Z_k))) \right|. \tag{E.69}$$

Applying Lemma E.2 conditionally on \mathcal{F}_{k-1} , Y_{k-1} , Y_{\perp} with C=1 (hence no constant appears) yields the desired result.

We can now combine Lemma E.1 and Lemma E.4 to obtain the main moment bound of this section.

Theorem E.1 (Causal symmetrization). Let $m \in \mathbb{N}$ and $(\mathcal{Z}, \mathcal{A})$ be a standard Borel measurable space. Let Z_1, \ldots, Z_m be \mathbb{Z} -valued random with natural filtration $(\mathcal{F}_i)_{i=0,\ldots,m}$. Let $h \colon \mathbb{R} \to \mathbb{R}$ be convex.

Let $g: \mathbb{Z}^m \to \mathbb{R}$ be measurable and coordinate-wise Lipschitz with respect to cost functions $\rho_i: \mathbb{Z} \times \mathbb{Z} \to [0, \infty)$ such that $\rho_i(z_i, z_i) = 0$ with constants $L_i \geq 0$: for any $z, z' \in \mathbb{Z}^m$ differing only in the i-th coordinate,

$$|g(z) - g(z')| \le L_i \rho_i(z_i, z_i').$$
 (E.70)

Set $S := g(Z_1, \ldots, Z_m)$ and

For each $i \in \{1, \ldots, m\}$, assume there exists a conditionally independent resample

$$Z_i' \sim \text{Law}(Z_i \mid \mathcal{F}_{i-1}), \quad Z_i' \perp \!\!\! \perp Z_i \mid \mathcal{F}_{i-1}.$$
 (E.71)

Let $\varepsilon_{1:m}$, $\varepsilon'_{1:m}$ be independent Rademacher variables, independent of all Z, Z' and \mathcal{F}_m . Assume there exist constants $c_{ik} \geq 0$, measurable functions $d_{ik} : \mathcal{Z} \to [0, \infty)$ and $J \subset \{1, \ldots, m\}$ such that, the following conditions hold:

(i) For any i < k, there exists $j(i) \in J$, such that, for any $z_{1:i-1} \in \mathbb{Z}^{i-1}$ and $z_i, z_i' \in \mathbb{Z}$,

$$W_{\rho_k}\left(\text{Law}(Z_k \mid Z_{1:i} = z_{1:i}), \text{Law}(Z_k \mid Z_{1:i-1} = z_{1:i-1}, Z_i = z_i')\right) \le c_{ik}\rho_i(z_i, z_i') + d_{ik}(z_{j(i)}) \, \mathbb{1}\{i \notin J\} \,. \quad (E.72)$$

- (ii) For any $i \notin J$, $\varepsilon_i \rho_i(Z_i, Z_i')$ is σ_i^2 -sub-Gaussian conditionally on \mathcal{F}_{i-1} .
- (iii) For any $j \in J$, Z_i is independent of \mathcal{F}_{i-1} .

Then, there are Gaussian random variables $G_j, G'_i \sim \mathcal{N}(0, 8\sigma_i^2)$ independent and independent of all $Z, Z', \varepsilon, \mathcal{F}_m$ such that

$$\mathbb{E}[h(S - \mathbb{E}[S])] \tag{E.73}$$

$$\leq \mathbb{E}\left[h\left(\sum_{i \notin J} \operatorname{Sym}_{j(i)}\left(\varepsilon_{i}\left(L_{i}|G_{i}| + \sum_{k>i} L_{k}c_{ik}|G_{i}| + L_{k}d_{ik}(Z_{j(i)})\right)\right) + \sum_{j \in J} \varepsilon_{j}\left(L_{j}\rho_{j}(Z_{j}, Z'_{j}) + \sum_{k>j} L_{k}c_{jk}\rho_{j}(Z_{j}, Z'_{j})\right)\right],\tag{E.74}$$

where we use the notation:

$$\operatorname{Sym}_{j(i)} \left(\varepsilon_i \left(L_i | G_i| + \sum_{k>i} L_k c_{ik} | G_i| + L_k d_{ik} (Z_{j(i)}) \right) \right) := \tag{E.75}$$

$$\varepsilon_{j(i)} \left(\varepsilon_i \left(L_i | G_i| + \sum_{k>i} L_k c_{ik} | G_i| + L_k d_{ik}(Z_{j(i)}) \right) - \varepsilon_i' \left(L_i | G_i'| + \sum_{k>i} L_k c_{ik} | G_i'| + L_k d_{ik}(Z_{j(i)}) \right) \right). \tag{E.76}$$

E.2 Technical lemmas

We will make use of the following elementary lemma.

Lemma E.5. Let Z be a real-valued random variable. Assume there exist $c \ge 1$, $f, g: \mathbb{R} \to \mathbb{R}_+$ non-decreasing and $p \ge 2$ integer such that, for any integer $q \in [2, p]$,

$$\mathbb{E}[|Z|^q]^{1/q} \le f(q) + c^{1/q}g(q) \tag{E.77}$$

Then, for any $\delta \in (0, e^{-2}]$, with probability at least $1 - \delta$,

$$|Z| \le \begin{cases} ef(\log(1/\delta) + 1) + g(\log(1/\delta) + 1)e & \text{if } \delta \ge ce^{-p} \\ \frac{f(p) + c^{1/p}g(p)}{\delta^{1/p}} & \text{if } \delta < ce^{-p} \end{cases}.$$
 (E.78)

Proof. By Markov's inequality, for any integer $q \in [2, p]$,

$$\mathbb{P}(|Z| \ge t) \le \frac{\mathbb{E}[|Z|^q]}{t^q} \le \left(\frac{f(q) + c^{1/q}g(q)}{t}\right)^q. \tag{E.79}$$

Setting the right-hand side to δ and solving for t gives

$$t = \frac{f(q) + c^{1/q}g(q)}{\delta^{1/q}},$$
 (E.80)

If $\delta < ce^{-p}$, we can take q = p to obtain the second case of the result. If $\delta \ge ce^{-p}$, we take q the smallest integer such that $q \ge \log(c/\delta)$. Note that q is in [2, p] and $q \le \log(c/\delta) + 2$.

Since $c \ge 1$ and $\delta \le 1$, we have $\log(c/\delta) \ge 0$ and thus $\left(\frac{c}{\delta}\right)^{1/q} \le \left(\frac{c}{\delta}\right)^{1/\log(c/\delta)} = e$. Plugging this into (E.80) gives the bound in the first case.

We state the following lemma about sub-Gaussian random variables that will be useful later.

Lemma E.6. Let $X \in \mathbb{R}^m$ be a σ^2 -sub-Gaussian random variable, i.e., for any $\lambda > 0$,

$$\log \mathbb{E}[e^{\lambda \|X - \mathbb{E}[X]\|^2}] \le \frac{\sigma^2 \lambda^2}{2}. \tag{E.81}$$

Then, for X' an i.i.d. copy of X and ε a Rademacher random variable independent of X, X', the random variable $\varepsilon \|X - X'\|$ is sub-Gaussian with parameter at most $4\sigma^2$.

Proof. Since $Z := \varepsilon ||X - X'||$ is symmetric, it suffices to bound Z^2 as

$$Z^{2} = \|X - X'\|^{2} \le 2\|X - \mathbb{E}[X]\|^{2} + 2\|X' - \mathbb{E}[X]\|^{2}, \tag{E.82}$$

by Young's inequality. Using the independence of X and X' yields the result.

We will require the following chaining lemma for processes with L^p -Lipschitz increments. This result is a variant of the famous Dudley's entropy integral bound for sub-Gaussian processes, adapted to the L^p -Lipschitz setting.

This lemma is a direct consequence of the general chaining theory of Talagrand (2022) (see Talagrand (2022, Thm. B.2.3) with $\phi(x) = x^p$). Let us also mention Dirksen (2015) refined these ideas in the context of subpexponential processes while Latala and Tkocz (2015) further developed these tools for processes with heavier tails but still admitting a control over all moments. In our setting, the increments are assumed to be controlled only in L^p , which requires a different treatment of the maximal inequalities at each scale.

Lemma E.7 (Dudley-type entropy integral under L^p increments). Let $(X_t)_{t\in T}$ be a real-valued process indexed by a pseudometric space (T,d). Assume T is totally bounded with diameter $\Delta := \operatorname{diam}_d(T) \in (0,\infty)$ and that for some p > 1 and L > 0,

$$||X_t - X_s||_p \le L d(t, s) \qquad \forall s, t \in T. \tag{E.83}$$

Then

$$\mathbb{E}\left[\sup_{s,t\in T} (X_t - X_s)\right] \leq C L \int_0^{\Delta} (\mathcal{N}(T,d,\varepsilon))^{1/p} d\varepsilon, \tag{E.84}$$

where $\mathcal{N}(T,d,\varepsilon)$ is the ε -covering number and $C<\infty$ is an absolute constant.

E.3 Concentration bounds for ICL

We now apply the moment symmetrization results to derive concentration bounds for ICL in the dependent data setting. These concentration bounds will then be translated into generalization bounds in the next subsection. Let us recall ICL notations.

We denote by $\Theta \subset \mathbb{R}^d$ the space of tasks θ and by $\pi(\theta)$ the density of the pretraining task distribution. Given a task θ , the data is generated according to a task-specific distribution with density $p(\cdot | \theta)$. The training data is then generated by first sampling a task θ from the task distribution π , and then sampling data points $(x_t)_{t\geq 1}$ according to

$$x_{t+1} \sim p_{t+1}(\cdot | x_{1:t}, \theta)$$
 (E.85)

where $x_{1:t} = (x_1, ..., x_t)$.

Given a dataset of tasks $\theta_1, \ldots, \theta_N$ and associated samples $x_{1:T}^{(1)}, \ldots, x_{1:T}^{(N)}$, a model f is trained by minimizing the next-sample prediction loss

$$\widehat{L}(f, (\theta_n, x_{1:T}^n)_{n \le N}) = \frac{1}{NT} \sum_{n=1}^N \sum_{t=1}^T \ell_t(f(x_{1:t-1}^n), x_t^n),$$
(E.86)

where $\ell_t : \mathcal{X} \times \mathcal{X} \to [0, +\infty)$ is a loss function at step t.

We now provide a detailed version of Assumption 2.

Assumption 5 (Weak dependence). We assume that there are deterministic coefficients $(A_t)_{t\geq 1}$ and $(B_{s,t})_{t\geq s\geq 1}$ such that, for any $t\geq s\geq 1$, $\theta,\theta'\in\Theta$, any $x_{1:(s-1)}\in\mathcal{X}^{s-1}$, and any $x_t,x_t'\in\mathcal{X}$,

$$W_1(p_t(dx_t \mid \theta), p_t(dx_t' \mid \theta')) \le A_t \|\theta - \theta'\| \tag{E.87}$$

$$W_1(p_t(dx_t \mid x_{1:s}, \theta), p_t(dx_t' \mid x_{1:(s-1)}, x_s', \theta)) \le B_{s,t} \|\theta\|.$$
(E.88)

In the second assumption, the Wasserstein distance between the conditional distributions of x_t given x_s and x_s' is assumed to be controlled by the norm of the task θ . This is a slight difference with Assumption 2 where we assumed a dependence on $1 + \|\theta\|$. This is however without loss of generality as we can always consider $\tilde{\theta} = (1, \theta) \in \mathbb{R}^{d+1}$ and redefine the task distribution accordingly and this cosmetic change simplifies the presentation. We could also consider a dependence on $\|x_s - x_s'\|$, see Theorem E.1, but we omit this for simplicity.

Assumption 6 (Finite moments of the task distribution). There exists $q \ge 2$ integer such that $\mathbb{E}[\|\theta\|^q] < +\infty$.

Our theory could be extended to more general assumptions on the distributions of sample, but, for simplicity, we will make the following sub-Gaussian assumption on the data, conditionally on the past data and the task. Hence, this assumption does not restrict the task distribution in any way.

Assumption 7 (Sub-Gaussian data). There exists $\sigma > 0$ such that, for any $t \ge 1$, $\theta \in \Theta$, and any $x_{1:(t-1)} \in \mathcal{X}^{t-1}$, $x_t \sim \mathrm{p}_t(\cdot \mid x_{1:(t-1)}, \theta)$ is σ^2 -sub-Gaussian, i.e.,, for any $\lambda > 0$,

$$\log \mathbb{E}_{x_t \sim p_t(\cdot | x_{1:(t-1)}, \theta)} \left[e^{\lambda \|x_t - \mathbb{E}[x_t]\|^2} \right] \le \frac{\sigma^2 \lambda^2}{2} . \tag{E.89}$$

Assumption 8 (Lipschitz model and loss). The models $f \in \mathcal{F}$ are uniformly Lipschitz in the following sense: there exists $L_T > 0$ such that, for any $f \in \mathcal{F}$, any $x_{1:T}, x'_t$,

$$\frac{1}{T} \sum_{s=1}^{T} \|f(x_{1:s-1}) - f(x_{1:t-1}, x'_t, x_{t+1:s-1})\| \le L_T \|x_t - x'_t\|, \tag{E.90}$$

The losses ℓ_t are uniformly 1-Lipschitz: for any $t \ge 1$, any $x, x' \in \mathcal{X}$,

$$|\ell_t(x, x') - \ell_t(x, x')| \le ||x - x'||$$
 (E.91)

We will consider the following assumption on the function class \mathcal{F} .

Assumption 9. Assume that the hypothesis class \mathcal{F} is bounded for w.r.t. some distance dist on \mathcal{F} and that, the following extended Lipschitz condition holds: for any $f, f' \in \mathcal{F}$, any $x_{1:T}$, any $t \geq 1$, any x'_t , for any $f \in \mathcal{F}$, any $x_{1:T}, x'_t$,

$$\frac{1}{T} \sum_{s=1}^{T} \| f(x_{1:s-1}) - f(x_{1:t-1}, x'_t, x_{t+1:s-1}) - \left(f'(x_{1:s-1}) - f'(x_{1:t-1}, x'_t, x_{t+1:s-1}) \right) \|$$
 (E.92)

$$\leq M_T \|x_t - x_t'\| \operatorname{dist}(f, f'). \tag{E.93}$$

Note that Assumption 8 is implied of Assumption 9 when the constant function equal to zero is in \mathcal{F} with $L_T = M_T \sup_{f \in \mathcal{F}} \operatorname{dist}(f, 0)$.

We denote by $||X||_h$ the L^h norm of a random variable X, i.e., $||X||_h = (\mathbb{E}[||X||^h])^{1/h}$.

Lemma E.8. For any $r \in [2, q]$ integer, under Assumptions 5–8, we have

$$\left\| \sup_{f \in \mathcal{F}} \left\{ \mathbb{E} \left[\widehat{L}(f, (\theta_n, x_{1:T}^n)_{n \le N}) \right] - \widehat{L}(f, (\theta_n, x_{1:T}^n)_{n \le N}) \right\}$$
(E.94)

$$-\mathbb{E}\left[\sup_{f\in\mathcal{F}}\left\{\mathbb{E}\left[\widehat{L}(f,(\theta_n,x_{1:T}^n)_{n\leq N})\right]-\widehat{L}(f,(\theta_n,x_{1:T}^n)_{n\leq N})\right\}\right]_r$$
(E.95)

$$\leq c\sigma L_T \sqrt{\frac{Tr}{N}}$$
(E.96)

$$+ c\sqrt{r} \frac{L_T}{\sqrt{N}} \sqrt{\sum_{t=1}^{T} \left(\sum_{s>t} B_{t,s}\right)^2} \|\theta_1\|_2 + cr^{3/2} \frac{L_T}{N^{1-1/r}} \sqrt{\sum_{t=1}^{T} \left(\sum_{s>t} B_{t,s}\right)^2} \|\theta_1\|_q$$
 (E.97)

$$+ c\sqrt{r} \frac{L_T}{\sqrt{N}} \left(\sum_{t=1}^T A_t \right) \|\theta_1 - \mathbb{E}[\theta_1]\|_2 + cr \frac{L_T}{N^{1-1/r}} \left(\sum_{t=1}^T A_t \right) \|\theta_1 - \mathbb{E}[\theta_1]\|_q,$$
 (E.98)

where c > 0 is a universal constant.

Proof. We apply Theorem E.1 with

$$(Z_1, \dots, Z_m) = (\theta_1, x_1^{(1)}, \dots, x_T^{(1)}, \dots, \theta_N, x_1^{(N)}, \dots, x_T^{(N)}),$$
(E.99)

and

$$g(\theta_1, x_{1:T}^{(1)}, \dots, \theta_N, x_{1:T}^{(N)})$$
 (E.100)

$$= \sup_{f \in \mathcal{F}} \left\{ \mathbb{E} \left[\widehat{L}(f, (\theta_n, x_{1:T}^n)_{n \le N}) \right] - \widehat{L}(f, (\theta_n, x_{1:T}^n)_{n \le N}) \right\}$$
(E.101)

$$= \sup_{f \in \mathcal{F}} \frac{1}{NT} \left\{ \mathbb{E} \left[\sum_{n=1}^{N} \sum_{t=1}^{T} \ell_{t}(f(x_{1:t-1}^{n}), x_{t}^{n}) \right] - \sum_{n=1}^{N} \sum_{t=1}^{T} \ell_{t}(f(x_{1:t-1}^{n}), x_{t}^{n}) \right\}.$$
 (E.102)

By Assumption 8, f is coordinate-wise Lipschitz with respect to x_t^n with constant $L_{N,T} := L_T/N$ and formally constant with respect to θ_n .

By Lemma E.6 and Assumption 7, $\varepsilon_t^n \|x_t^n - x_t'^n\|$ is $4\sigma^2$ -sub-Gaussian conditionally on $x_{1:(t-1)}$, θ_n , for ε_t^n a Rademacher variable independent of all data.

We now apply Theorem E.1 with $h(x) = |x|^r$ for r integer such that $2 \le r \le q$ and J corresponding to the indices of the tasks $\theta_1, \ldots, \theta_N$. We obtain that

$$\|f - \mathbb{E}[f]\|_r \tag{E.103}$$

$$\leq \left\| \sum_{n=1}^{N} \sum_{t=1}^{T} \operatorname{Sym}_{n} \left(\varepsilon_{t}^{n} \left(L_{N,T} |G_{t}^{n}| + \sum_{s>t} L_{N,T} B_{t,s} \|\theta_{n}\| \right) \right) + \sum_{n=1}^{N} \sum_{t=1}^{T} L_{N,T} \varepsilon_{n} A_{t} \|\theta_{n} - \theta_{n}'\| \right\|_{r},$$
(E.104)

where

$$\operatorname{Sym}_{n}\left(\varepsilon_{t}^{n}\left(L_{N,T}|G_{t}^{n}|+\sum_{s>t}L_{N,T}B_{t,s}\|\theta_{n}\|\right)\right) := \tag{E.105}$$

$$\varepsilon_{n}\left(\varepsilon_{t}^{n}\left(L_{N,T}|G_{t}^{n}|+\sum_{s>t}L_{N,T}B_{t,s}\|\theta_{n}\|\right)-\varepsilon_{t}^{n\prime}\left(L_{N,T}|G_{t}^{n\prime}|+\sum_{s>t}L_{N,T}B_{t,s}\|\theta_{n}\|\right)\right),\tag{E.106}$$

and $G_t^n, G_t'^n \sim \mathcal{N}(0, 32\sigma^2)$ independent of all data and Rademacher variables.

Using Minkowski's inequality, we have

$$\|f - \mathbb{E}[f]\|_r \tag{E.107}$$

$$\leq \left\| \sum_{n=1}^{N} \varepsilon_n \sum_{t=1}^{T} L_{N,T}(\varepsilon_t^{\ n} | G_t^{\ n} | - \varepsilon_t^{\ n'} | G_t^{\ n'} |) \right\|_r \tag{E.108}$$

$$+ \left\| \sum_{n=1}^{N} \varepsilon_{n} \left(\|\theta_{n}\| \sum_{t=1}^{T} L_{N,T} \sum_{s>t} B_{t,s} \varepsilon_{t}^{n} - \|\theta_{n}'\| \sum_{t=1}^{T} L_{N,T} \sum_{s>t} B_{t,s} \varepsilon_{t}^{n\prime} \right) \right\|_{L^{\infty}}$$
(E.109)

$$+ \left\| \sum_{n=1}^{N} \varepsilon_{n} \|\theta_{n} - \theta_{n}'\| \sum_{t=1}^{T} L_{N,T} A_{t} \right\|_{r}. \tag{E.110}$$

We now bound each term (E.108)–(E.110) separately.

We begin with (E.108). By independence of the Rademacher variables and the Gaussian variables, we have that (E.108) can be rewritten as

$$(E.108) = \sqrt{2}L_{N,T} \left\| \sum_{n=1}^{N} \sum_{t=1}^{T} G_t^n \right\|_{T}$$
(E.111)

$$=8\sigma L_{N,T}\sqrt{NT}\|G\|_{r}, \qquad (E.112)$$

where $G \sim \mathcal{N}(0,1)$. Using standard bounds on subGaussian random variables, we have that $||G||_r \leq c\sqrt{r}$ for some universal constant c > 0 (see e.g. Vershynin (2018, Chap. 2)). Hence, we have

$$(E.108) \le c\sigma L_{N,T} \sqrt{NTr}, \qquad (E.113)$$

for some universal constant c > 0.

We now turn to (E.109). By Boucheron et al. (2005, Thm. 15.11), applied to each independent and zero-mean term

$$\varepsilon_n \left(\|\theta_n\| \sum_{t=1}^T \varepsilon_t^n \sum_{s>t} B_{t,s} - \|\theta_n'\| \sum_{t=1}^T \varepsilon_t^{n'} \sum_{s>t} B_{t,s} \right), \tag{E.114}$$

we have

$$(E.109) \le c\sqrt{r}L_{N,T}\sqrt{N} \left\| \|\theta_1\| \sum_{t=1}^{T} \varepsilon_t^{\ 1} \sum_{s>t} B_{t,s} - \|\theta_1'\| \sum_{t=1}^{T} \varepsilon_t^{\ 1'} \sum_{s>t} B_{t,s} \right\|_2$$
(E.115)

+
$$crL_{N,T}N^{1/r} \left\| \|\theta_1\| \sum_{t=1}^{T} \varepsilon_t^{-1} \sum_{s>t} B_{t,s} - \|\theta_1'\| \sum_{t=1}^{T} \varepsilon_t^{-1'} \sum_{s>t} B_{t,s} \right\|_r$$
, (E.116)

where c > 0 is a universal constant.

Using Minkowski's inequality again, we have

$$(E.109) \le c\sqrt{r}L_{N,T}\sqrt{N} \left\| \|\theta_1\| \sum_{t=1}^{T} \varepsilon_t^{-1} \sum_{s>t} B_{t,s} \right\|_{2}$$
(E.117)

+
$$crL_{N,T}N^{1/r} \left\| \|\theta_1\| \sum_{t=1}^{T} \varepsilon_t^{-1} \sum_{s>t} B_{t,s} \right\|_{r}$$
 (E.118)

$$\leq c\sqrt{r}L_{N,T}\sqrt{N}\|\theta_1\|_2 \left\| \sum_{t=1}^T \varepsilon_t^{-1} \sum_{s>t} B_{t,s} \right\|_2 \tag{E.119}$$

+
$$crL_{N,T}N^{1/r}\|\theta_1\|_r \left\| \sum_{t=1}^T \varepsilon_t^{-1} \sum_{s>t} B_{t,s} \right\|_r$$
, (E.120)

where we used that θ_1 and $(\varepsilon_t^{\ 1})_{t\geq 1}$ are independent. Now, $\sum_{t=1}^T \varepsilon_t^{\ 1} \sum_{s>t} B_{t,s}$ is a zero-mean sub-Gaussian random variable with parameter $\sum_{t=1}^T \left(\sum_{s>t} B_{t,s}\right)^2$ by Hoeffding's lemma (see e.g. Wainwright (2019, Exercise 2.4)) and we have, for some universal constant c>0, for any integer h

$$\left\| \sum_{t=1}^{T} \varepsilon_{t}^{1} \sum_{s>t} B_{t,s} \right\|_{h} \le c\sqrt{h} \left(\sum_{t=1}^{T} \left(\sum_{s>t} B_{t,s} \right)^{2} \right)^{1/2}. \tag{E.121}$$

Plugging this into (E.120) with h = 2 and h = r gives

$$(E.109) \le c\sqrt{r}L_{N,T}\sqrt{N}\sqrt{\sum_{t=1}^{T}\left(\sum_{s>t}B_{t,s}\right)^{2}}\|\theta_{1}\|_{2} + cr^{3/2}L_{N,T}N^{1/r}\sqrt{\sum_{t=1}^{T}\left(\sum_{s>t}B_{t,s}\right)^{2}}\|\theta_{1}\|_{r}$$
 (E.122)

$$\leq c\sqrt{r}L_{N,T}\sqrt{N}\sqrt{\sum_{t=1}^{T}\left(\sum_{s>t}B_{t,s}\right)^{2}}\|\theta_{1}\|_{2} + cr^{3/2}L_{N,T}N^{1/r}\sqrt{\sum_{t=1}^{T}\left(\sum_{s>t}B_{t,s}\right)^{2}}\|\theta_{1}\|_{q}$$
 (E.123)

(E.124)

where we used that $r \leq q$ to obtain the last inequality.

Finally, we proceed similarly for (E.110). By Boucheron et al. (2005, Thm. 15.11) applied to each independent and zero-mean term

$$\varepsilon_n \|\theta_n - \theta_n'\| \sum_{t=1}^T L_{N,T} A_t , \qquad (E.125)$$

we have

$$(E.110) \le c\sqrt{r}L_{N,T}\sqrt{N}\left(\sum_{t=1}^{T}A_{t}\right)\left\|\theta_{1}-\theta_{1}'\right\|_{2} + crL_{N,T}N^{1/r}\left(\sum_{t=1}^{T}A_{t}\right)\left\|\theta_{1}-\theta_{1}'\right\|_{r}$$
(E.126)

$$\leq c\sqrt{r}L_{N,T}\sqrt{N}\left(\sum_{t=1}^{T}A_{t}\right)\|\theta_{1} - \mathbb{E}[\theta_{1}]\|_{2} + crL_{N,T}N^{1/r}\left(\sum_{t=1}^{T}A_{t}\right)\|\theta_{1} - \mathbb{E}[\theta_{1}]\|_{q}, \tag{E.127}$$

where we use Minkowski's inequality and the fact that $r \leq q$ to obtain the last inequality.

Combining (E.113), (E.124), and (E.127) and replacing $L_{N,T}$ by L_T/N gives the result.

Proposition E.1 (Concentration bound for ICL). Under Assumptions 5–8, for any $\delta \in (0, e^{-2}]$, with probability at least $1 - \delta$,

$$\left| \sup_{f \in \mathcal{F}} \left\{ \mathbb{E} \left[\widehat{L}(f, (\theta_n, x_{1:T}^n)_{n \le N}) \right] - \widehat{L}(f, (\theta_n, x_{1:T}^n)_{n \le N}) \right\} - \mathbb{E} \left[\sup_{f \in \mathcal{F}} \left\{ \mathbb{E} \left[\widehat{L}(f, (\theta_n, x_{1:T}^n)_{n \le N}) \right] - \widehat{L}(f, (\theta_n, x_{1:T}^n)_{n \le N}) \right\} \right] \right|$$
(E.128)

is bounded by

(a) If $\delta \geq Ne^{-q}$,

$$c\sigma \frac{L_T}{\sqrt{N}} \sqrt{T(\log(N/\delta) + 1)} \tag{E.129}$$

$$+ c\sqrt{(\log(N/\delta) + 1)} \frac{L_T}{\sqrt{N}} \sqrt{\sum_{t=1}^T \left(\sum_{s>t} B_{t,s}\right)^2} \|\theta_1\|_2 + c(\log(N/\delta) + 1)^{3/2} \frac{L_T}{N} \sqrt{\sum_{t=1}^T \left(\sum_{s>t} B_{t,s}\right)^2} \|\theta_1\|_q \qquad (E.130)$$

$$+ c\sqrt{(\log(N/\delta) + 1)} \frac{L_T}{\sqrt{N}} \left(\sum_{t=1}^T A_t \right) \|\theta_1 - \mathbb{E}[\theta_1]\|_2 + c(\log(N/\delta) + 1) \frac{L_T}{N} \left(\sum_{t=1}^T A_t \right) \|\theta_1 - \mathbb{E}[\theta_1]\|_q$$
 (E.131)

(b) If $\delta < Ne^{-q}$,

$$\frac{1}{\delta^{1/q}} \left(c\sigma L_{N,T} \sqrt{\frac{Tq}{N}} \right) \tag{E.132}$$

$$+ c\sqrt{q} \frac{L_T}{\sqrt{N}} \sqrt{\sum_{t=1}^{T} \left(\sum_{s>t} B_{t,s}\right)^2} \|\theta_1\|_2 + cq^{3/2} \frac{L_T}{N^{1-1/q}} \sqrt{\sum_{t=1}^{T} \left(\sum_{s>t} B_{t,s}\right)^2} \|\theta_1\|_q$$
 (E.133)

$$+ c\sqrt{q} \frac{L_T}{\sqrt{N}} \left(\sum_{t=1}^T A_t \right) \|\theta_1 - \mathbb{E}[\theta_1]\|_2 + cq \frac{L_T}{N^{1-1/q}} \left(\sum_{t=1}^T A_t \right) \|\theta_1 - \mathbb{E}[\theta_1]\|_q \right)$$
 (E.134)

Proof. We apply Lemma E.5 to the moment bound from Lemma E.8.

For Lemma E.5, we use:

$$f(r) = c\sigma L_T \sqrt{\frac{Tr}{T}} + c\sqrt{r} \frac{L_T}{\sqrt{N}} \sqrt{\sum_{t=1}^{T} \left(\sum_{s>t} B_{t,s}\right)^2} \|\theta_1\|_2 + c\sqrt{r} \frac{L_T}{\sqrt{N}} \left(\sum_{t=1}^{T} A_t\right) \|\theta_1 - \mathbb{E}[\theta_1]\|_2$$
 (E.135)

$$g(r) = cr^{3/2} \frac{L_T}{N^{1-1/r}} \sqrt{\sum_{t=1}^T \left(\sum_{s>t} B_{t,s}\right)^2} \|\theta_1\|_q + cr \frac{L_T}{N^{1-1/r}} \left(\sum_{t=1}^T A_t\right) \|\theta_1 - \mathbb{E}[\theta_1]\|_q.$$
 (E.136)

Applying Lemma E.5 then gives the desired concentration bound.

E.4 Complexity bounds for ICL

We now derive bounds for the analogue of the Rademacher complexity term in our setting. We will again rely on Theorem E.1.

Lemma E.9. Under Assumptions 5–9, we have

$$\mathbb{E}\left[\sup_{f\in\mathcal{F}}\mathbb{E}\left[\widehat{L}(f,(\theta_n,x_{1:T}^n)_{n\leq N})\right] - \widehat{L}(f,(\theta_n,x_{1:T}^n)_{n\leq N})\right]$$
(E.137)

$$\leq c\mathcal{I}(\mathcal{F}, \operatorname{dist}, q) \left(\sigma M_T \sqrt{\frac{Tq}{N}} \right)$$
 (E.138)

$$+ c\sqrt{q} \frac{M_T}{\sqrt{N}} \sqrt{\sum_{t=1}^T \left(\sum_{s>t} B_{t,s}\right)^2} \|\theta_1\|_2 + q^{3/2} \frac{M_T}{N^{1-1/q}} \sqrt{\sum_{t=1}^T \left(\sum_{s>t} B_{t,s}\right)^2} \|\theta_1\|_q$$
 (E.139)

$$+ \sqrt{q} \frac{M_T}{\sqrt{N}} \left(\sum_{t=1}^T A_t \right) \|\theta_1 - \mathbb{E}[\theta_1]\|_2 + cq \frac{M_T}{N^{1-1/q}} \left(\sum_{t=1}^T A_t \right) \|\theta_1 - \mathbb{E}[\theta_1]\|_q \right), \tag{E.140}$$

where c > 0 is a universal constant and where the Dudley-type integral $\mathcal{I}_{dist}(\mathcal{F})$ is defined as

$$\mathcal{I}(\mathcal{F}, \operatorname{dist}, q) = \int_0^{\Delta} (\mathcal{N}(\mathcal{F}, \operatorname{dist}, u))^{1/q} du, \quad \text{with } \Delta = \operatorname{diam}_{\operatorname{dist}}(\mathcal{F}) = \sup_{f, f' \in \mathcal{F}} \operatorname{dist}(f, f').$$
 (E.141)

Proof. The main idea of the proof is to use Lemma E.7 and to rely on Theorem E.1 to control the moments of the increments of the process $\sup_{f \in \mathcal{F}} \widehat{L}(f, (\theta_n, x_{1:T}^n)_{n \leq N}) - \mathbb{E}\Big[\widehat{L}(f, (\theta_n, x_{1:T}^n)_{n \leq N})\Big]$. Fix $f, f' \in \mathcal{F}$. We apply Theorem E.1 with

$$(Z_1, \dots, Z_m) = (\theta_1, x_1^{(1)}, \dots, x_T^{(1)}, \dots, \theta_N, x_1^{(N)}, \dots, x_T^{(N)}),$$
 (E.142)

and

$$g(\theta_1, x_{1:T}^{(1)}, \dots, \theta_N, x_{1:T}^{(N)})$$
 (E.143)

$$= \mathbb{E}\left[\widehat{L}(f, (\theta_n, x_{1:T}^n)_{n \le N})\right] - \widehat{L}(f, (\theta_n, x_{1:T}^n)_{n \le N}) \tag{E.144}$$

$$-\left(\mathbb{E}\Big[\widehat{L}(f',(\theta_n,x_{1:T}^n)_{n\leq N})\Big] - \widehat{L}(f',(\theta_n,x_{1:T}^n)_{n\leq N})\right)$$
(E.145)

and proceed as in the proof of Lemma E.8 except that g is now $M_T \operatorname{dist}(f, f')$ coordinate-wise Lipschitz by Assumption 9 to obtain that:

$$\left\| \widehat{L}(f, (\theta_n, x_{1:T}^n)_{n \le N}) - \mathbb{E} \left[\widehat{L}(f, (\theta_n, x_{1:T}^n)_{n \le N}) \right] - \left(\widehat{L}(f', (\theta_n, x_{1:T}^n)_{n \le N}) - \mathbb{E} \left[\widehat{L}(f', (\theta_n, x_{1:T}^n)_{n \le N}) \right] \right) \right\|_{q}$$
(E.146)

$$\leq \operatorname{dist}(f, f') \left(c \sigma M_T \sqrt{\frac{Tq}{N}} \right)$$
 (E.147)

$$+ c\sqrt{q} \frac{M_T}{\sqrt{N}} \sqrt{\sum_{t=1}^T \left(\sum_{s>t} B_{t,s}\right)^2} \|\theta_1\|_2 + cq^{3/2} \frac{M_T}{N^{1-1/q}} \sqrt{\sum_{t=1}^T \left(\sum_{s>t} B_{t,s}\right)^2} \|\theta_1\|_q$$
 (E.148)

$$+ c\sqrt{q} \frac{M_T}{\sqrt{N}} \left(\sum_{t=1}^T A_t \right) \|\theta_1 - \mathbb{E}[\theta_1]\|_2 + cq \frac{M_T}{N^{1-1/q}} \left(\sum_{t=1}^T A_t \right) \|\theta_1 - \mathbb{E}[\theta_1]\|_q \right). \tag{E.149}$$

Applying Lemma E.7 then gives that

$$\mathbb{E}\left[\sup_{f,f'\in\mathcal{F}}\mathbb{E}\left[\widehat{L}(f,(\theta_n,x_{1:T}^n)_{n\leq N})\right] - \widehat{L}(f,(\theta_n,x_{1:T}^n)_{n\leq N}) - \left(\mathbb{E}\left[\widehat{L}(f',(\theta_n,x_{1:T}^n)_{n\leq N})\right] - \widehat{L}(f',(\theta_n,x_{1:T}^n)_{n\leq N})\right)\right]$$
(E.150)

is bounded by the RHS of the statement of the lemma. To conclude, it suffices to notice that, for any $f_0 \in \mathcal{F}$ fixed,

$$\mathbb{E}\left[\sup_{f\in\mathcal{F}}\mathbb{E}\left[\widehat{L}(f,(\theta_{n},x_{1:T}^{n})_{n\leq N})\right] - \widehat{L}(f,(\theta_{n},x_{1:T}^{n})_{n\leq N})\right]$$

$$= \mathbb{E}\left[\sup_{f\in\mathcal{F}}\mathbb{E}\left[\widehat{L}(f,(\theta_{n},x_{1:T}^{n})_{n\leq N})\right] - \widehat{L}(f,(\theta_{n},x_{1:T}^{n})_{n\leq N}) - \left(\mathbb{E}\left[\widehat{L}(f_{0},(\theta_{n},x_{1:T}^{n})_{n\leq N})\right] - \widehat{L}(f_{0},(\theta_{n},x_{1:T}^{n})_{n\leq N})\right)\right]$$
(E.151)
$$(E.152)$$

$$\leq \mathbb{E}\left[\sup_{f,f'\in\mathcal{F}}\mathbb{E}\left[\widehat{L}(f,(\theta_n,x_{1:T}^n)_{n\leq N})\right] - \widehat{L}(f,(\theta_n,x_{1:T}^n)_{n\leq N}) - \left(\mathbb{E}\left[\widehat{L}(f',(\theta_n,x_{1:T}^n)_{n\leq N})\right] - \widehat{L}(f',(\theta_n,x_{1:T}^n)_{n\leq N})\right)\right],\tag{E.153}$$

which concludes the proof.

E.5 Generalization bounds for ICL

Putting together the concentration bound from Proposition E.1 and the complexity bound from Lemma E.9, we obtain the following generalization bound for ICL:

Theorem E.2 (Generalization bound for ICL). Under Assumptions 5–9, for any $\delta \in (0, e^{-2}]$, for any $\delta \in (0, Ne^{-q}]$, with probability at least $1 - \delta$, the generalization gap

$$\sup_{f \in \mathcal{F}} \mathbb{E}\Big[\widehat{L}(f, (\theta_n, x_{1:T}^n)_{n \le N})\Big] - \widehat{L}(f, (\theta_n, x_{1:T}^n)_{n \le N})$$
(E.154)

is bounded by

(a) If $\delta \geq Ne^{-q}$,

$$c\sigma\sqrt{\frac{T}{N}}\left(L_T\sqrt{(\log(N/\delta)+1)} + M_T\mathcal{I}(\mathcal{F}, \operatorname{dist}, q)\sqrt{q}\right)$$
(E.155)

$$+ c \left(L_T \sqrt{(\log(N/\delta) + 1)} + M_T \mathcal{I}(\mathcal{F}, \operatorname{dist}, q) \sqrt{q} \right) \frac{1}{\sqrt{N}} \sqrt{\sum_{t=1}^T \left(\sum_{s>t} B_{t,s} \right)^2} \|\theta_1\|_2$$
 (E.156)

$$+ c \left((\log(N/\delta) + 1)^{3/2} L_T + q^{3/2} N^{1/q} M_T \mathcal{I}(\mathcal{F}, \text{dist}, q) \right) \frac{1}{N} \sqrt{\sum_{t=1}^T \left(\sum_{s>t} B_{t,s} \right)^2} \|\theta_1\|_q$$
 (E.157)

$$+ c \left(L_T \sqrt{(\log(N/\delta) + 1)} + M_T \mathcal{I}(\mathcal{F}, \operatorname{dist}, q) \sqrt{q} \right) \frac{1}{\sqrt{N}} \left(\sum_{t=1}^T A_t \right) \|\theta_1 - \mathbb{E}[\theta_1]\|_2$$
 (E.158)

$$+ c \left((\log(N/\delta) + 1) L_T + q N^{1/q} M_T \mathcal{I}(\mathcal{F}, \operatorname{dist}, q) \right) \frac{1}{N} \left(\sum_{t=1}^T A_t \right) \|\theta_1 - \mathbb{E}[\theta_1]\|_q$$
 (E.159)

(b) If $\delta < Ne^{-q}$,

$$\left(\frac{L_T}{\delta^{1/q}} + M_T \mathcal{I}(\mathcal{F}, \text{dist}, q)\right) \left(c\sigma \sqrt{\frac{Tq}{N}}\right)$$
(E.160)

$$+ c\sqrt{q} \frac{L_T}{\sqrt{N}} \sqrt{\sum_{t=1}^{T} \left(\sum_{s>t} B_{t,s}\right)^2} \|\theta_1\|_2 + cq^{3/2} \frac{L_T}{N^{1-1/q}} \sqrt{\sum_{t=1}^{T} \left(\sum_{s>t} B_{t,s}\right)^2} \|\theta_1\|_q$$
 (E.161)

$$+ c\sqrt{q} \frac{L_T}{\sqrt{N}} \left(\sum_{t=1}^T A_t \right) \|\theta_1 - \mathbb{E}[\theta_1]\|_2 + cq \frac{L_T}{N^{1-1/q}} \left(\sum_{t=1}^T A_t \right) \|\theta_1 - \mathbb{E}[\theta_1]\|_q \right), \tag{E.162}$$

where c > 0 is a universal constant and where the Dudley-type integral $\mathcal{I}_{dist}(\mathcal{F})$ is defined as

$$\mathcal{I}(\mathcal{F}, \operatorname{dist}, q) = \int_0^{\Delta} (\mathcal{N}(\mathcal{F}, \operatorname{dist}, u))^{1/q} du, \quad \text{with } \Delta = \operatorname{diam}_{\operatorname{dist}}(\mathcal{F}) = \sup_{f, f' \in \mathcal{F}} \operatorname{dist}(f, f').$$
 (E.163)

Proof. The result is obtained by combining Proposition E.1 and Lemma E.9: we write the decomposition

$$\sup_{f \in \mathcal{F}} \left\{ \mathbb{E} \left[\widehat{L}(f, (\theta_n, x_{1:T}^n)_{n \le N}) \right] - \widehat{L}(f, (\theta_n, x_{1:T}^n)_{n \le N}) \right\}$$
(E.164)

$$= \mathbb{E}\left[\sup_{f \in \mathcal{F}} \left\{ \mathbb{E}\left[\widehat{L}(f, (\theta_n, x_{1:T}^n)_{n \le N})\right] - \widehat{L}(f, (\theta_n, x_{1:T}^n)_{n \le N})\right\}\right]$$
(E.165)

$$+ \sup_{f \in \mathcal{F}} \left\{ \mathbb{E} \left[\widehat{L}(f, (\theta_n, x_{1:T}^n)_{n \leq N}) \right] - \widehat{L}(f, (\theta_n, x_{1:T}^n)_{n \leq N}) \right\} - \mathbb{E} \left[\sup_{f \in \mathcal{F}} \left\{ \mathbb{E} \left[\widehat{L}(f, (\theta_n, x_{1:T}^n)_{n \leq N}) \right] - \widehat{L}(f, (\theta_n, x_{1:T}^n)_{n \leq N}) \right\} \right], \tag{E.166}$$

and we bound (E.165) using Lemma E.9 and (E.166) with high probability using Proposition E.1.

F Additional details on examples

F.1 Example: Volterra equation model

We discuss the Volterra equation model to explicit the dependence of the generalization bounds on the memory decay parameter $\alpha > 0$.

Setup. Let $(W_t)_{t\geq 1}$ be noise sequence taking values in \mathbb{R}^d . Given a Lipschitz drift b, $mathbb{R}^d \to \mathbb{R}^d$ with Lipschitz constant $L\geq 0$, we consider the discretized Volterra equation: for $t\geq 0$,

$$X_{t+1} = \sum_{u=1}^{t} K(t, u) \left(b(X_u) + W_u \right), \qquad K(t, u) = \frac{1}{(t - u + 1)^{\alpha}}, \quad \alpha > 0.$$
 (F.1)

When applying the generalization framework, we would consider the augmented sequence $(X_1, W_1, X_2, W_2, \ldots)$. To satisfy the weak dependence assumption Assumption 5, we need to bound the effect of perturbations in either the state or the noise or the drift. We begin with perturbations in the state or noise, and we discuss drift perturbations at the end of this section. For perturbations in the state or noise, we will obtain bounds on the Wasserstein distance between the conditional laws of X_t and X_t' given the past, where X_t and X_t' are two versions of the process (F.1) that differ by a perturbation at some time s < t.

The coefficient α will play a key role in the dependence structure through the sums:

$$H_{\alpha}(n) = \sum_{r=1}^{n} \frac{1}{r^{\alpha}}.$$
 (F.2)

We also use $\zeta(\alpha) = \sum_{r=1}^{\infty} r^{-\alpha}$ for $\alpha > 1$ and we have the following bounds on $H_{\alpha}(n)$

$$H_{\alpha}(n) \leq \begin{cases} 1 + \log n, & \alpha = 1, \\ \zeta(\alpha), & \alpha > 1. \end{cases}$$
 (F.3)

We will make use of the following technical lemma.

Lemma F.1. Let $(a_n)_{n\geq 0}$ be nonnegative numbers and suppose that for $n\geq 1$,

$$a_n \le L \sum_{r=1}^n r^{-\alpha} a_{n-r} + g_n,$$
 (F.4)

with non-decreasing $(g_n)_{n\geq 1}$ and given $a_0\geq 0$. Define, for $N\geq 1$,

$$\lambda_N := \begin{cases} L(1 + \log N) & \text{if } \alpha = 1, \\ L\zeta(\alpha) & \text{if } \alpha > 1. \end{cases}$$
 (F.5)

Then, for all $1 \le n \le N$,

$$a_n \le \lambda_N^n a_0 + \sum_{j=1}^n g_j \lambda_N^{n-j}. \tag{F.6}$$

Proof. Let $A_n := \max_{0 \le m \le n} a_m$. From (F.4), $a_n \le L \sum_{r=1}^n r^{-\alpha} A_{n-r} + g_n \le L H_{\alpha}(n) A_{n-1} + g_n$, so $A_n \le L H_{\alpha}(n) A_{n-1} + g_n$ since $(g_n)_n$ is non-decreasing. Bounding $H_{\alpha}(n)$ using (F.3) gives $A_n \le \lambda_N A_{n-1} + g_n$ for all $1 \le n \le N$. Iterating this inequality yields the result.

State perturbation. Fix $s \ge 1$ and let $\mathcal{F}_s := \sigma(X_1, \dots, X_s, W_1, \dots, W_s)$ on which we condition. Assume the two systems agree up to s-1, and at time s we have

$$X_s' = X_s - h$$

with $h \neq 0$. For $t \geq s$, define $\Delta_t := X_t - X_t'$. Subtracting (F.1) for the two evolutions (they share (W_u)) gives for $t \geq s$:

$$\Delta_{t+1} = \sum_{u=s}^{t} \frac{b(X_u) - b(X_u')}{(t-u+1)^{\alpha}}, \qquad \|\Delta_{t+1}\| \le L \sum_{u=s}^{t} \frac{\|\Delta_u\|}{(t-u+1)^{\alpha}}.$$
 (F.7)

Set n := t - s + 1, $a_n := \mathbb{E}(\|\Delta_{s+n}\| \mid \mathcal{F}_s)$ and $a_0 = \|\Delta_s\| = \|h\|$. Applying Lemma F.1 with $g_n = 0$ yields, for $n \le N$,

$$a_n \le \lambda_N^n \|h\|, \tag{F.8}$$

We now bound the Wasserstein distance between the conditional laws of X_{s+n} and X'_{s+n} given \mathcal{F}_s by using the synchronous coupling between X_{s+n} and X'_{s+n} (which share the same noise sequence $(W_u)_{u>s}$):

$$W_1\big(\mathcal{L}(X_{s+n}\mid\mathcal{F}_s),\ \mathcal{L}(X_{s+n}'\mid\mathcal{F}_s)\big)\leq \mathbb{E}\big(\|X_{s+n}-X_{s+n}'\|\mid\mathcal{F}_s\big)\leq \lambda_N^n\,\|h\|.$$

Therefore, for any horizon $T \ge s + 1$,

$$\sup_{s+1 \le t \le T} W_1(\mathcal{L}(X_t \mid \mathcal{F}_s), \ \mathcal{L}(X_t' \mid \mathcal{F}_s)) \le \|h\| \lambda_{T-s}^{T-s} = \begin{cases} \|h\| (L(1 + \log(T-s))^{T-s} & \text{if } \alpha = 1, \\ \|h\| (L\zeta(\alpha))^{T-s} & \text{if } \alpha > 1. \end{cases}$$
(F.9)

The behaviour of the bound crucially depends on α and L: if $\alpha > 1$ and $L\zeta(\alpha) < 1$, the effect of the perturbation decays exponentially fast with T - s; if $\alpha > 1$ and $L\zeta(\alpha) > 1$, the effect of the perturbation grows exponentially fast with T - s. In both case, higher values of α (faster memory decay) lead to better dependence properties.

Noise perturbation. Fix $s \ge 1$ and let $\mathcal{F}_{s-1} := \sigma(X_1, \ldots, X_{s-1}, W_1, \ldots, W_{s-1})$. Assume the two systems agree up to time s except that at time s we have

$$W_s' = W_s + \eta$$

with $\eta \neq 0$, and $W'_u = W_u$ for $u \neq s$. Again define $\Delta_t := X_t - X'_t$ for $t \geq s$. Subtracting the two recursions gives for $t \geq s$:

$$\Delta_{t+1} = \sum_{u=s}^{t} \frac{b(X_u) - b(X_u')}{(t-u+1)^{\alpha}} + \frac{W_s - W_s'}{(t-s+1)^{\alpha}}.$$
 (F.10)

Taking norms and using Lipschitzness,

$$\|\Delta_{t+1}\| \le L \sum_{u=s}^{t} \frac{\|\Delta_{u}\|}{(t-u+1)^{\alpha}} + \frac{\|\eta\|}{(t-s+1)^{\alpha}}.$$

Set n := t - s + 1 and $a_n := \mathbb{E}(\|\Delta_{s+n}\| \mid \mathcal{F}_{s-1})$. Note $a_0 = 0$ (since $X_s = X_s'$). Apply Lemma F.1 with $g_n := \|\eta\| n^{-\alpha}$ to obtain, for $n \le N$,

$$a_n \le \sum_{j=1}^n \|\eta\| j^{-\alpha} \lambda_N^{n-j} \le \|\eta\| \times \frac{\lambda_N^n - 1}{\lambda_N - 1},$$
 (F.11)

where we consider $\lambda_N \neq 1$ for simplicity.

Bounding the Wasserstein distance as before yields, for any horizon $T \ge s + 1$,

$$\sup_{s+1 \le t \le T} W_1 \left(\mathcal{L}(X_t \mid \mathcal{F}_{s-1}), \ \mathcal{L}(X_t' \mid \mathcal{F}_{s-1}) \right) \le \begin{cases} \|\eta\| \frac{(L(1 + \log(T - s)))^{T - s} - 1}{L(1 + \log(T - s)) - 1}, & \text{if } \alpha = 1, \\ \|\eta\| \frac{(L\zeta(\alpha))^{T - s} - 1}{L\zeta(\alpha) - 1}, & \text{if } \alpha > 1. \end{cases}$$
(F.12)

Drift perturbation. To consider drift perturbations, we write the drift as b_{θ} where θ is a parameter. In addition to assuming that b_{θ} is uniformly L-Lipschitz for all θ , we also assume that it is M-Lipschitz in θ uniformly in x, that is, for all $x, x' \in \mathbb{R}^d$ and θ, θ' ,

$$||b_{\theta}(x) - b_{\theta'}(x')|| \le L ||x - x'|| + M ||\theta - \theta'||.$$
 (F.13)

Consider θ, θ' and the two systems with drifts b_{θ} and $b_{\theta'}$ respectively:

$$X_{t+1} = \sum_{u=1}^{t} K(t, u) \left(b_{\theta}(X_u) + W_u \right), \tag{F.14}$$

$$X'_{t+1} = \sum_{u=1}^{t} K(t, u) \left(b_{\theta'}(X'_u) + W_u \right). \tag{F.15}$$

As before, we will bound the Wasserstein distance between X_t and X_t' by using the synchronous coupling. Assuming that the two sequences share the same noise sequence (W_u) , we define $\Delta_t = X_t - X_t'$ and obtain, using (F.13), for $t \leq T$

$$\|\Delta_{t+1}\| \le L \sum_{u=1}^{t} \frac{\|\Delta_{u}\|}{(t-u+1)^{\alpha}} + M\|\theta - \theta'\|H_{\alpha}(T).$$
 (F.16)

Setting $a_n = \|\Delta_n\|$ and $g_n = M\|\theta - \theta'\|H_\alpha(T)$ with $a_0 = 0$, we can apply Lemma F.1 as before to obtain, for $t \leq T$,

$$W_{1}(\mathcal{L}(X_{t}), \mathcal{L}(X_{t}')) \leq M \|\theta - \theta'\| \begin{cases} (1 + \log T) \frac{(L(1 + \log T))^{t} - 1}{L(1 + \log T) - 1}, & \text{if } \alpha = 1, \\ \zeta(\alpha) \frac{(L\zeta(\alpha))^{t} - 1}{L\zeta(\alpha) - 1}, & \text{if } \alpha > 1 \end{cases}$$
(F.17)

where we used (F.3) to bound $H_{\alpha}(T)$.

F.2 Examples for task selection assumptions

In this section, we check that the examples of Section 3.1 in the main text satisfy Assumptions 3 and 4. These are lengthy but mostly straightforward calculations, which we sketch to illustrate how to verify the assumptions in practice.

Example F.1 (Linear regression). We consider the linear regression example of Section 3.1 in the main text and check that it satisfies Assumptions 3 and 4. Fix a true task $\theta^* \in \mathbb{R}^d$. For t = 1, ..., T, consider $q_t \sim \mathcal{N}(0, \sigma_q^2 I_d)$ and noise $\epsilon_t \sim \mathcal{N}(0, \sigma_\epsilon^2)$ i.i.d., and $y_t = q_t^{\mathsf{T}} \theta^* + \epsilon_t$, $z_t = (q_t, y_t)$, $X = \{z_t\}_{t=1}^T$. Define $Q \in \mathbb{R}^{T \times d}$ has rows q_t^{T} and $Y = (y_t)_{t=1}^T$, and, for any parameter $\theta \in \mathbb{R}^d$,

$$\ell_T(\theta) := \log p_T(X \mid \theta) = -\frac{1}{2\sigma^2} ||Y - Q\theta||_2^2 + \text{const},$$

where the constant term depends on Q but not on θ

Let us begin with the tail behavior. Both q_t and $y_t = q_t^T \theta^* + \epsilon_t$ are sub-Gaussian; hence for some c > 0 and all $R \ge 1$,

$$\mathbb{P}(\exists t \le n, ||z_t|| \ge R) \le \operatorname{poly}(n) e^{-cR^2} \le \operatorname{poly}(n) e^{-R}.$$

For the tail condition on the likelihood, let $\Delta = \theta - \theta_0$ and $r_0 := Y - Q\theta_0$. Then

$$\ell_T(\theta) - \ell_T(\theta_0) = -\frac{1}{2\sigma_{\epsilon}^2} (\|Q\Delta\|_2^2 - 2\Delta^{\mathsf{T}} Q^{\mathsf{T}} r_0)$$

Now, by e.g., Wainwright (2019, Thm. 6.1), for T large enough, there is c > 0 constant such that, with probability at least $1 - e^{-cT}$, $\|Q\Delta\| \ge c\sqrt{T} \|\Delta\|$ and $\|Q^{\top}r_0\| \le c^{-1}\sqrt{T} \|r_0\|$. Hence, uniformly over $\|\theta\| \ge R$ (so $\|\Delta\| \ge R - \|\theta_0\|$).

$$\ell_T(\theta) - \ell_T(\theta_0) \le -\frac{c^2 T}{2\sigma_z^2} \|\Delta\|^2 + \frac{c^{-1} \sqrt{T}}{\sigma_z^2} \|\Delta\| \|r_0\|.$$

For all R larger than a constant multiple of $||r_0||/\sqrt{T} + ||\theta_0||$, the right-hand side is negative; thus $\sup_{\|\theta\| \ge R} p_T(X \mid \theta) < p_T(X \mid \theta_0)$. Since $||r_0||$ is sub-Gaussian and the norm bounds above hold with probability at least $1 - e^{-cn} \ge 1 - e^{-cR}$ for $R \ge T$, we obtain, for all $R \ge T$,

$$\mathbb{P}\bigg(\sup_{\|\theta\|\geq R} \mathrm{p}_T(X\mid\theta) \ \geq \ \mathrm{p}_T(X\mid\theta_0)\bigg) \ \leq \ \mathrm{poly}(T)\,e^{-R}.$$

We now consider the moment condition. Then, for any reference θ_0 ,

$$\sup_{\theta} \frac{\mathrm{p}_T(X \mid \theta)}{\mathrm{p}_T(X \mid \theta_0)} = \exp\left(\sup_{\theta} \{\ell_T(\theta) - \ell_T(\theta_0)\}\right) \le \exp\left(\frac{1}{2\sigma_{\epsilon}^2} \|Y - Q\theta_0\|_2^2\right),$$

Therefore, we have

$$\log^{2} \sup_{\theta} \frac{p_{T}(X \mid \theta)}{p_{T}(X \mid \theta_{0})} \leq C (\|Q(\theta^{*} - \theta_{0})\|_{2}^{2} + \|\epsilon\|_{2}^{2})^{2},$$

and using Gaussian moment bounds

$$\mathbb{E}\left[\log^2 \sup_{\theta} \frac{\mathrm{p}_T(X \mid \theta)}{\mathrm{p}_T(X \mid \theta_0)}\right] \leq \mathrm{poly}(n) \left(1 + \|\theta^* - \theta_0\|_2^4\right) = \mathrm{poly}(n).$$

We finally check the local regularity condition. For any t and θ, θ' ,

$$\log \frac{\mathbf{p}_{t}(y_{t} \mid q_{1:t}, y_{1:t-1}, \theta)}{\mathbf{p}_{t}(y_{t} \mid q_{1:t}, y_{1:t-1}, \theta')} = -\frac{1}{2\sigma_{\epsilon}^{2}} [(y_{t} - \theta^{\mathsf{T}} q_{t})^{2} - (y_{t} - {\theta'}^{\mathsf{T}} q_{t})^{2}].$$

Assuming that $||q_{1:t}||_{\infty}$, $|y_{1:t}| \le R$ and $||\theta||$, $||\theta'|| \le R$ (with $R \ge 1$) and using that $(a-b)^2 - (a-c)^2 = (c-b)(2a-b-c)$, we have

$$\left|\log \frac{\mathrm{p}_t(y_t \mid q_{1:t}, y_{1:t-1}, \theta)}{\mathrm{p}_t(y_t \mid q_{1:t}, y_{1:t-1}, \theta')}\right| = \frac{1}{2\sigma_\epsilon^2} \left|(\theta - \theta')^\top q_t\right| \left|2y_t - (\theta + \theta')^\top q_t\right| \leq \frac{1}{\sigma_\epsilon^2} R^3 \|\theta - \theta'\|,$$

so the condition holds.

Example F.2 (Ornstein–Uhlenbeck process). We consider the Ornstein–Uhlenbeck (OU) process example of Section 3.1 in the main text and check that it satisfies Assumptions 3 and 4. For simplicity, we consider the one-dimensional case d=1; the extension to d>1 with diagonal diffusion is straightforward. We consider tasks $\theta=(\mu,\tau)$ where $\mu\in\mathbb{R}$ and $\tau\in[\overline{\tau},\underline{\tau}]$ with $0<\overline{\tau}\leq\underline{\tau}<\infty$. Given θ , the Ornstein–Uhlenbeck (OU) SDE

$$dX_t = \tau(\mu - X_t) dt + \sigma dW_t$$

is observed at regular times $t_r = r \Delta_t$ (r = 1, ..., n). We write $x_r := X_{t_r}$ and $X = \{x_r\}_{r=1}^n$. The Markov transition is Gaussian with mean

$$m_{\theta}(x) := \mu + e^{-\tau \Delta_t}(x - \mu) = e^{-\tau \Delta_t}x + (1 - e^{-\tau \Delta_t})\mu$$

and variance $v_{\theta} \coloneqq \operatorname{Var}(x_r \mid x_{r-1}, \theta) = \sigma^2 \frac{1 - e^{-2\tau \Delta_t}}{2\tau}$. For any path $x_{1:n}$, define $\ell_n(\theta) \coloneqq \log p_n(X \mid \theta)$. Recall $\theta = (\mu, \tau)$ with $\tau \in [\overline{\tau}, \underline{\tau}]$, discretization step Δ_t , and

$$m_{\theta}(x) = \mu + \rho_{\tau}(x - \mu) = \rho_{\tau}x + (1 - \rho_{\tau})\mu, \qquad v_{\theta} = \sigma^{2}\frac{1 - \rho_{\tau}^{2}}{2\tau}, \qquad \rho_{\tau} := e^{-\tau\Delta_{t}}.$$

Fix a reference $\theta_0 = (\mu_0, \tau_0)$, write $m_0 := m_{\theta_0}$, $v_0 := v_{\theta_0}$, and let $X = (x_1, \dots, x_n)$ with x_r the OU samples at times $r\Delta_t$. The one–step densities are Gaussian, hence

$$\log \frac{p_n(X \mid \theta)}{p_n(X \mid \theta_0)} = \sum_{r=1}^n \left\{ -\frac{1}{2} \log \frac{v_\theta}{v_0} - \frac{\left(x_r - m_\theta(x_{r-1})\right)^2}{2v_\theta} + \frac{\left(x_r - m_0(x_{r-1})\right)^2}{2v_0} \right\}. \tag{F.18}$$

Let us begin with the tail behavior. Each one-step innovation $x_r - m_{\theta}(x_{r-1})$ is Gaussian with variance v_{θ} and

$$0 < v_{\min} \le v_{\theta} \le v_{\max} < \infty, \quad v_{\min} := \sigma^2 \frac{1 - e^{-2\underline{\tau}\Delta_t}}{2\tau}, \ v_{\max} := \sigma^2 \frac{1 - e^{-2\overline{\tau}\Delta_t}}{2\overline{\tau}}.$$

Moreover, if x_{r-1} satisfies $|x_{r-1}| \le R$, then $m_{\theta}(x_{r-1})$ also satisfies $|m_{\theta}(x_{r-1})| \le \rho_{\underline{\tau}}R + (1 - \rho_{\underline{\tau}})|\mu|$. Hence, there exists c > 0 depending only on $(\Delta_t, \overline{\tau}, \underline{\tau}, \sigma)$ and the law of x_0 such that, for all $R \ge 1$,

$$\mathbb{P}\Big(\exists r \le n, |x_r| \ge R\Big) \tag{F.19}$$

$$\leq \mathbb{P}\Big(\exists r \leq n, |x_r - m_{\theta}(x_{r-1})| \geq (1 - \rho_{\underline{\tau}})R - |\mu|\Big) \tag{F.20}$$

$$\leq \operatorname{poly}(n) e^{-cR^2} \leq \operatorname{poly}(n) e^{-R},$$
 (F.21)

for R large enough compared to $|\mu|$.

Let us continue with the tail condition on the likelihood. We have the bound

$$\left|\sum_{r=1}^{n} -\frac{1}{2}\log\frac{v_{\theta}}{v_{0}}\right| \le \frac{n}{2}\log\frac{v_{\text{max}}}{v_{\text{min}}} =: C_{\text{var}} n.$$
 (F.22)

For each r, abbreviate $m := m_{\theta}(x_{r-1})$ and $m_0 := m_0(x_{r-1})$. Using $v_{\theta} \ge v_{\min}$ and $v_0 \ge v_{\min}$

$$-\frac{(x_r-m)^2}{2v_\theta}+\frac{(x_r-m_0)^2}{2v_0}\leq \frac{1}{2v_{\min}}\Big((x_r-m_0)^2-(x_r-m)^2\Big).$$

Expanding the square,

$$(x_r - m_0)^2 - (x_r - m)^2 = -(m - m_0)^2 + 2(x_r - m_0)(m - m_0)$$

Summing over r and applying Cauchy–Schwarz,

$$\sum_{r=1}^{n} \left(-\frac{(x_r - m)^2}{2v_\theta} + \frac{(x_r - m_0)^2}{2v_0} \right) \le -\frac{1}{2v_{\min}} \sum_{r=1}^{n} \Delta_r^2 + \frac{1}{v_{\min}} \left(\sum_{r=1}^{n} (x_r - m_0)^2 \right)^{1/2} \left(\sum_{r=1}^{n} \Delta_r^2 \right)^{1/2}, \tag{F.23}$$

where $\delta_r := m_{\theta}(x_{r-1}) - m_0(x_{r-1})$.

On events where $|x_{1:n}| \leq R$, we have the conditions

$$c\|\mu - \mu_0\| - C(1+R)|\delta_r| \le L(1+R)\|\theta - \theta_0\|$$

for constants c, C, L depending only on $(\overline{\tau}, \underline{\tau}, \Delta_t)$. Therefore, for $\|\mu - \mu_0\|$ larger than a constant multiple of (1 + R), we have

$$\sum_{r=1}^{n} \delta_{r}^{2} \geq n c \|\mu - \mu_{0}\|^{2} \quad \text{and} \quad \left(\sum_{r=1}^{n} \delta_{r}^{2}\right)^{1/2} \leq \sqrt{n} C(1+R) \|\theta - \theta_{0}\|, \tag{F.24}$$

for constants c, C depending only on $(\overline{\tau}, \underline{\tau}, \Delta_t)$.

Combining (F.22), (F.23), and (F.24),

$$\log \frac{p_n(X \mid \theta)}{p_n(X \mid \theta_0)} \le Cn - cn\|\mu - \mu_0\|^2 + \left(\sum_{r=1}^n (x_r - m_0(x_{r-1}))^2\right)^{1/2} \sqrt{n}C(1+R)\|\theta - \theta_0\|,$$
 (F.25)

for constants c, C depending only on $(\overline{\tau}, \underline{\tau}, \Delta_t)$.

Fix $R \ge 1$ and assume that $|x_{1:n}| \le R$: we have shown that it holds with probability at least $1 - \operatorname{poly}(n)e^{-cR^2}$. In that case, $\left(\sum_{r=1}^n (x_r - m_0(x_{r-1}))^2\right)^{1/2}$ in (F.25) is bounded $\mathcal{O}(\sqrt{n}R)$ so the RHS can be made negative for all sufficiently large $\|\theta\|$: more precisely, it is negative for $\|\theta\| \ge R'$ with $R' \ge C(1+R)^2$ for a constant C depending only on $(\overline{\tau}, \underline{\tau}, \Delta_t)$. Since the event we are considering holds with probability at least $1 - \operatorname{poly}(n)e^{-cR^2}$, it means

Moving to the moment condition, by Gaussian moment bounds, (F.18) readily implies

$$\mathbb{E}\left[\log^2 \sup_{\theta} \frac{p_n(X \mid \theta)}{p_n(X \mid \theta_0)}\right] \le C n^2 = \text{poly}(n),$$

that it holds with probability at least $1 - \text{poly}(n)e^{-R'}$. This proves the required tail bound with $R \leftarrow R'$.

which verifies the likelihood-ratio moment condition in Assumption 3.

Finally, we show the local regularity condition. For fixed $x_{1:r-1}$, the conditional density is

$$\log p_r(x_r \mid x_{1:r-1}, \theta) = -\frac{1}{2} \log(2\pi v_\theta) - \frac{\left(x_r - m_\theta(x_{r-1})\right)^2}{2v_\theta}.$$

On sets where $|x_{1:r}| \leq R$, $||\theta|| \leq R$ (so μ, τ bounded) and with $\tau \in [\overline{\tau}, \tau]$, the maps

$$\theta \mapsto m_{\theta}(x_{r-1}) = e^{-\tau \Delta_t} x_{r-1} + \left(1 - e^{-\tau \Delta_t}\right) \mu, \qquad \theta \mapsto v_{\theta} = \sigma^2 \frac{1 - e^{-2\tau \Delta_t}}{2\tau}$$

are smooth with bounded first derivatives: $|\partial_{\mu}m_{\theta}| \leq 1$, $|\partial_{\tau}m_{\theta}| \leq C_R$, $|\partial_{\tau}v_{\theta}| \leq C$, $|\partial_{\mu}v_{\theta}| = 0$. Since $x_r - m_{\theta}(x_{r-1})$ is also bounded by a constant multiple of R on these sets, we obtain, for all θ, θ' with $\|\theta\|, \|\theta'\| \leq R$,

$$\sup_{\substack{|x_{1:r}| \le R \\ \|\theta\| \|\|\theta'\| \le R}} \left| \log \frac{\mathrm{p}_r(x_r \mid x_{1:r-1}, \theta)}{\mathrm{p}_r(x_r \mid x_{1:r-1}, \theta')} \right| \le \mathrm{poly}(R) \|\theta - \theta'\| \, .$$